

External Sorting

External Sorting

- Problem: Sort 1GB of data with 1MB of RAM.
- When a file doesn't fit in memory, there are two stages in sorting:
 1. File is divided into several segments, each of which sorted separately
 2. Sorted segments are *merged*.

(Each stage involves reading and writing the file at least once)

How to sort segments

- Any efficient sorting algorithm (such as Quicksort) can be used to sort segments
- Each sorted segment is called a **run**.
- Each run will be the size of the available memory.

External sort for Large Files

Basic idea

1. Form **runs** (i.e. sorted segments):

- bring as many records as possible to main memory, sort them, save it into a small file on disk.
- Repeat this until we have read all records from the original file and written them as sorted segments (i.e. **runs**) to disk.

2. Do a **multiway merge** of the runs.

External sort for Large Files (cont.)

- How big is each run?
 - As big as the available memory.
- What is the time it takes to create all sorted segments?
 - Ignoring the seek time and assuming b blocks in the file

The time for creating the initial sorted segments is $2b*btt$ (read in the file as segments and write out the runs)

- *Note that the entire file has not been sorted yet.*
- *These are just sorted segments, and*
- *the size of each segment is limited to the size of the available memory used for this purpose.*

Multiway Merging

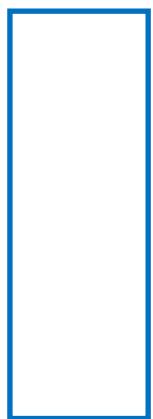
- **K-way merge:** we want to merge K sorted input lists to create a single sorted output list. (**K is the order of a K-way merge**)
- We will adapt the 2-way merge algorithm:
 - Instead of two lists, keep an array of lists: list[0], list[1], ... list[k-1]

Unsorted Example -- Step 1: Create Runs

file

21
110
27
5
108
9
109
11
25
107
2
4
10
22
3
1
4
21

Memory (buffer)
of size 6



Unsorted

file

21
110
27
5
108
9
109
11
25
107
2
4
10
22
3
1
4
21

Example -- Step 1: Create Runs

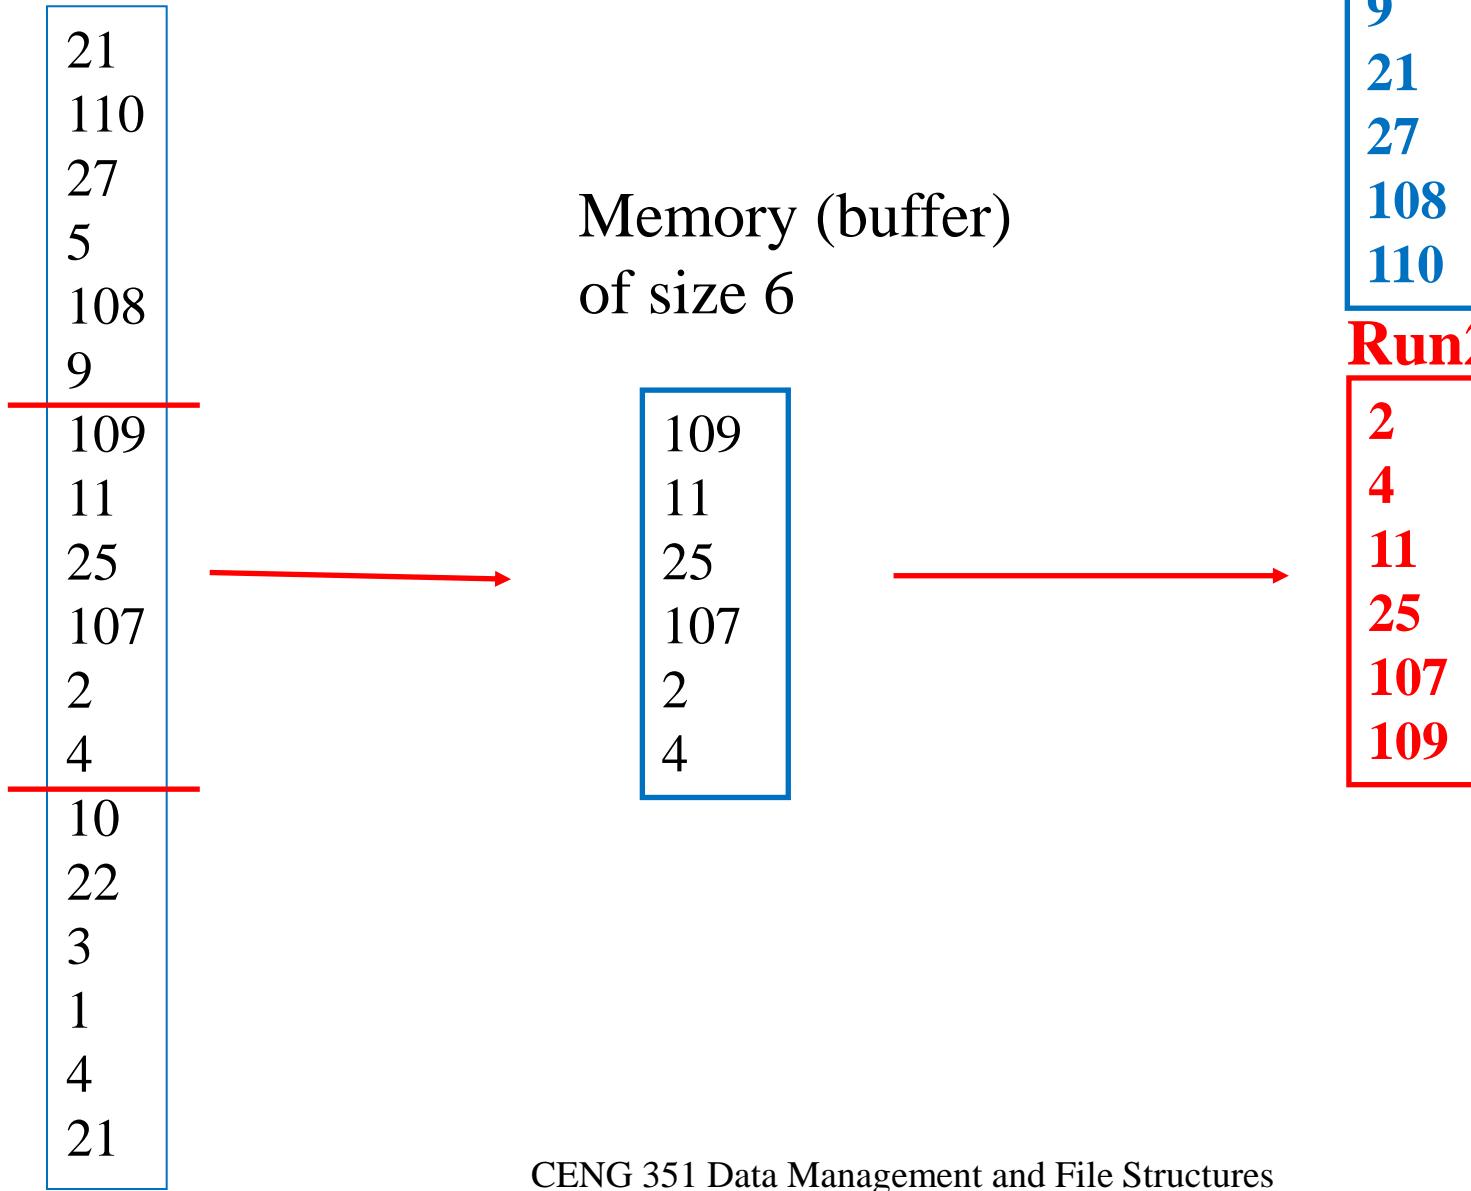
Run1

5
9
21
27
108
110

Memory (buffer)
of size 6

21
110
27
5
108
9

Unsorted file Example -- Step 1: Create Runs



Unsorted

file

21
110
27
5
108
9
109
11
25
107
2
4
10
22
3
1
4
21

Example -- Step 1: Create Runs

Memory (buffer)
of size 6

10
22
3
1
4
21

Run1

5
9
21
27
108
110

Run2

2
4
11
25
107
109

Run3

1
3
4
10
21
22

Run1

5
9
21
27
108
110

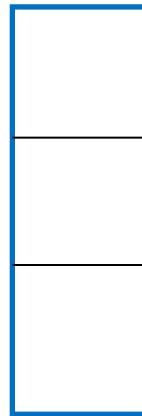
Run2

2
4
11
25
107
109

Run3

1
3
4
10
21
22

Memory (buffer)
of size 6



Run1

5
9
21
27
108
110

Run2

2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

Memory (buffer)
of size 6

5
9
21
27
108
110

Run1

5
9
21
27
108
110

Run2

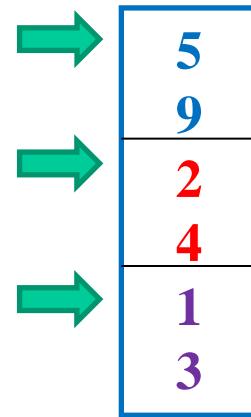
2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

Memory (buffer)
of size 6



Run1

5
9
21
27
108
110

Run2

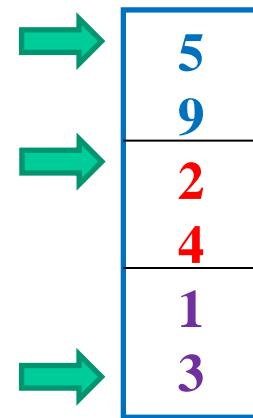
2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

Memory (buffer)
of size 6



1

Run1

5
9
21
27
108
110

Run2

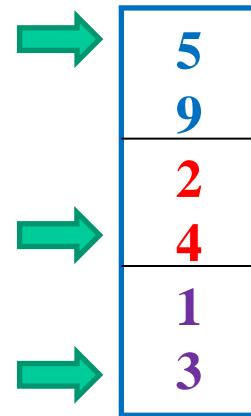
2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

Memory (buffer)
of size 6



Run1

5
9
21
27
108
110

Run2

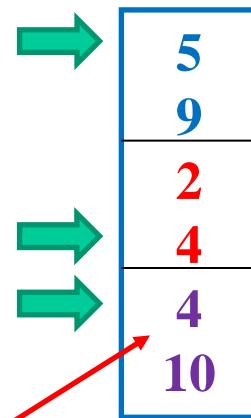
2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

Memory (buffer)
of size 6



Run1

5
9
21
27
108
110

Run2

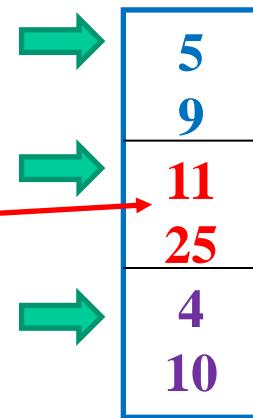
2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

Memory (buffer)
of size 6



Run1

5
9
21
27
108
110

Run2

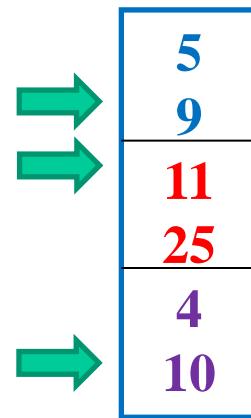
2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

Memory (buffer)
of size 6



1
2
3
4
4
5

Run1

5
9
21
27
108
110

Run2

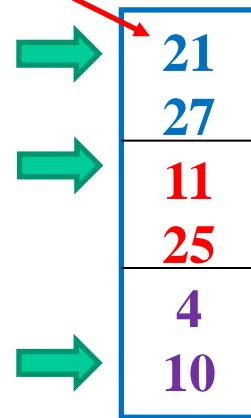
2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

Memory (buffer)
of size 6



1
2
3
4
4
5
9

Run1

5
9
21
27
108
110

Run2

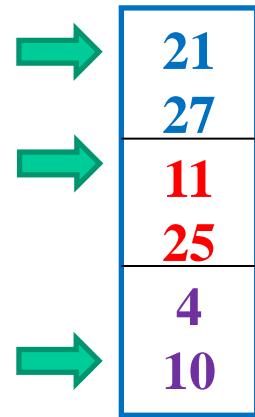
2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

Memory (buffer)
of size 6



Continues...

1
2
3
4
4
5
9

Run1

5
9
21
27
108
110

Run2

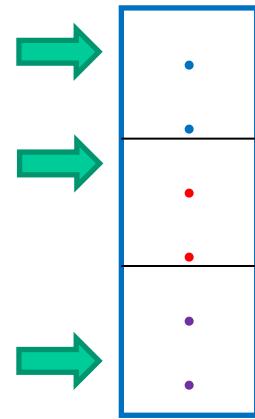
2
4
11
25
107
109

Run3

1
3
4
10
21
22

Example -- Step 2: K-way Merge

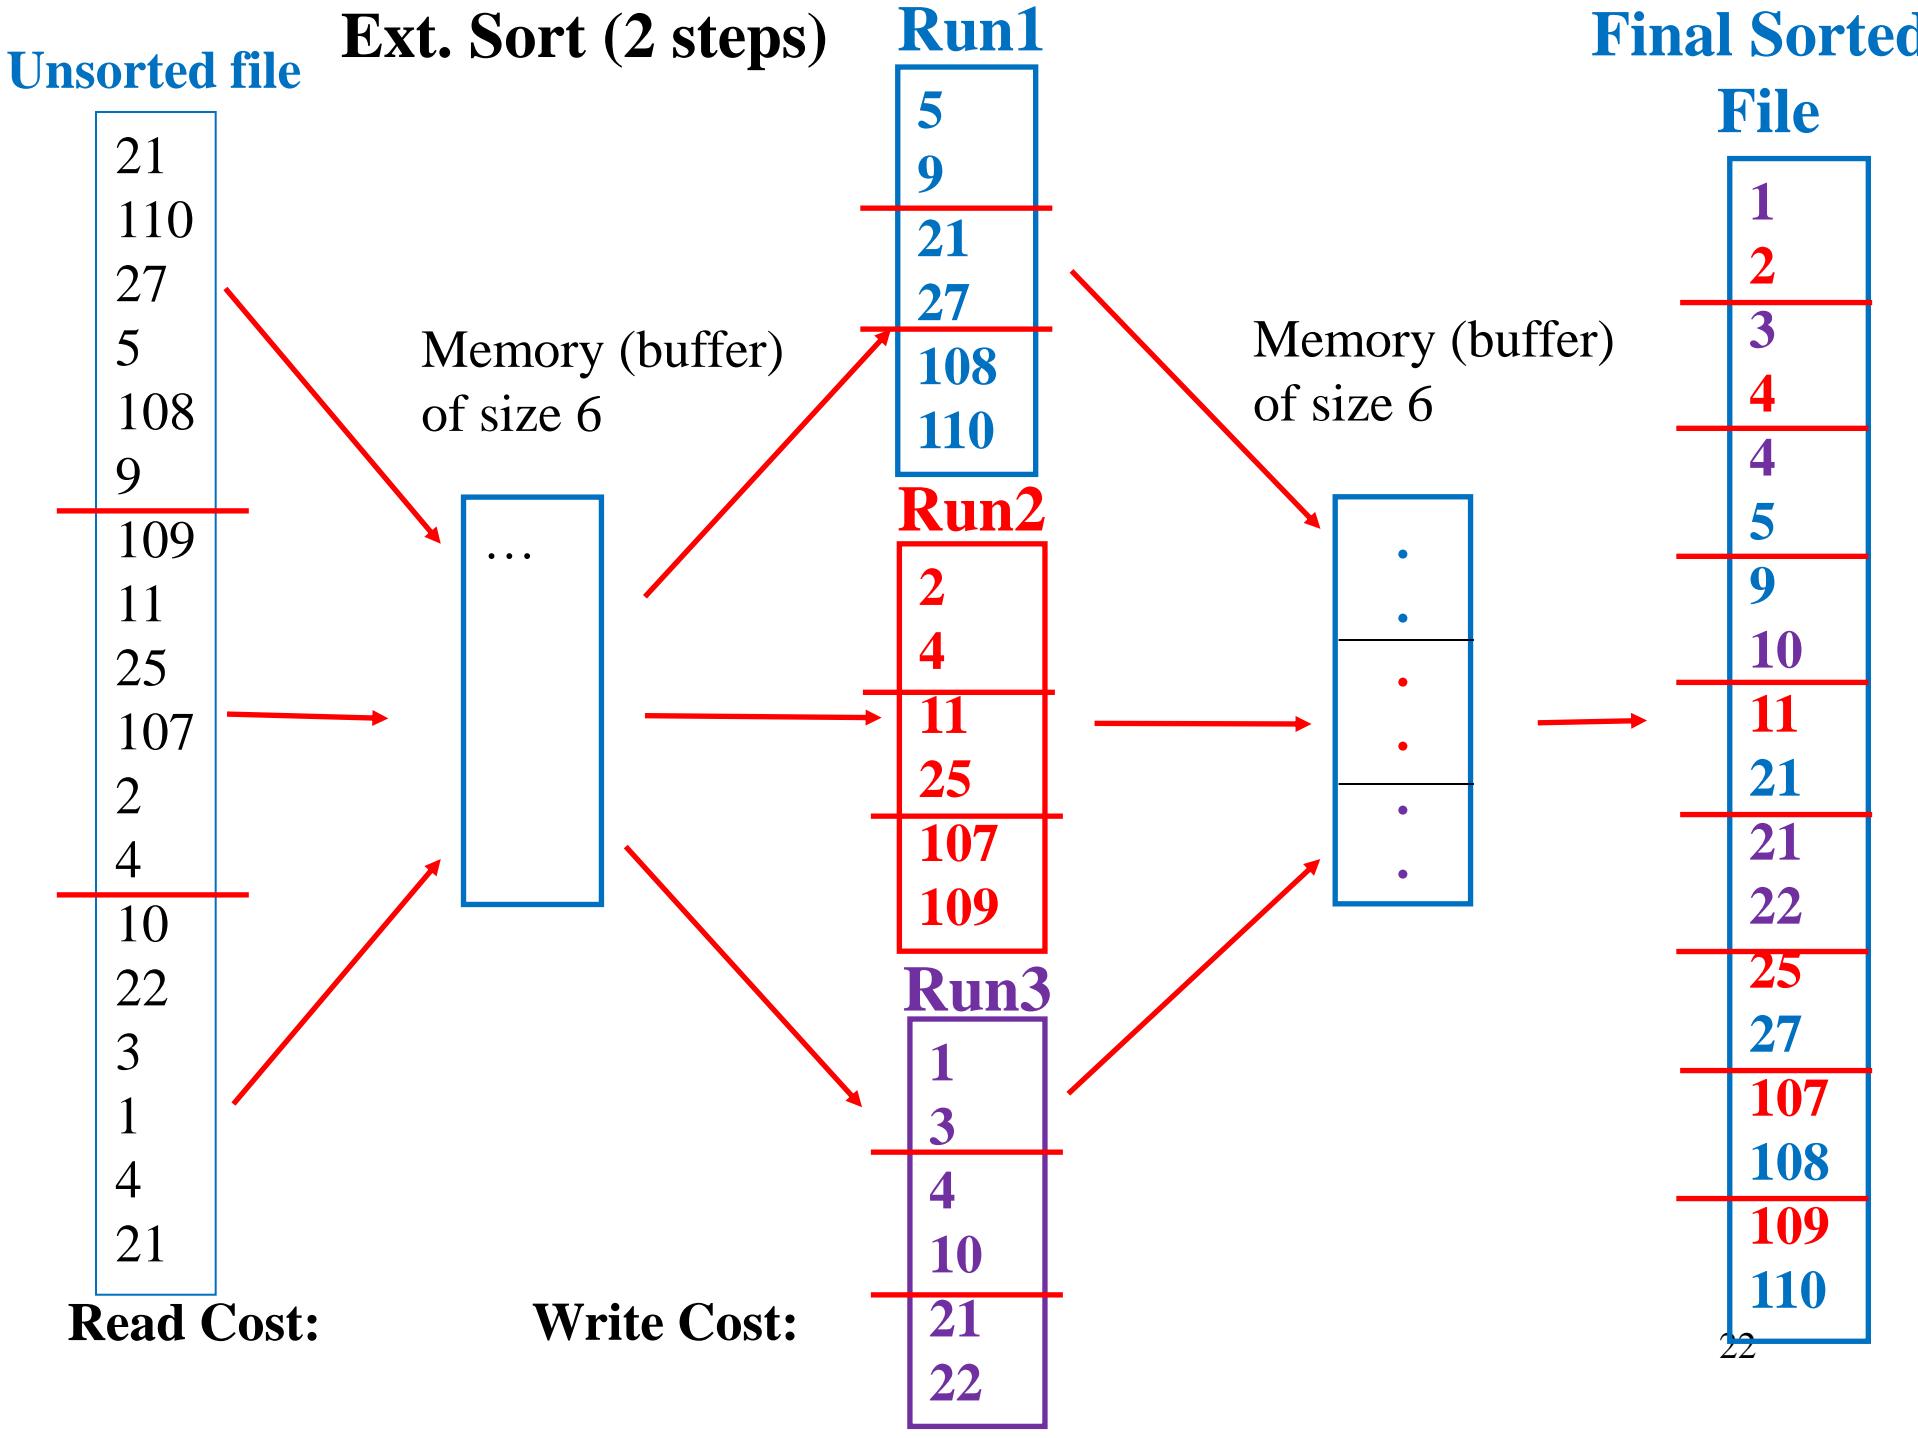
Memory (buffer)
of size 6

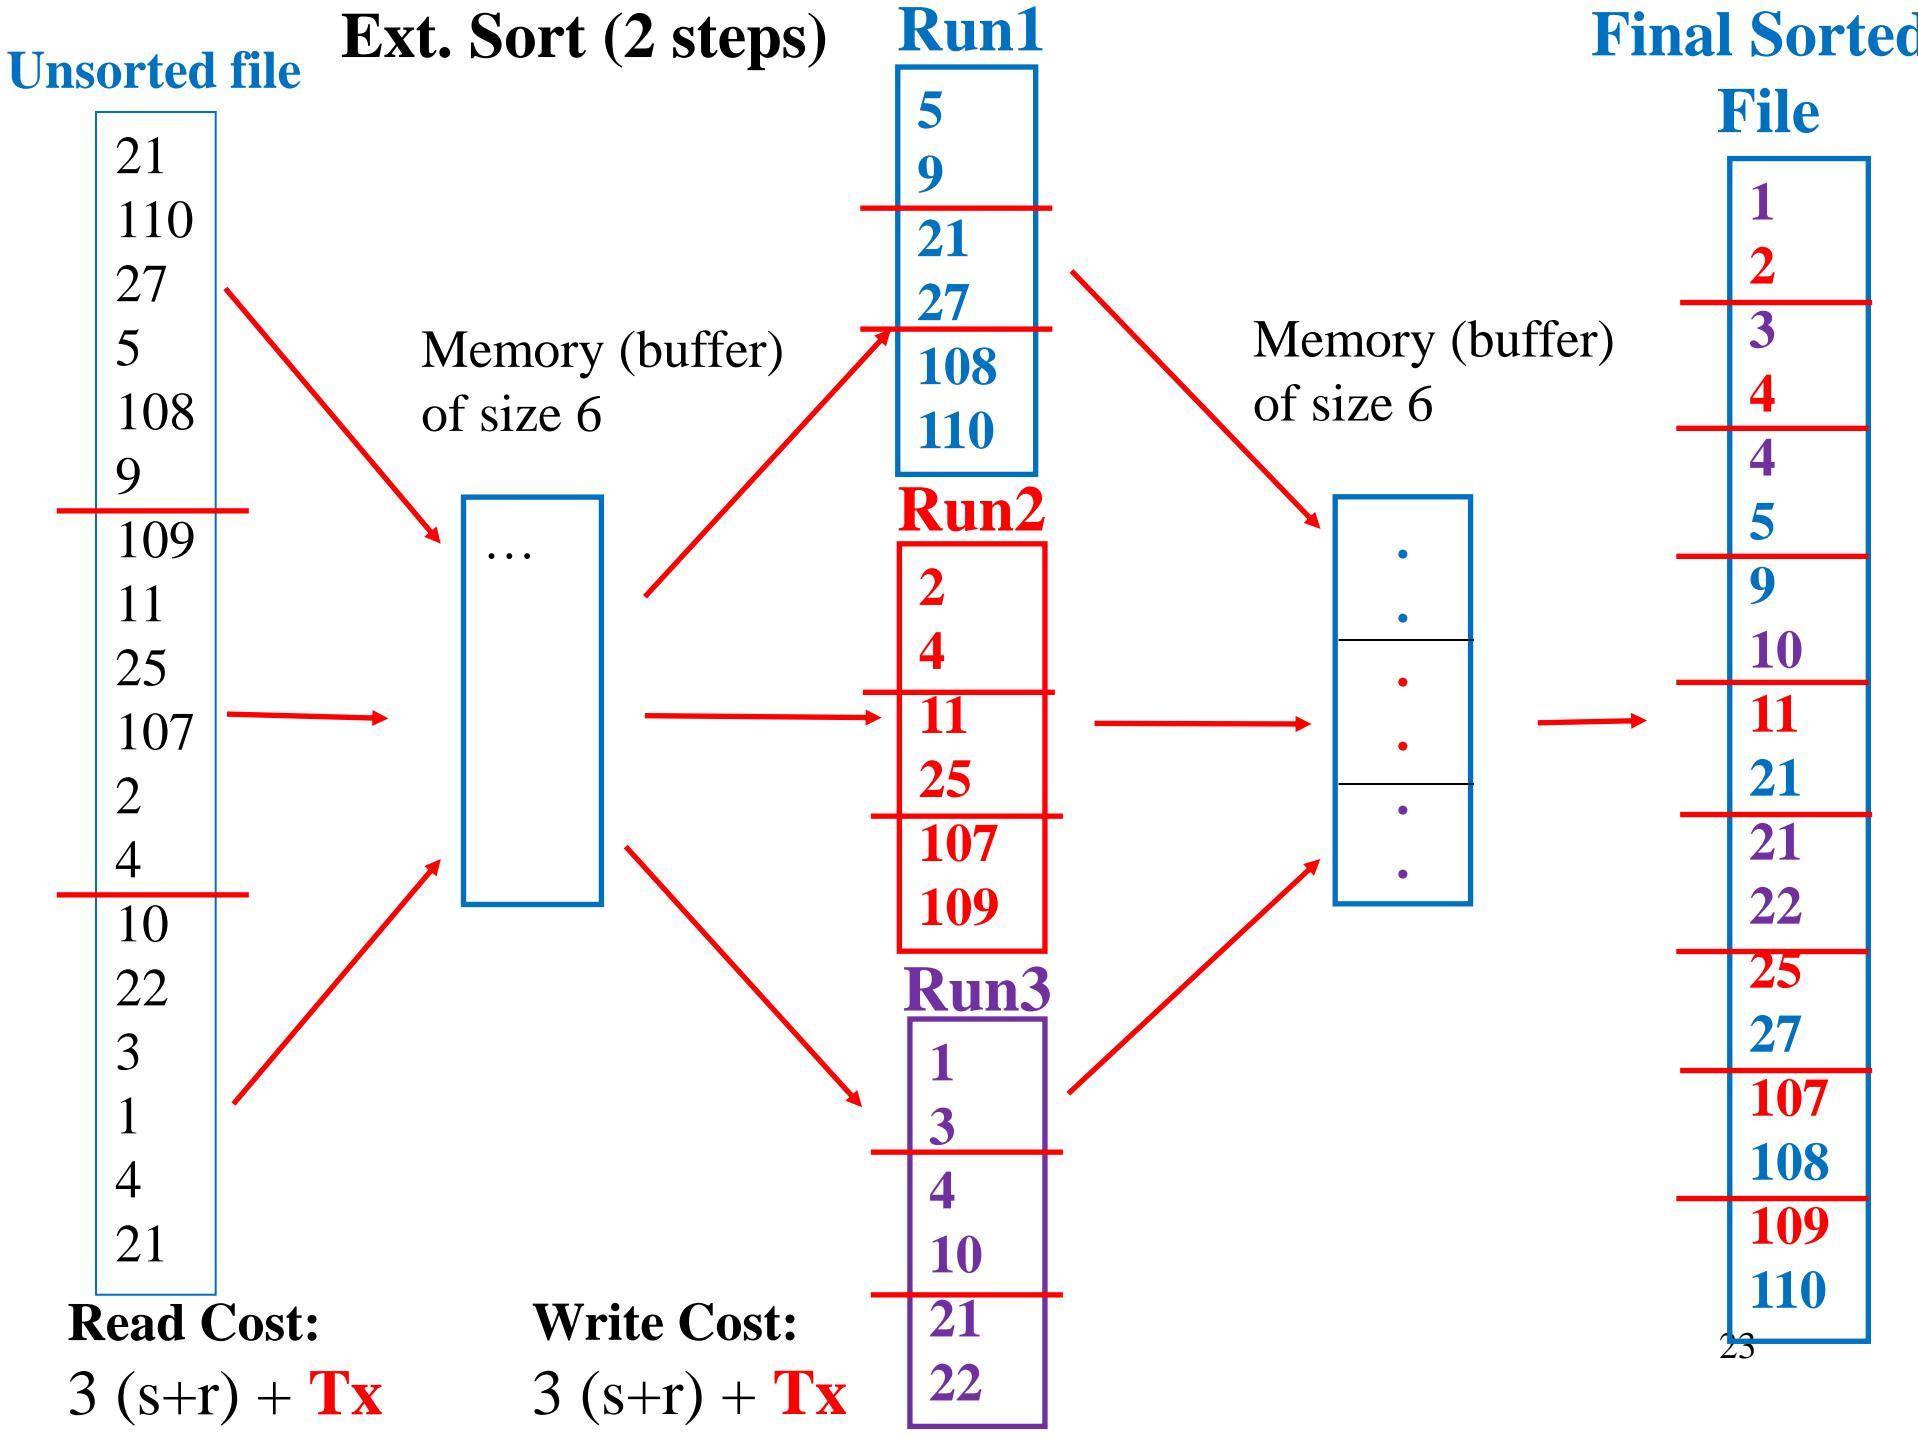


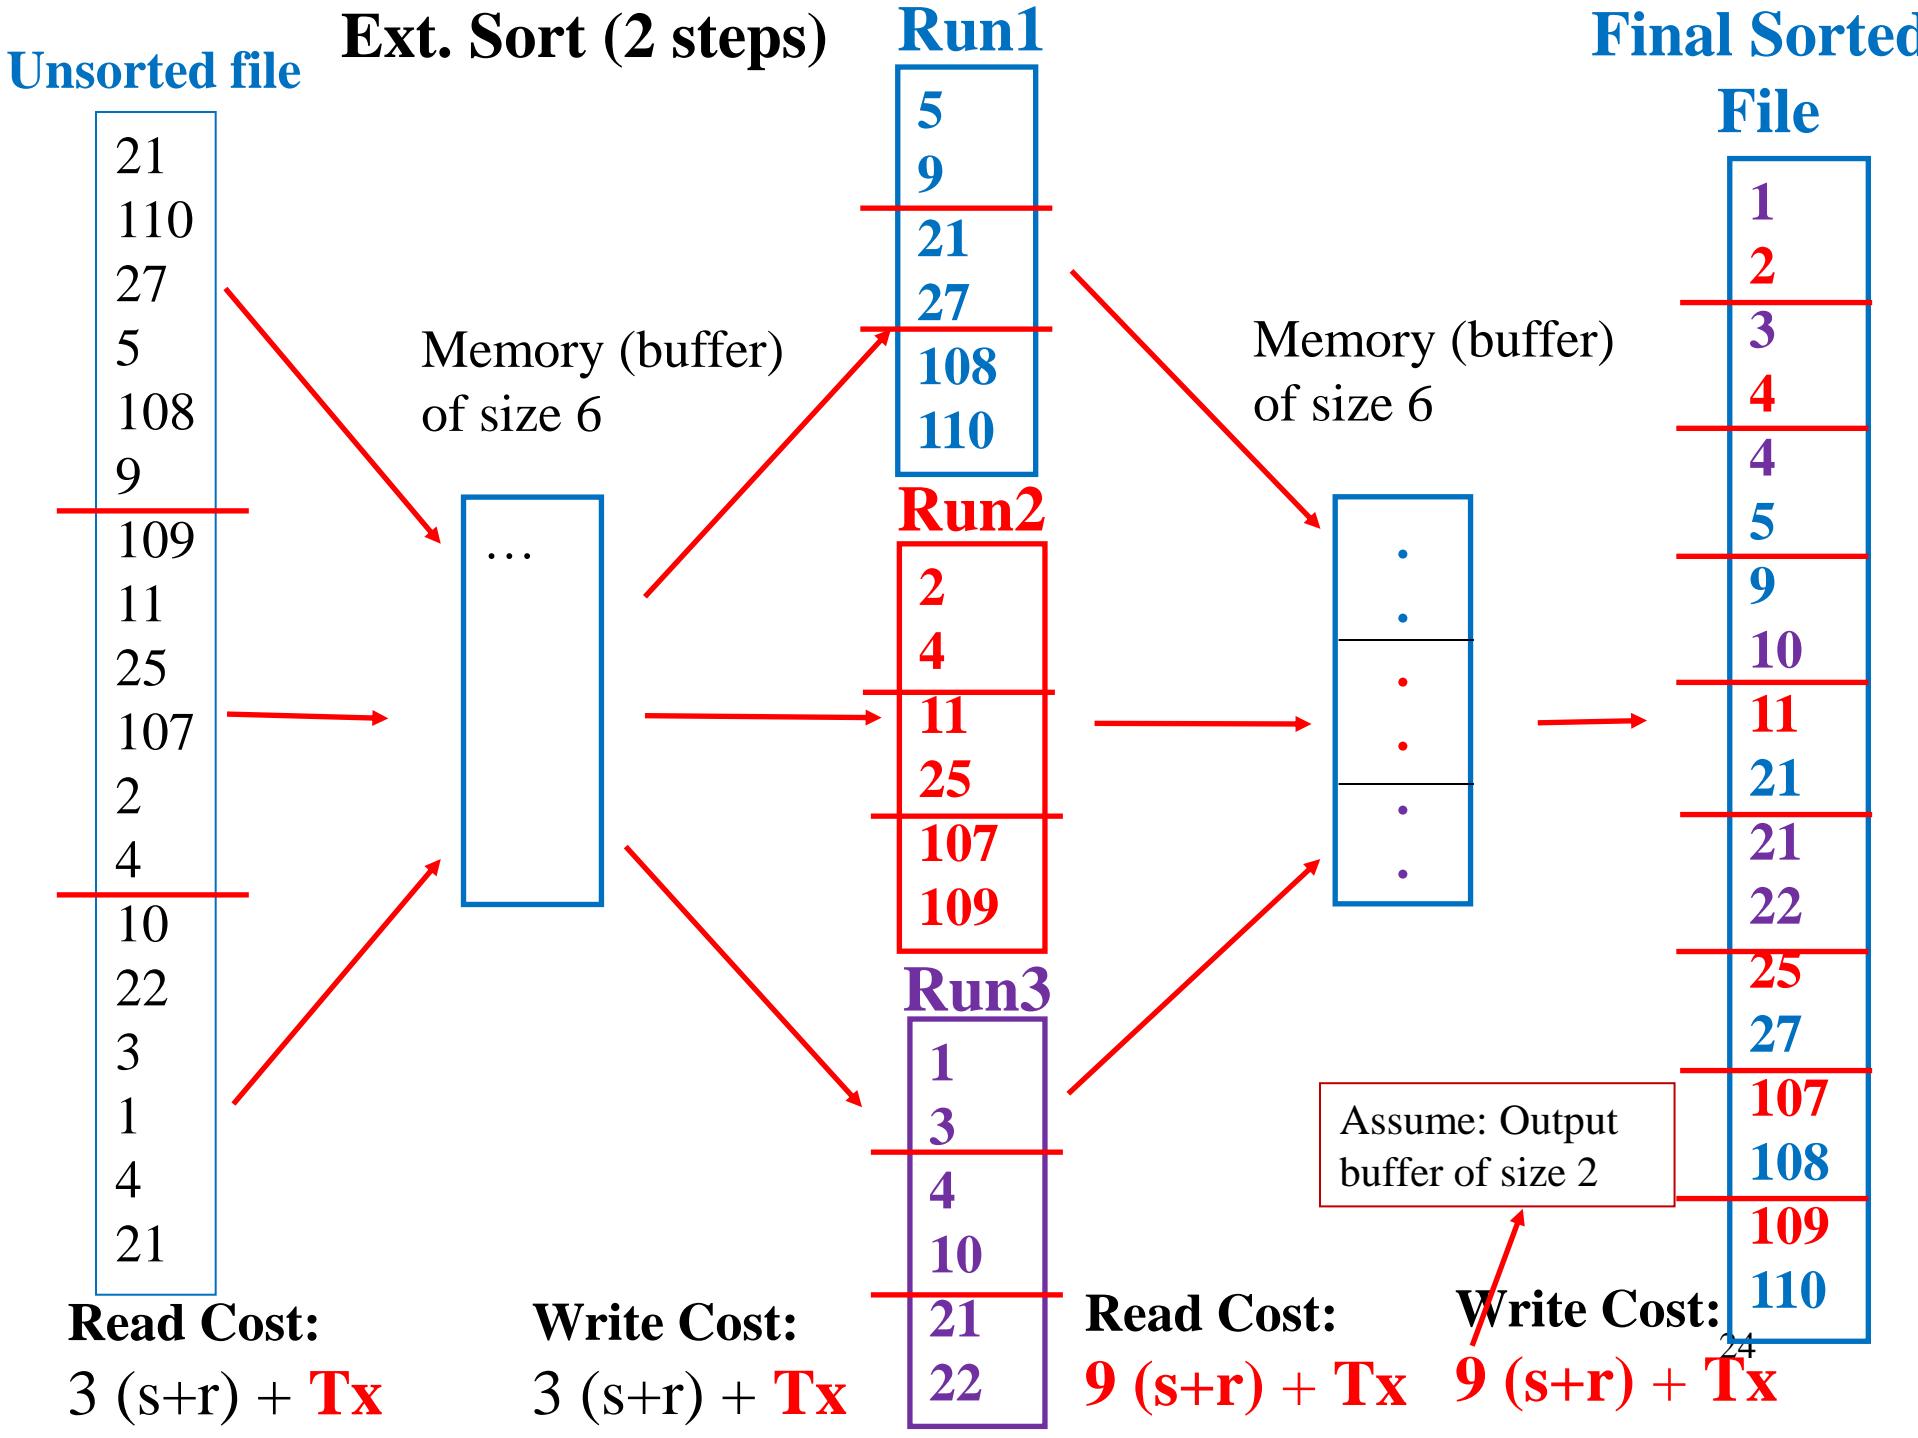
At the end

We obtain the
sorted file

1
2
3
4
4
5
9
10
11
21
21
22
25
27
107
108
109
110







Multiway Merging to sort Large Files

Let us consider the following example:

- File to be sorted:
 - 8,000,000 records
 - $R = 100$ bytes

⇒ Total file size = 800MB

- Memory available as a work area: 10MB (not counting memory used to hold program, other variables, O.S., I/O buffers etc.)

Step 1: Create Runs

Disk I/Os are performed in the following procedures:

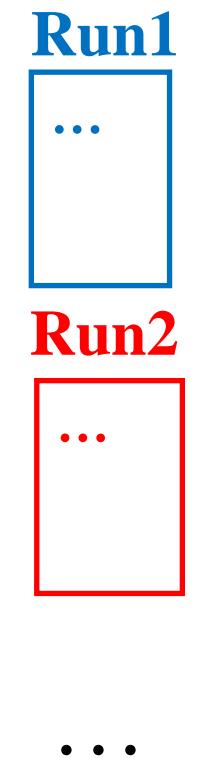
1. Reading records into main memory for sorting
2. Writing runs to disk.

These two steps are done as follows:

- Read 10MB, write 10MB (repeat this 80 times, because file is 800MB). So there will be **80 runs**
 - In terms of basic disk operations, this operation costs:
 - For reading: (**80 seeks +**) a total transfer time of **800 MB (T_x)**
 - Same for writing: (**80 seeks +**) a total transfer time of **800 MB (T_x)**

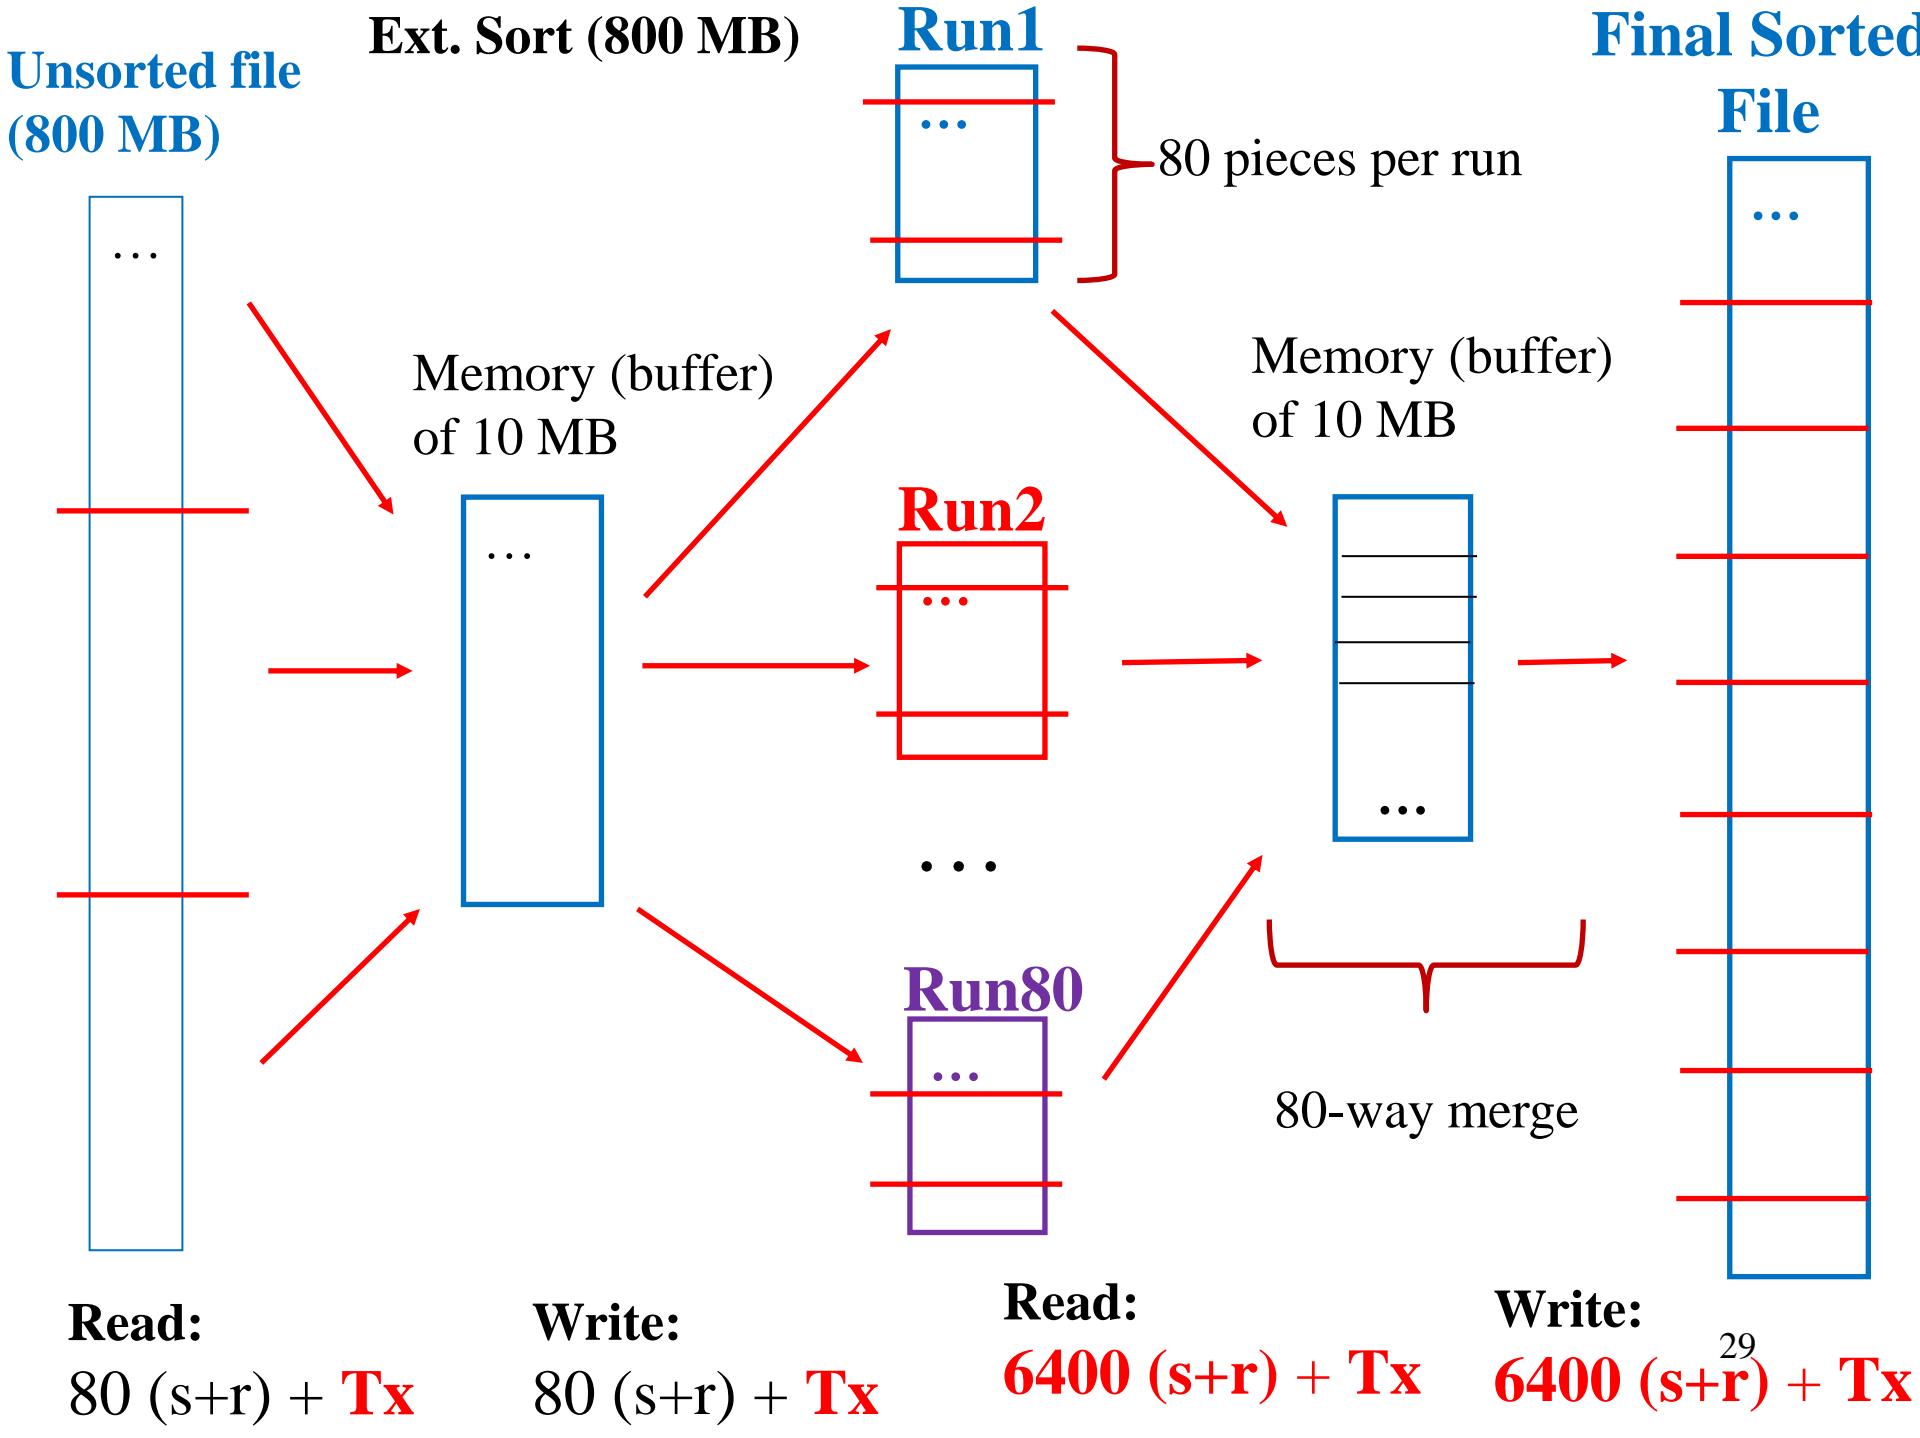
Unsorted
File (800 MB)

Ex. -- Step 1: Create Runs



Step 2: Multiway Merging

- Read runs into memory for merging.
 - Available memory is 10 M and each run size is 10MB.
 - There are 80 runs.
 - We need to do a **80-way merge**.
 - $10M / 80 = 125,000$ bytes are available for each list to be merged.
- Read one piece of each run.
 - How many pieces to be read for each run?
Size of run/size of chunk = $10,000,000/125,000 = 80$ pieces
- Sorting completed in one round:
 - For reading: a total transfer time of 800 MB (**T_x**)
 - Same for writing: a total transfer time of 800 MB (**T_x**)
 - In addition, time for a **high number** of seeks

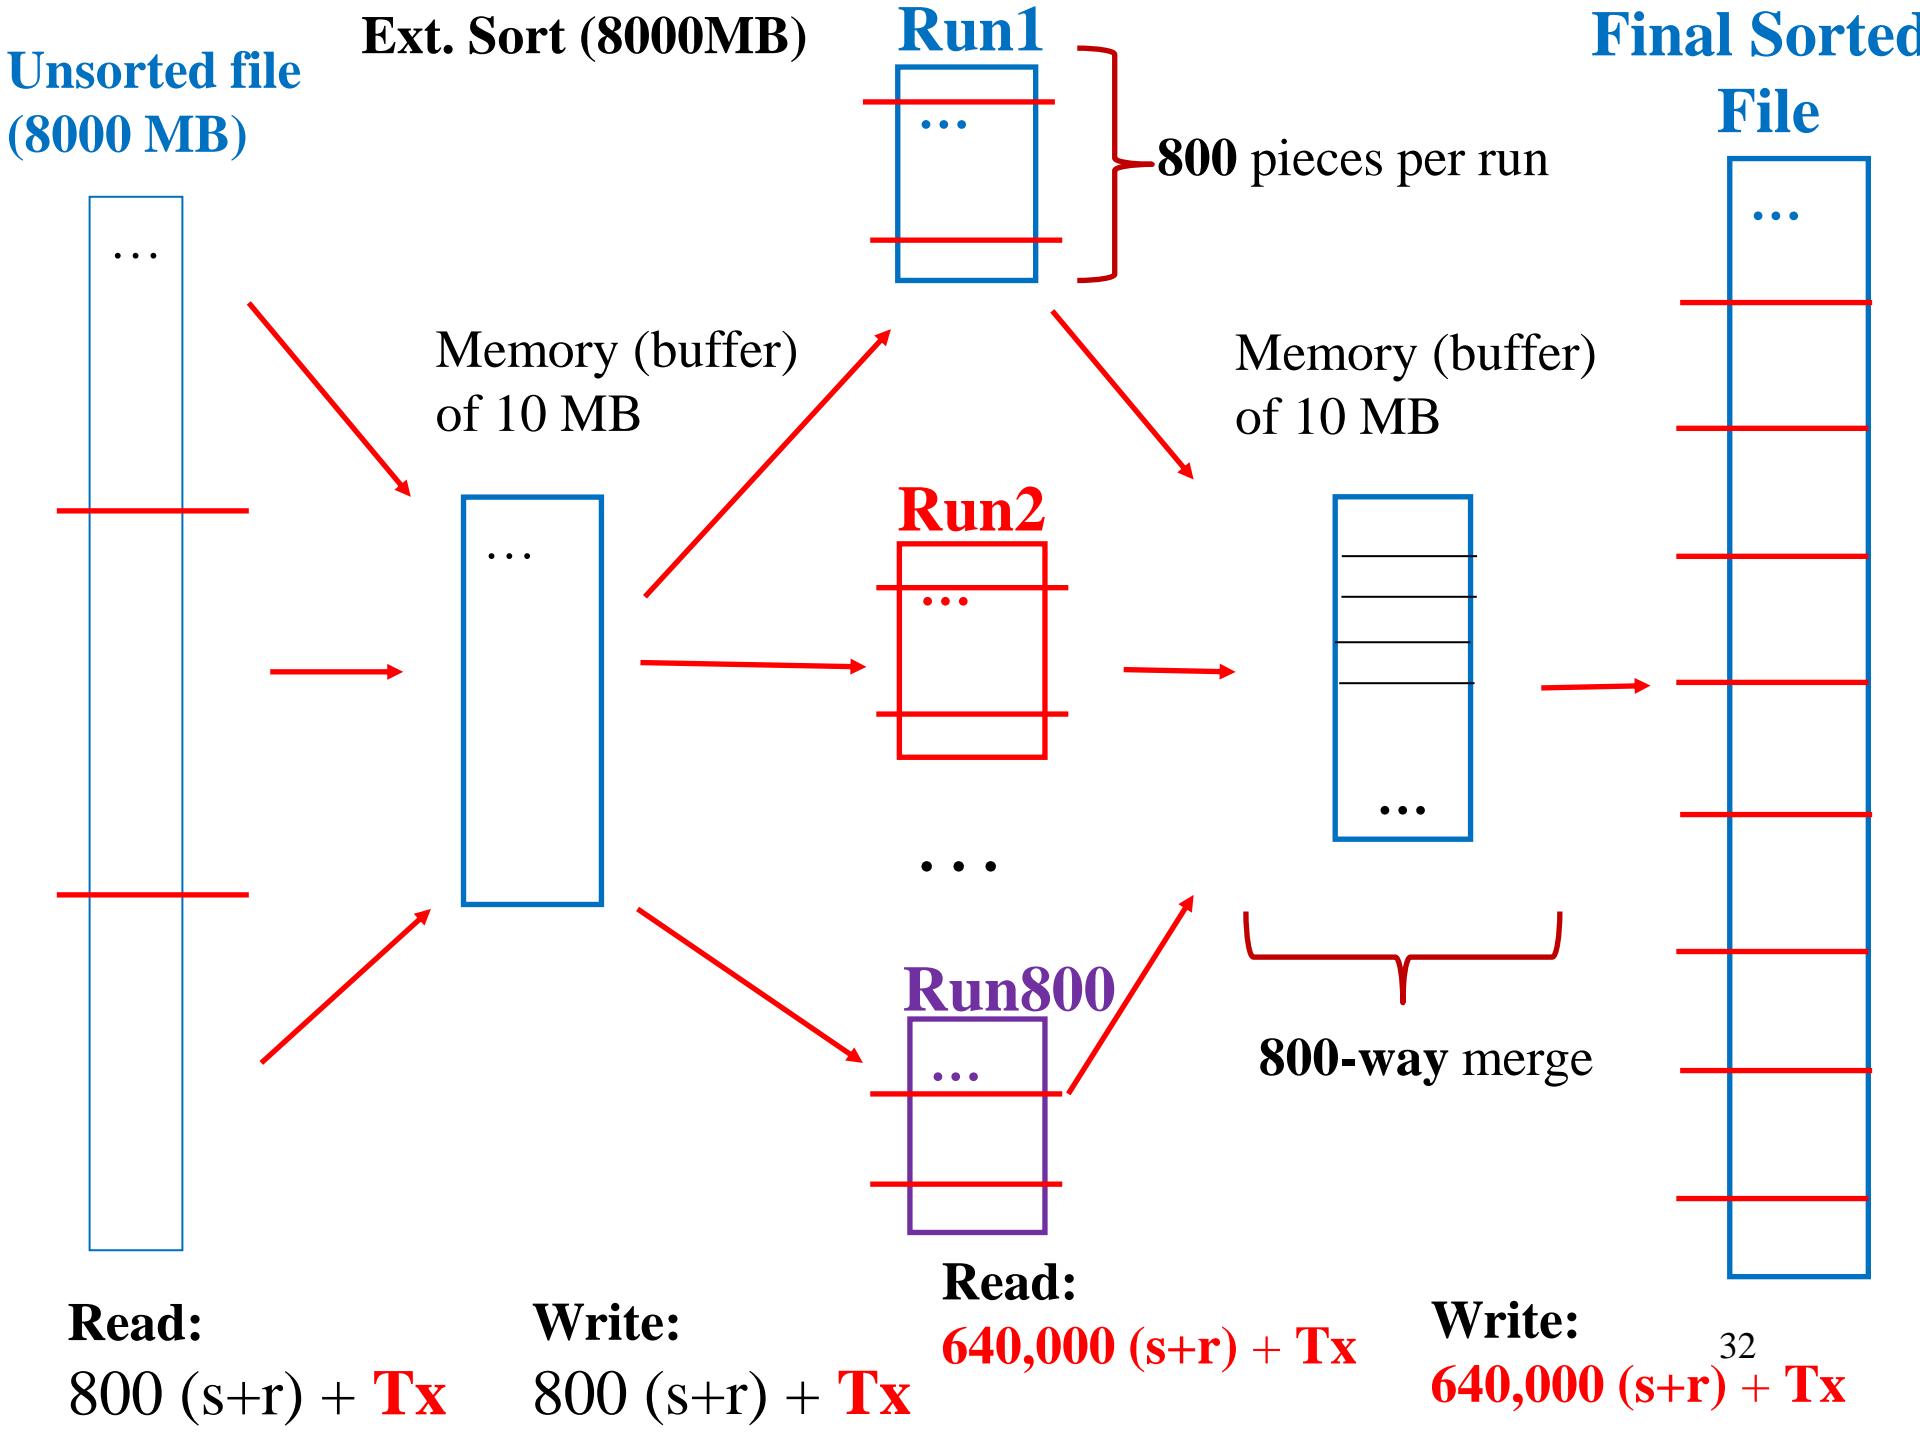


Step 2 (continues)

- Total # of seeks = Total # of pieces (counting all runs):
There are 80 runs and 80 pieces in each run:
 - ✓ 80^2 pieces = 6400 seeks for reading
- Writing the merged records to the big sorted file:
- The number of separate writes closely approximates reads:
 - ✓ 80^2 pieces = 6400 seeks for writing
- So we estimate two seeks (one for reading and one for writing) for each piece: $80 * 80$ pieces therefore
a total of $2 * 6400 = 12800$ seeks are necessary.

Sorting a File that is 10 times larger

- How is the time for merge phase affected if the file is 80 million records?
 - More runs: **800** runs
 - 800-way merge in 10MB memory
 - i.e. divide the memory into 800 pieces (lists).
 - Each piece (list) holds $1/800^{\text{th}}$ of a run
 - So, $800 \text{ runs} * 800 \text{ seeks/run} = 640,000 \text{ seeks}$
 - Total = $2 * 640,000 = \mathbf{1,280,000}$ seeks



The cost of increasing the file size

- In general, for a K-way merge of K runs, the buffer size reserved for each run is
 - $(1/K) * \text{size of memory space} = (1/K) * \text{size of each run}$
- So K seeks are required to read all of the records in each run and K seeks to write records.
- Since there are K runs, merge requires a total of **$2K^2$** seeks.
- Because K is directly proportional to N it also follows that the total cost is an **$O(N^2)$** operation.

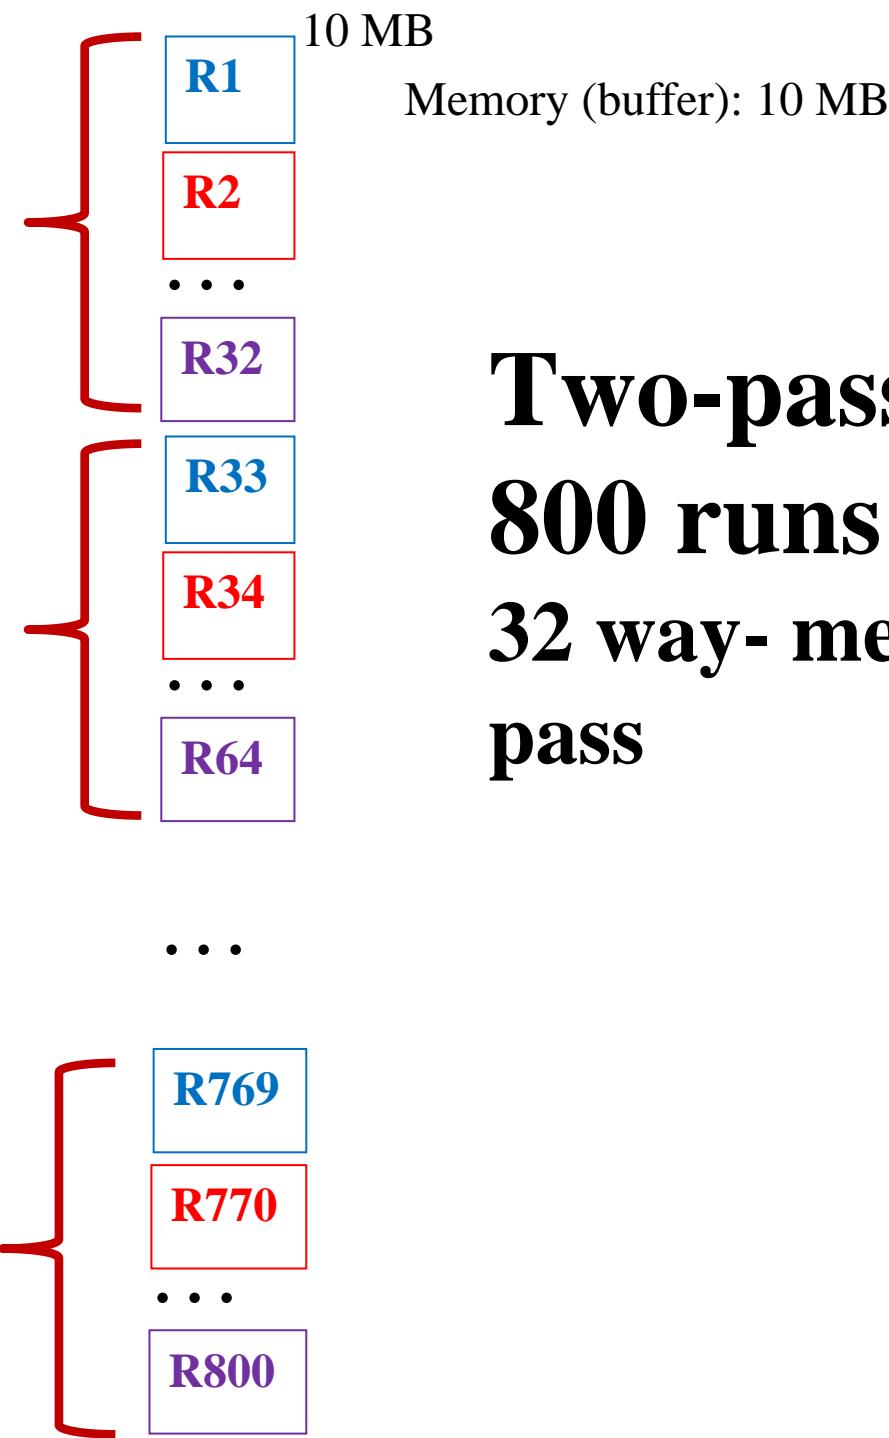
Multiple-pass multiway merges

- Instead of merging all runs at once, we break the original set of runs into small groups and merge the runs in these groups separately.
 - more memory space is available for each run; hence fewer seeks are required per run.
- When all of the smaller merges are completed, a second pass merges the new set of merged runs.

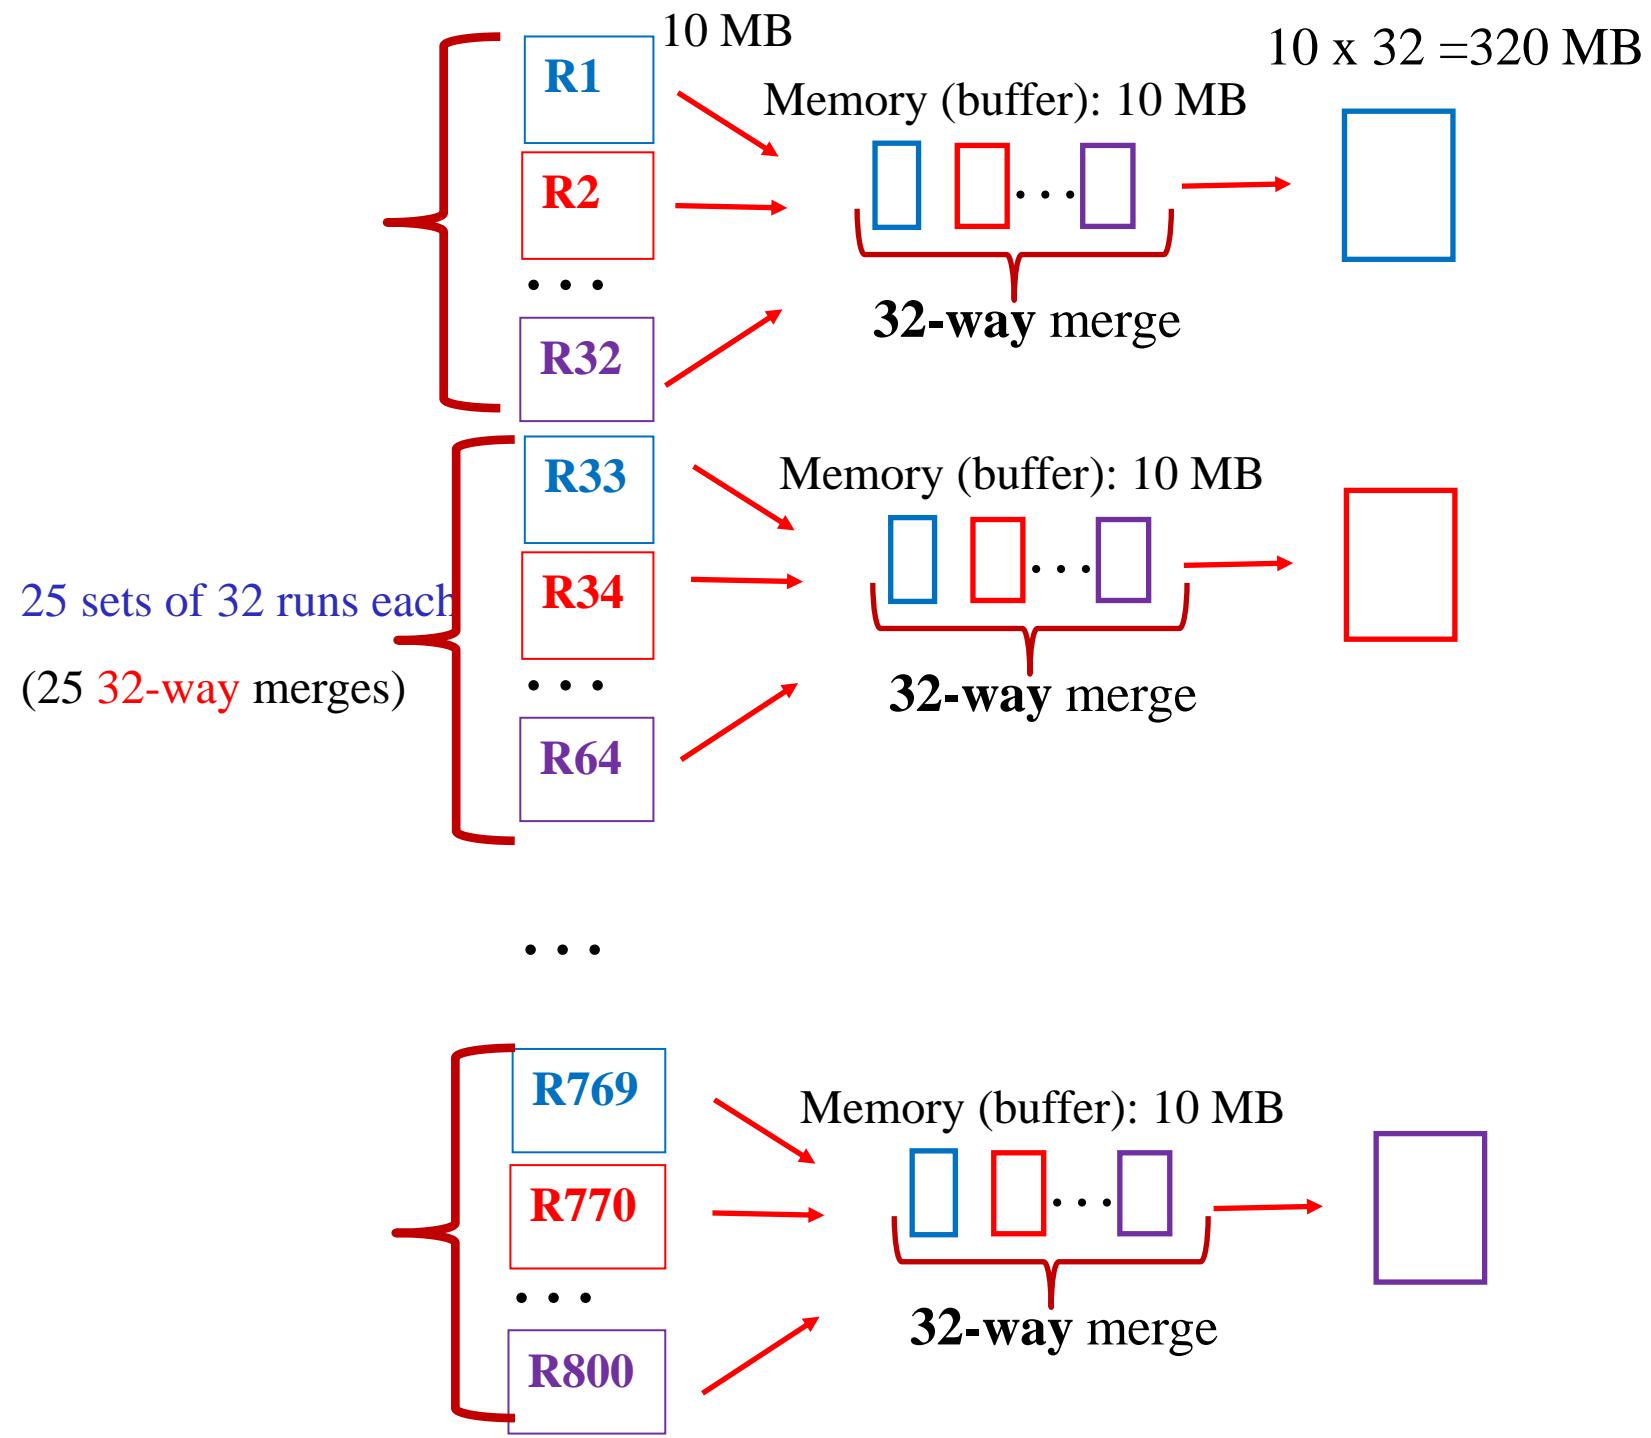


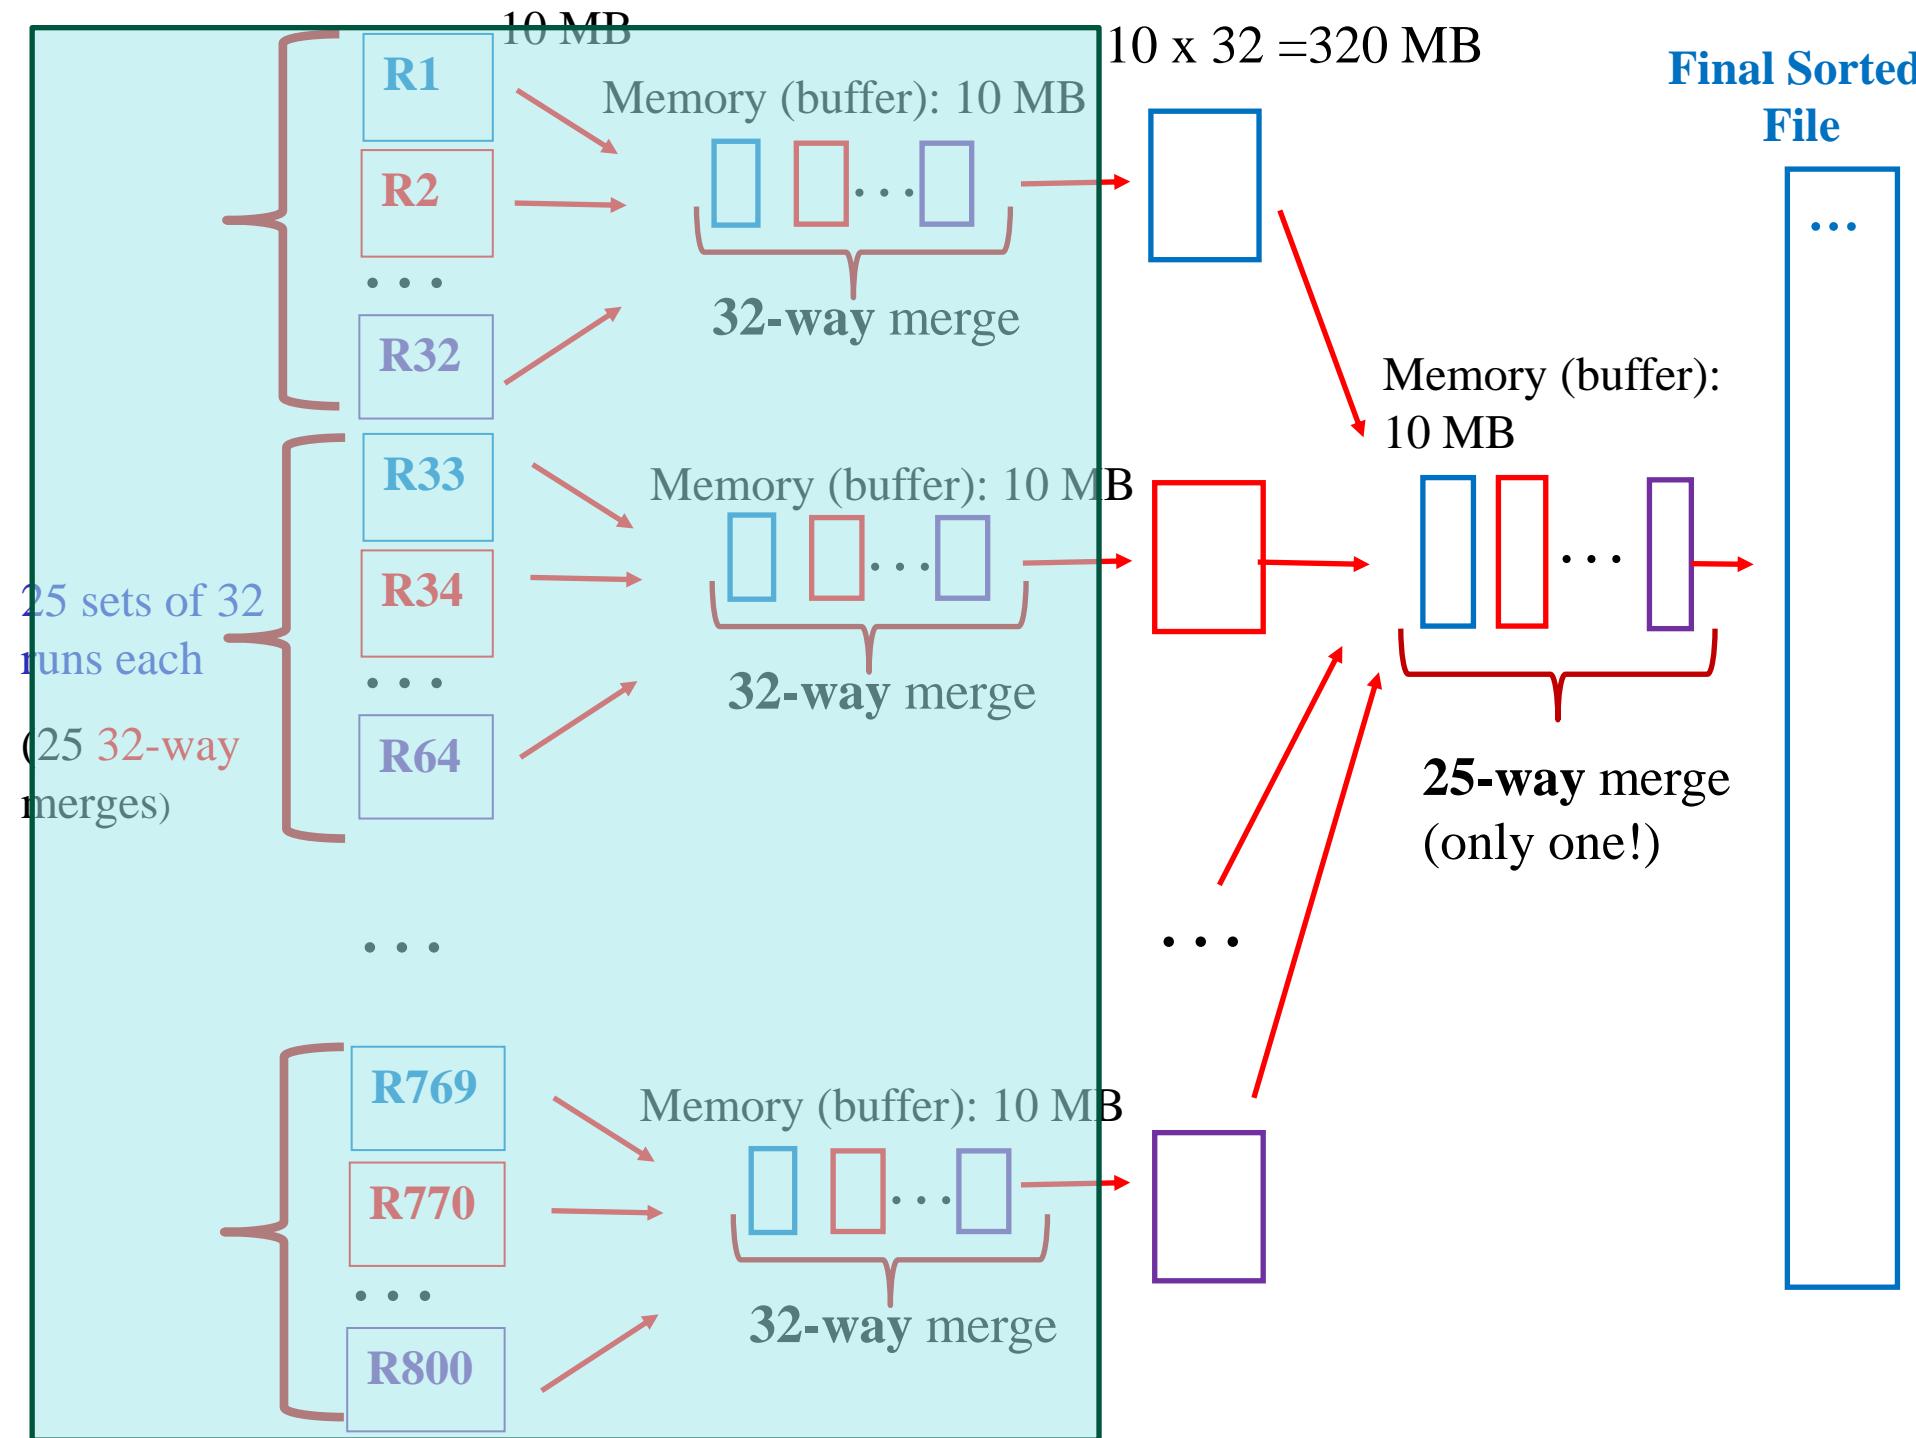
10 MB
Memory (buffer): 10 MB

Two-pass merge of 800 runs

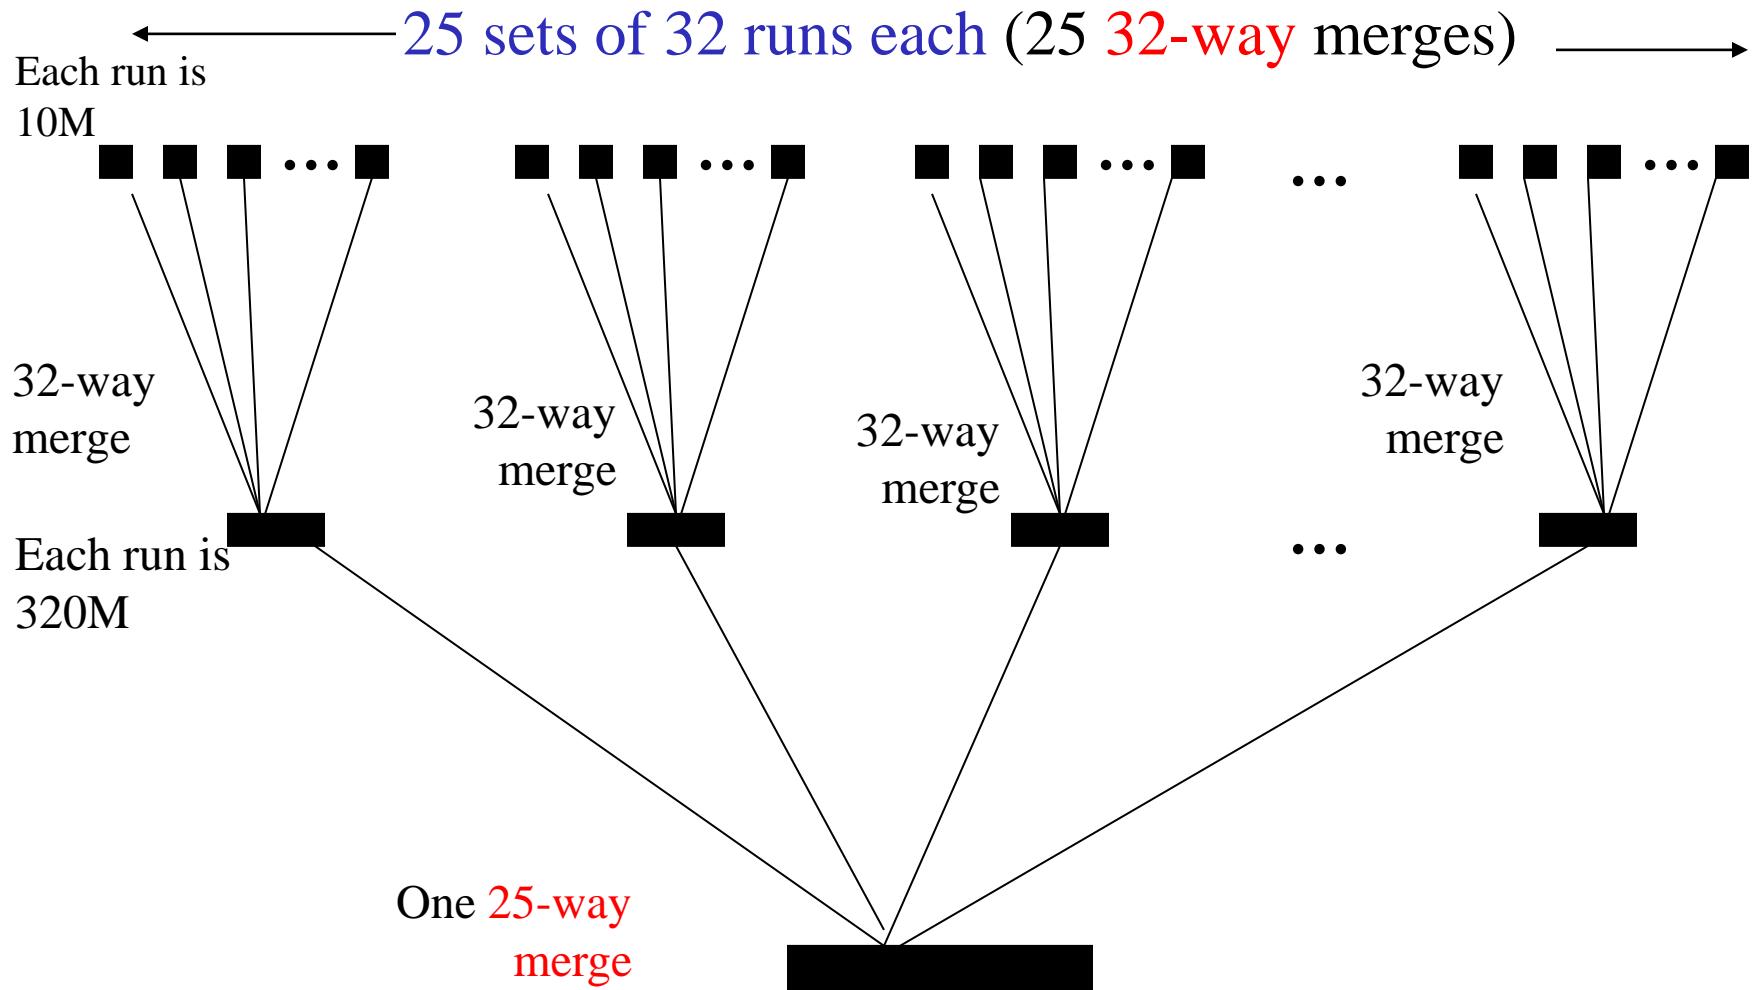


**Two-pass merge of
800 runs:
32 way- merge in the first
pass**



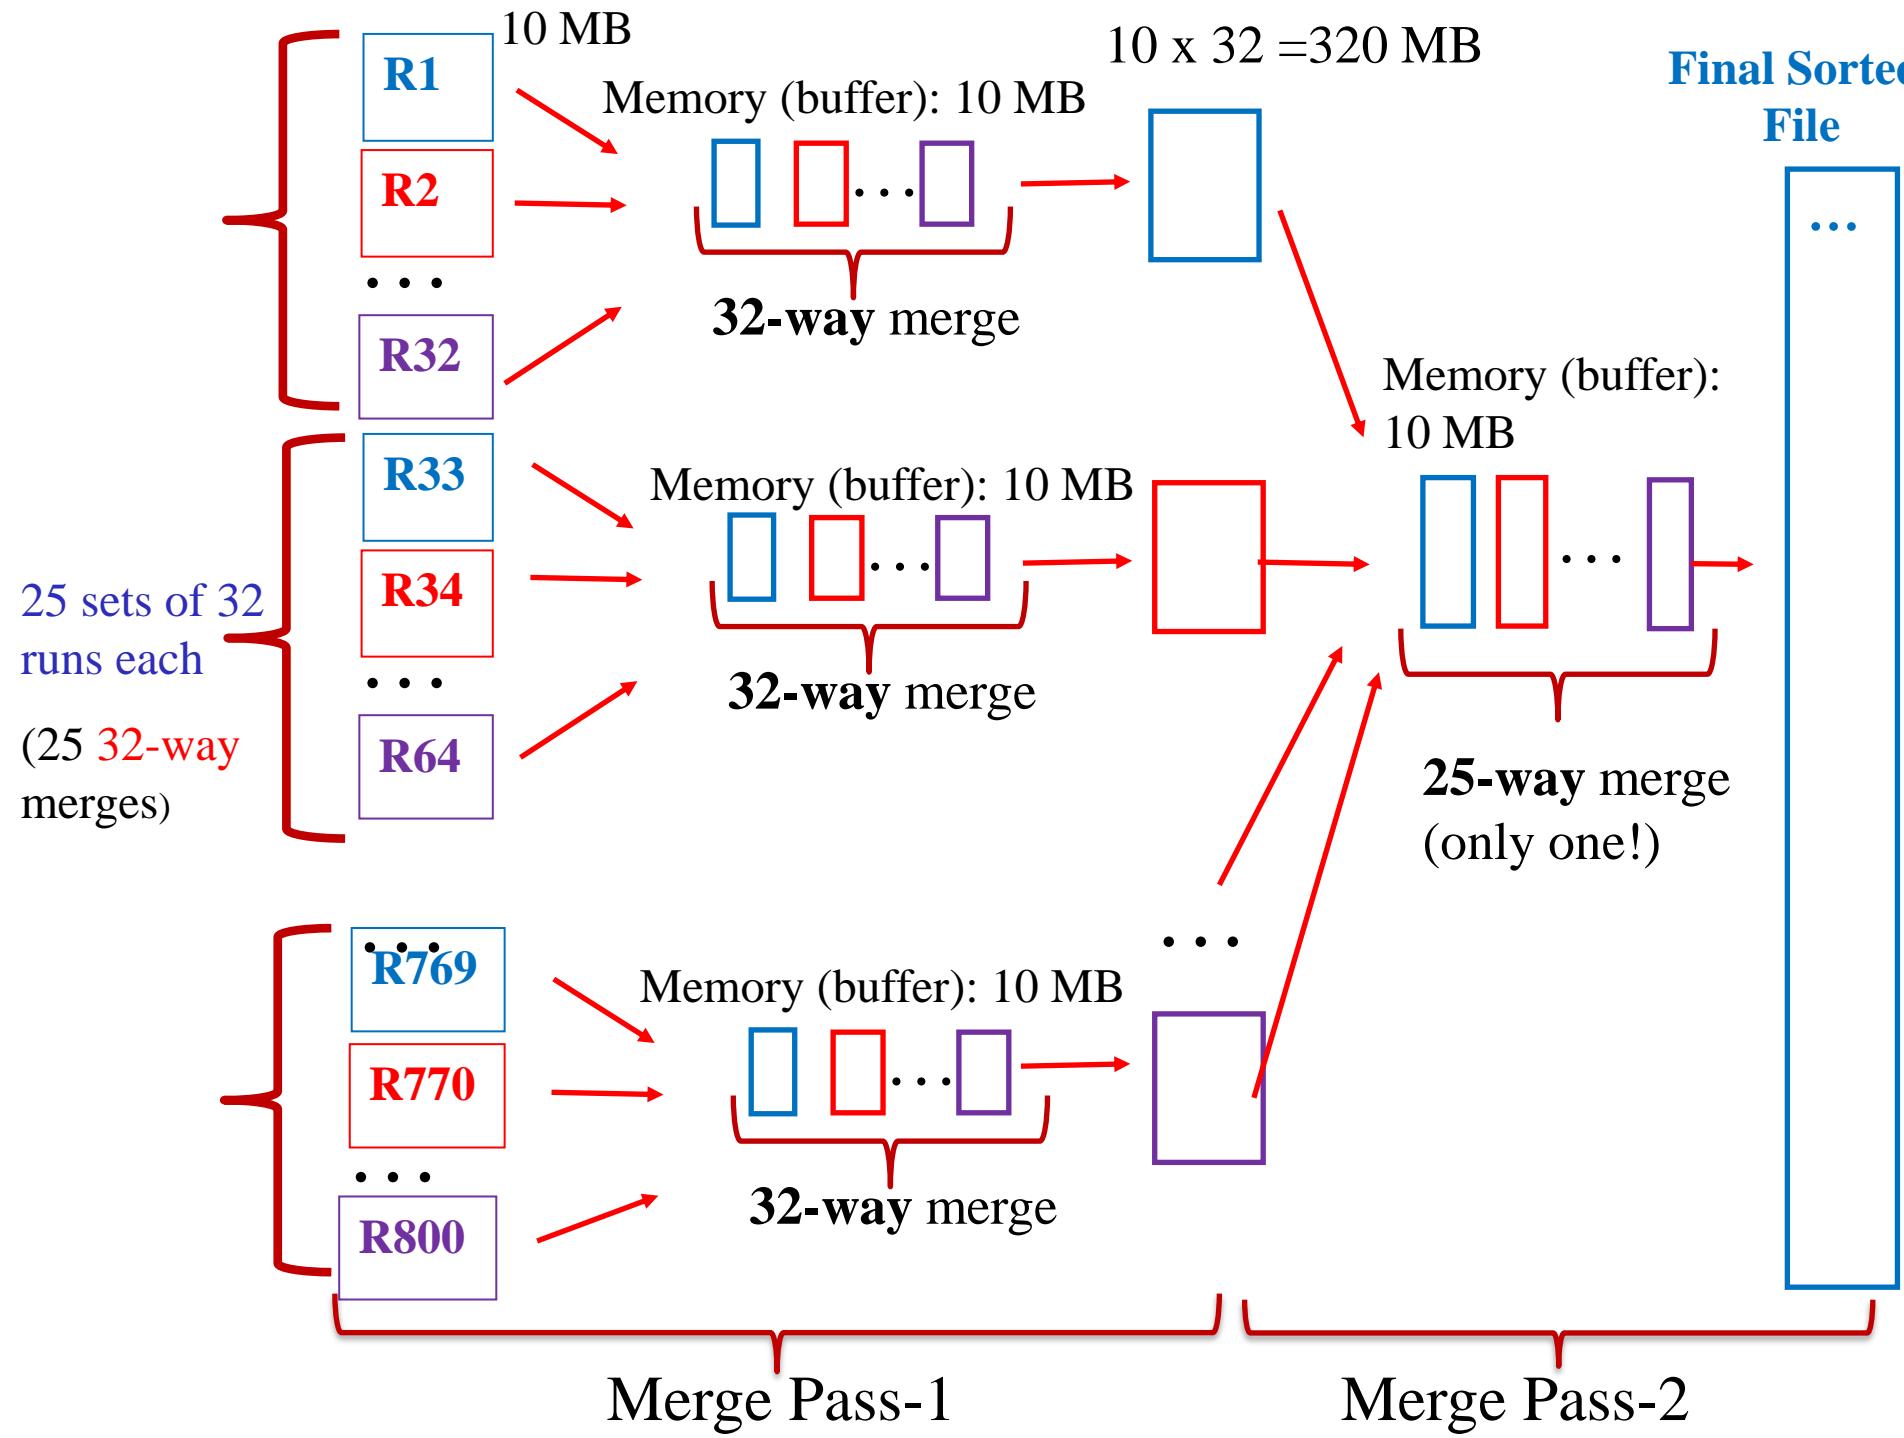


Two-pass merge of 800 runs

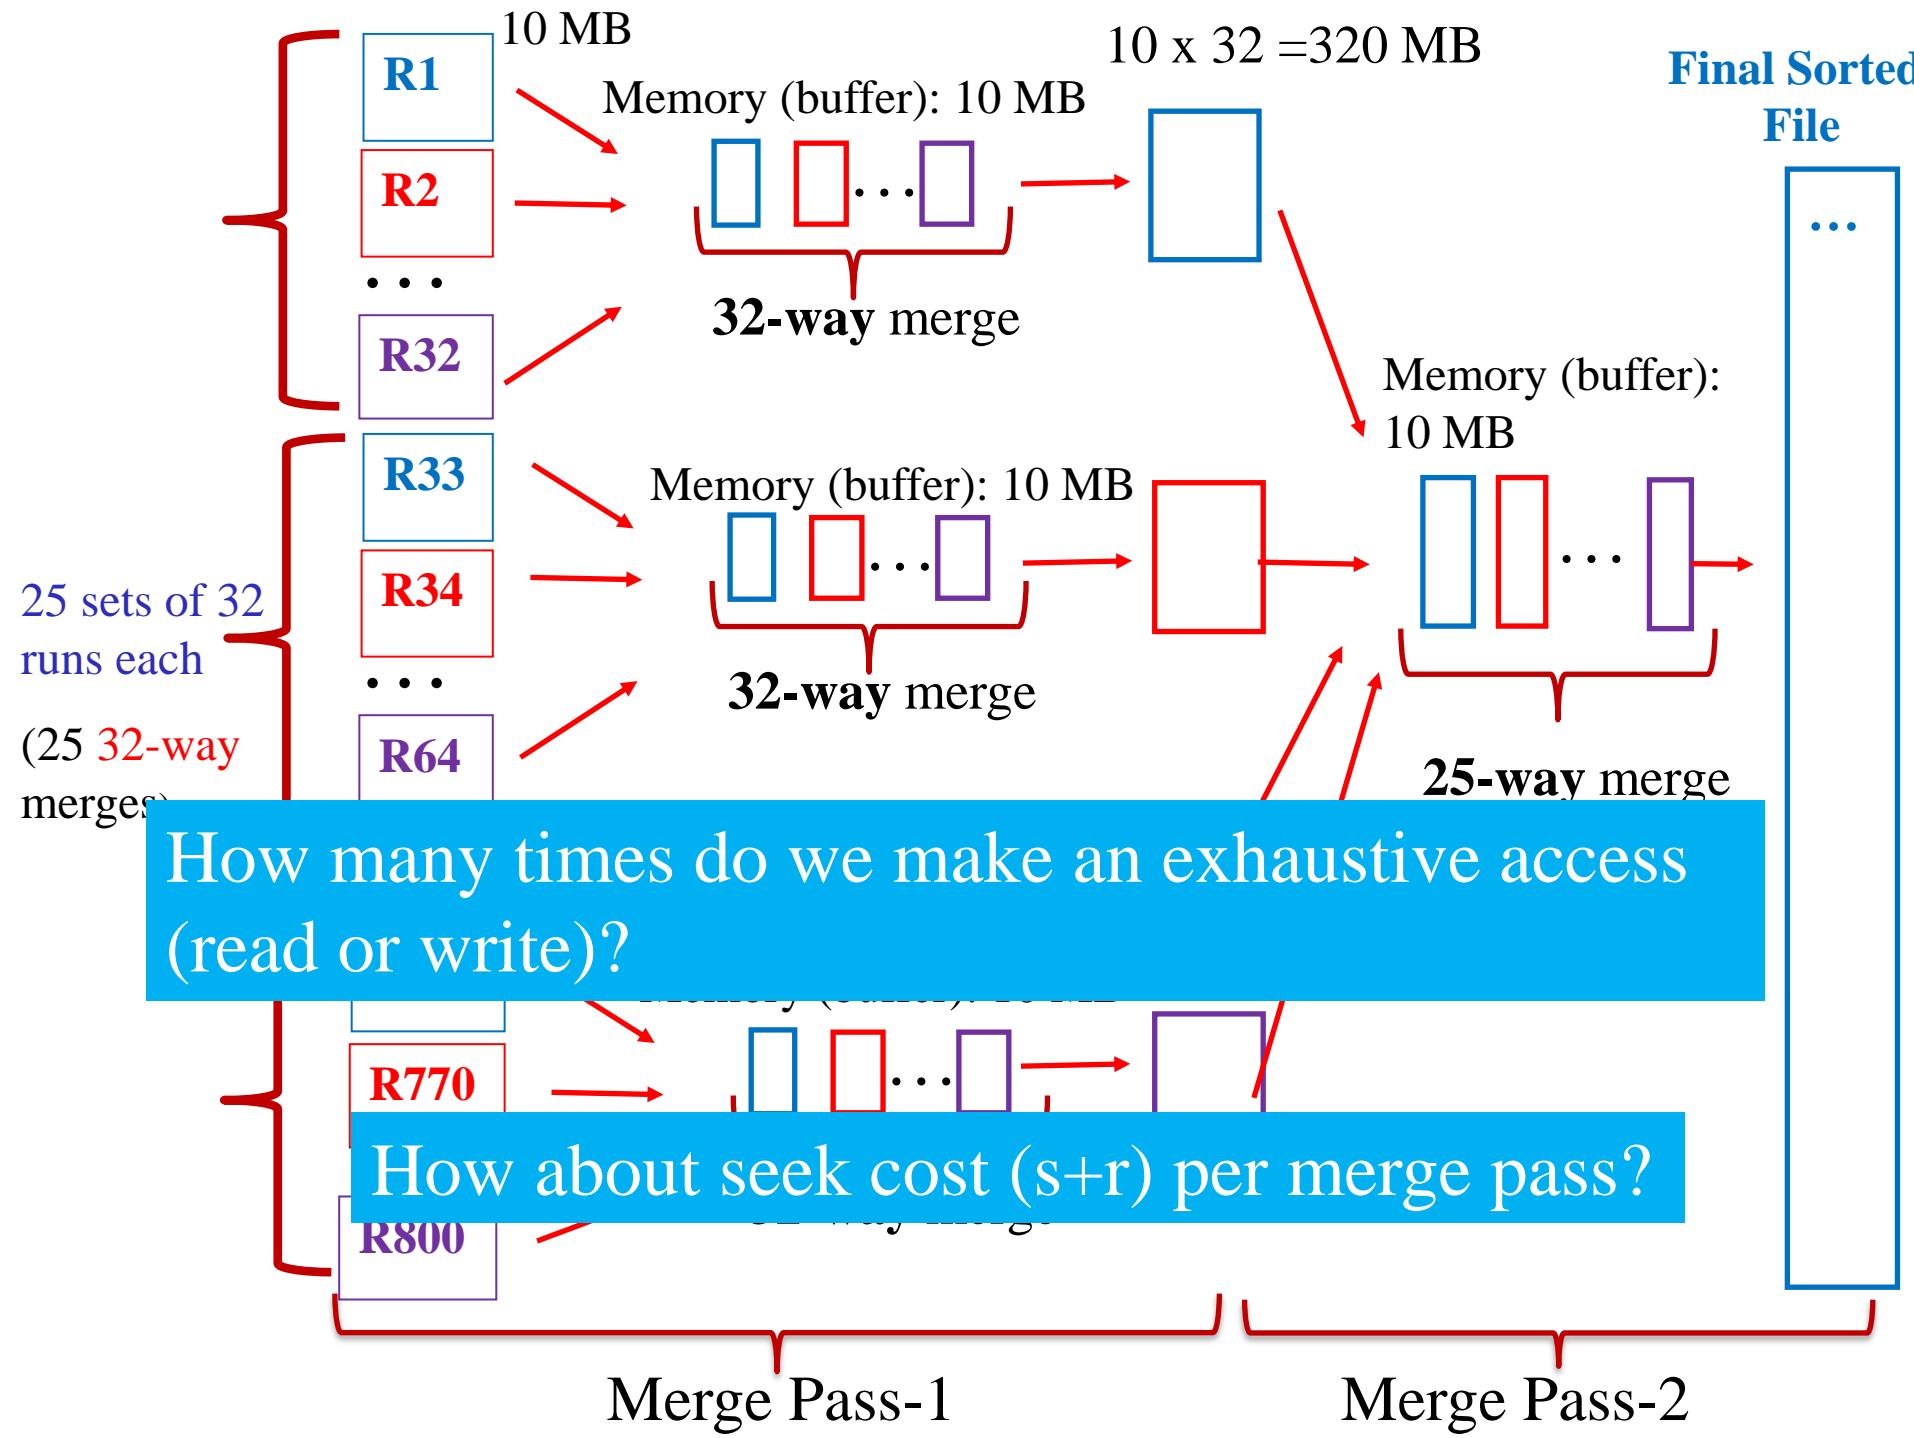


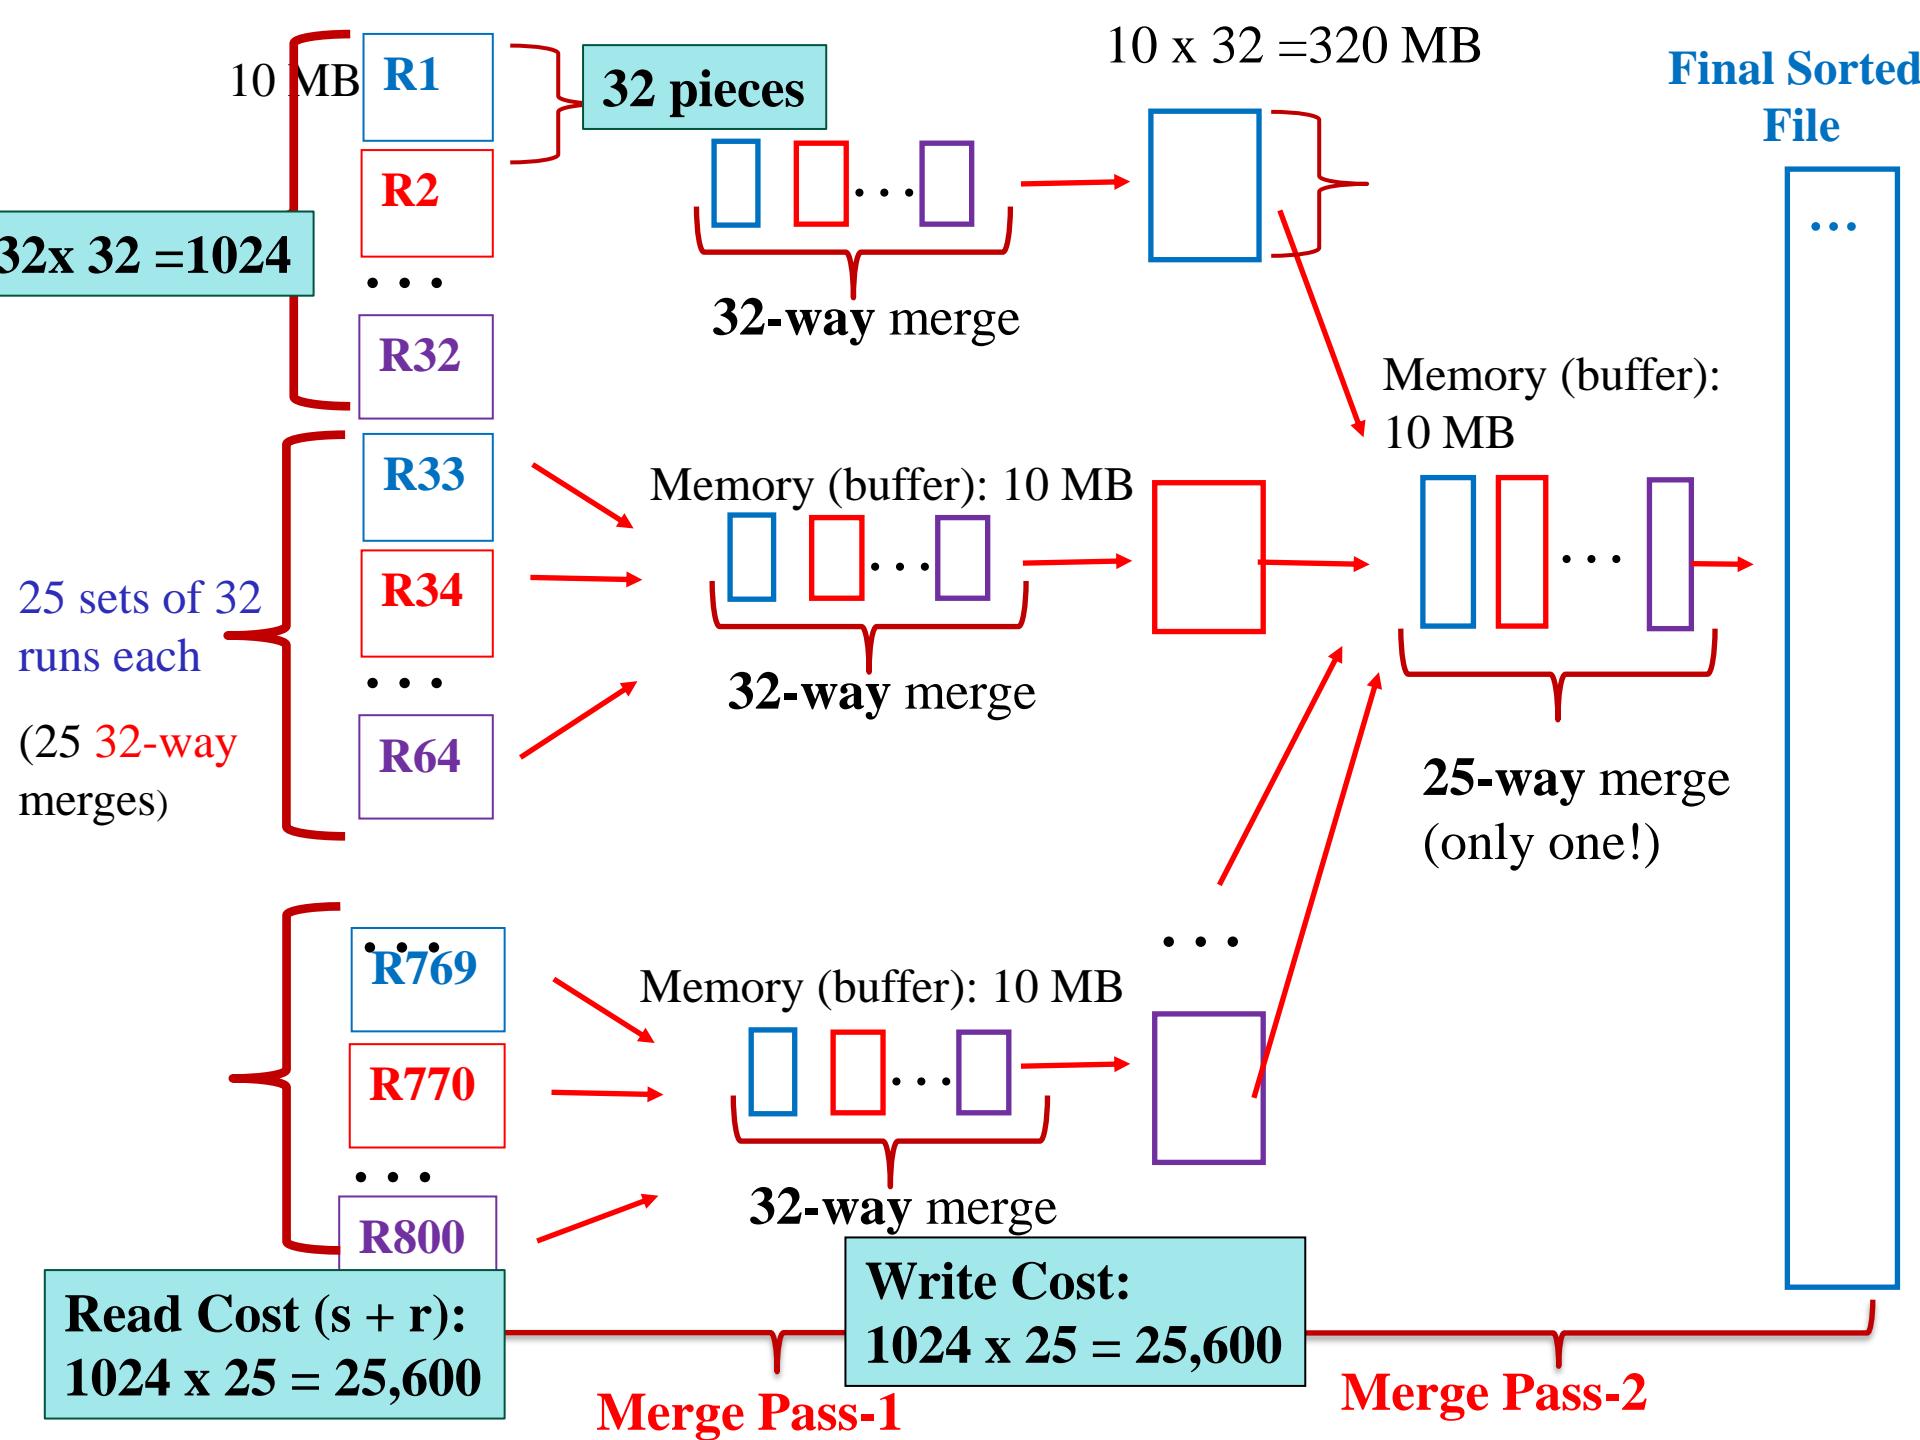
Two-pass merge of 800 runs

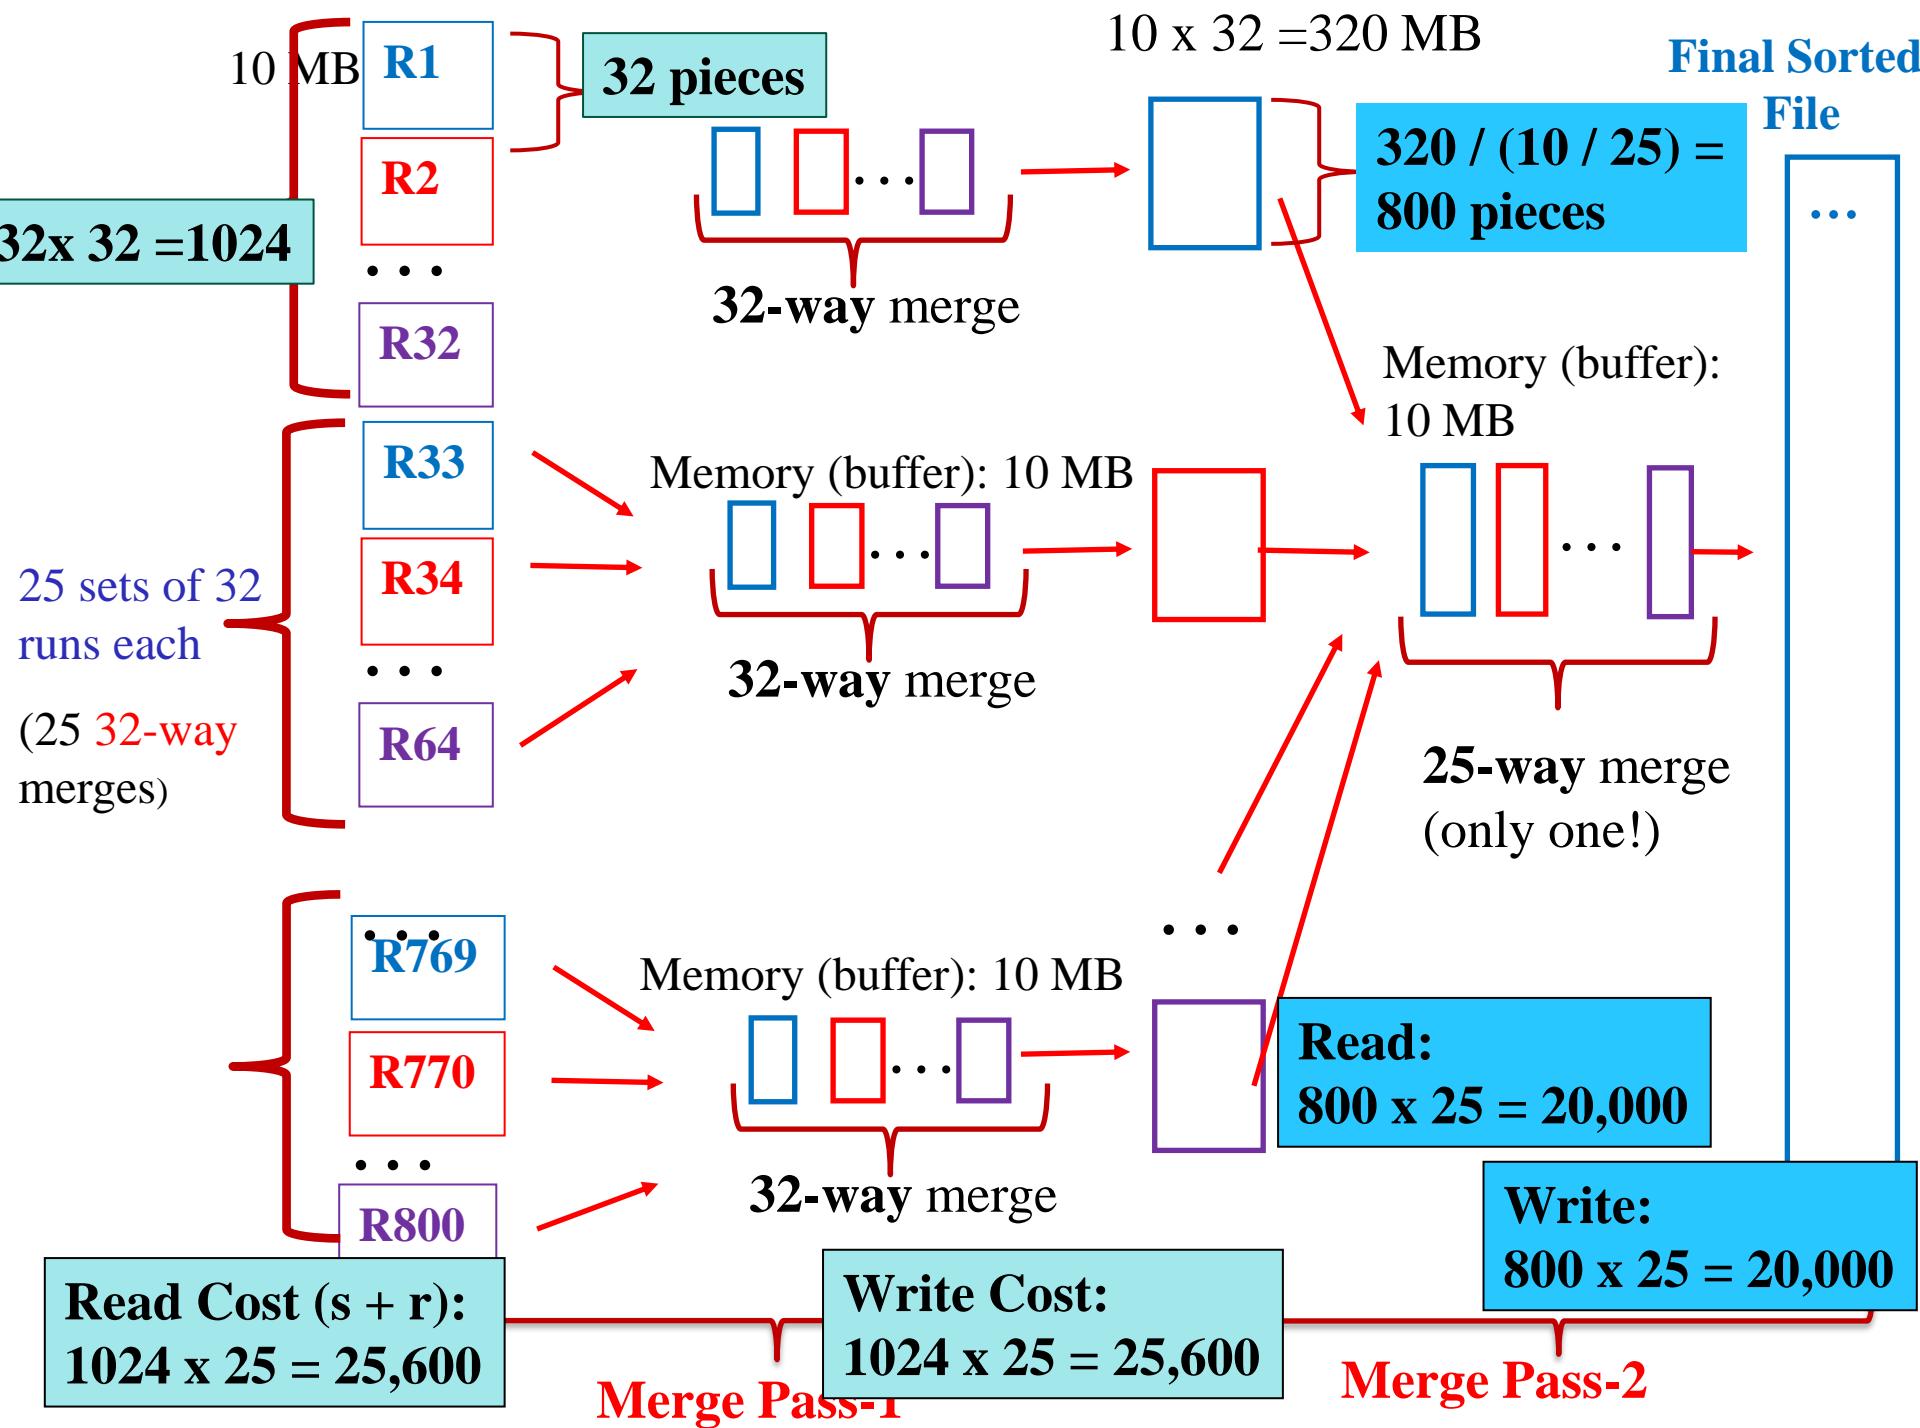
Final Sorted File



Final Sorted File







Cost of multi-pass merging

- 25 sets of 32 runs, followed by 25-way merge:
 - Disadvantage: we read every record twice.
 - Advantage: we can use larger lists and avoid a large number of disk seeks.
- Calculations:

First Merge Pass:

- List size = 1/32 run => $32 * 32 = 1024$ seeks
- For 25 32-way merges => $25 * 1024 = 25,600$ seeks

Second Merge Pass:

- For each 25 final runs, 1/25 memory space is allocated.
- So each input list is 0.4M and it can hold 4000 records (or 1/800 run)
- Hence, 800 seeks per run, so we end up making $25 * 800 = 20,000$ seeks.

Total number of seeks for reading in two passes:

$$25600 + 20000 = 45,600$$

- What about the total time for merge?
 - We now have to transmit all of the records 4 times instead of two.
 - We also write the records twice, requiring an extra 45,600 seeks.
 - **Total number of seeks : 91200 seeks**
- Still the trade is profitable.

External sorting of 800 run (8000 MB) file

With one-pass merge:

One-Pass Merge		
Create runs (Pass-0)	$2 \times 800 (s+r)$	+ 2 Tx
Merge Pass-1	2 x 640,000 (s+r)	+ 2 Tx

With two-pass merge:

Two-Pass Merge		
Create runs (Pass-0)	$2 \times 800 (s+r)$	+ 2 Tx
Merge Pass-1	2 x 25,600 (s+r)	+ 2 Tx
Merge Pass-2	2 x 20,000 (s+r)	+ 2 Tx

Total Cost

$$\begin{aligned}\text{Total # of block transfers} &= 2b * \text{btt} + (\lceil \log_k m \rceil) * 2b * \text{btt} \\ &= 2b * \text{btt} * (\lceil \log_k m \rceil + 1)\end{aligned}$$

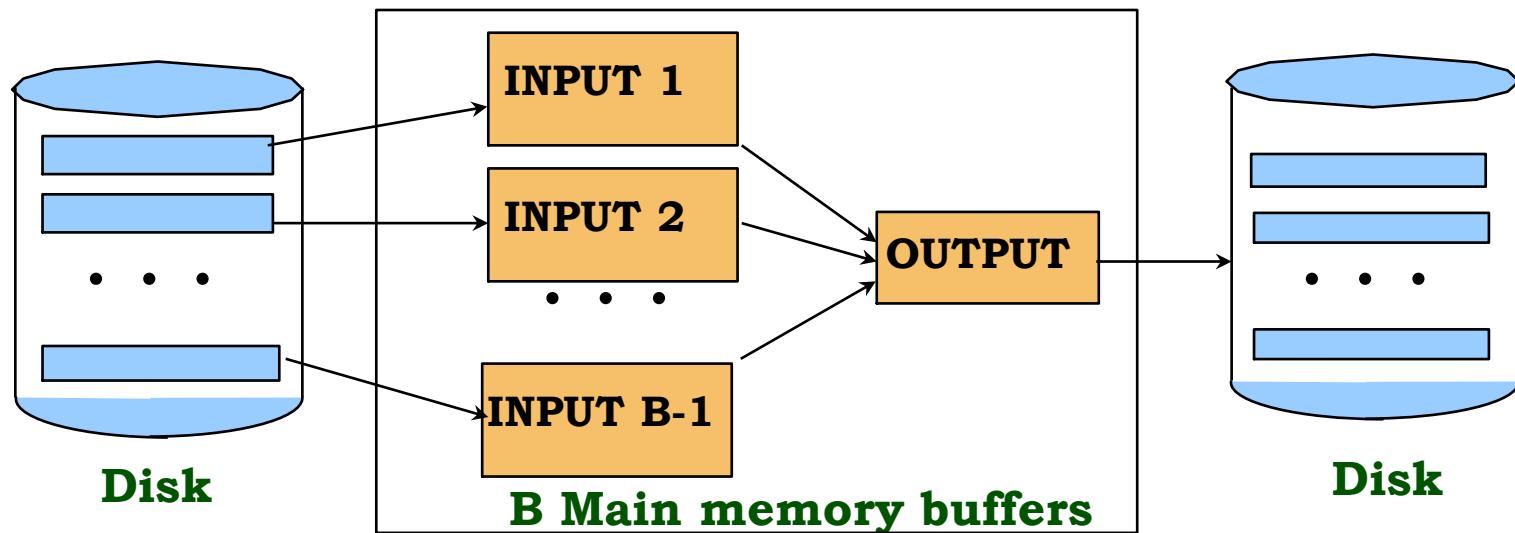
(when **k-way multi-pass merge** is used for **m runs**.)

Total cost including # of seeks:

$$= \underbrace{2b * \text{btt}}_{\text{creating sorted runs}} + \underbrace{(\lceil \log_k m \rceil) * [2k * m * (s+r) + 2b * \text{btt}]}_{\substack{\text{k-way multi-pass merge} \\ (\text{handled in } \lceil \log_k m \rceil \text{ passes (steps)})}}$$

DBMS view (from Raghu's book)

- Sort a file with N blocks (pages)
- Available memory B blocks (buffer pages):
 - Pass 0: use B buffer pages. Produce $\lceil N / B \rceil$ sorted runs of B pages each.
 - Pass 1, 2, ..., etc.: merge $B-1$ runs.



Cost of External Merge Sort

- Number of passes: $1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$
- Cost = $2N * (\# \text{ of passes})$
- E.g., with 5 buffer pages, to sort 108 page file:
 - Pass 0: $\lceil 108 / 5 \rceil = 22$ sorted runs of 5 pages each (last run is only 3 pages)
 - Pass 1: $\lceil 22 / 4 \rceil = 6$ sorted runs of 20 pages each (last run is only 8 pages)
 - Pass 2: $\lceil 6 / 4 \rceil = 2$ sorted runs, 80 pages and 28 pages
 - Pass 3: Sorted file of 108 pages

Number of Passes of External Sort

N	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

Problem

Assume you are given a file of 800,000 records, 400 bytes each, a block size of 4000 bytes with btt of 1 msec, and the available memory size is 8MB. You are asked to sort this file using 4-way merge with one disk drive. Assume $s = 5$ ms, $r = 4$ ms, $btt = 1$ ms.

1. How many runs (sorted segments) are created ?
2. What is the size of each run?
3. Estimate the cost of creating these runs.

Solution

1. How many runs (sorted segments) are created?

Number of runs = *file size/ size of avail.memory* =
 $800000 * 400 / 8000000 = 40$ sorted segments

2. What is the size of each run?

Size of each run = size of available memory = 8MB

3. Estimate the cost of creating these runs

- Number of blocks = $b = 800000 * 400 / 4000 = 80000$ blocks
- Cost of creating the runs = $2 * b * btt = 2 * 80000 * (1ms) = 160000 ms$

Problem (cont.)

4. What is the number of passes necessary to merge the sorted segments using 4-way merge?
5. Estimate the cost of one pass of merging.
6. Estimate the total cost of sorting the file, including the creation of sorted runs and the multi- step merging.

Solution (cont.)

4. What is the number of passes necessary to merge the sorted segments using 4-way merge?

$$\text{Number of passes} = \lceil \log_4 40 \rceil = 3$$

5. Estimate the cost of one pass of merging.

Cost of one pass of merging =

$$= 2 * k * (\# \text{ of runs}) * (s+r) + 2 * b * \text{btt}$$

where k is the number in k -way merging (= no of pieces in each run)

$$= 2 * 4 * 40 * (9 \text{ ms}) + 2 * 80000 * (1\text{ms})$$

$$= 2880 \text{ ms} + 160000 \text{ ms}$$

$$= 162880 \text{ ms}$$

Solution (cont.)

6. Estimate the total cost of sorting the file, including the creation of sorted runs and the multi- step merging.

Total cost of sorting the file =

$$\# \text{ of passes} * \text{cost for one pass of merging} + \text{cost of creating runs}$$

(Pass-0)

$$= 3 * (162880 \text{ ms}) + 160000 \text{ ms} = 648640 \text{ ms}$$