

# **Indexing**

# What is an Index?

- A simple index is a table containing an ordered list of keys and reference fields.
  - e.g. the index of a book
- In general, indexing is another way to handle the searching problem.

# Uses of an index

1. An index lets us **impose order on a file** without rearranging the file.
2. Indexes provide **multiple access paths** to a file.
  - e.g. library catalog providing search for author, book and title
3. An index can provide **keyed access to variable-length record files**.

# A simple index for a pile file

	<b>Label</b>	<b>ID</b>	<b>Title</b>	<b>Composer</b>	<b>Artist</b>
17	LON	2312	Symphony No.9	Beethoven	Giulini
62	RCA	2626	Romeo and Juliet	Prokofiev	Maazel
117	WAR	23699	Nebraska  ...		
152	ANG	3795	Violin Concerto  ...		

Address of record  
(i.e. Byte offset)

Primary key = (**Label**, **ID**)

- Index is sorted (in main memory).
- Records appear in file in the order they are entered.

# Index array:

Key	Reference
ANG3795	152
LON2312	17
RCA2626	62
WAR23699	117

- How to search for a recording with given LABEL ID?
  - **Binary search** in the index and then seek for the record in position given by the reference field.

# Operations to maintain an indexed file

- **Create** the original empty index and data files.
- **Load** the index file into memory before using it.
- **Rewrite** the index file **from memory** after using it.
- **Add** data **records** to the data file.
- **Delete records** from the data file
- **Update records** in the data file.
- **Update the index** to reflect changes in the data file

# Rewrite the index file from memory

- When the data file is closed, the index in memory needs to be written to the index file.
- An important issue to consider is what happens if the rewriting does not take place (e.g. power failures, turning machine off, etc.)
- Two important safeguards:
  - Keep a status flag in the header of the index file.
  - If the program detects the index is out of date it calls a procedure that reconstructs the index from the data file.

# Record Addition

1. Append the new record to the end of the data file.
  2. Insert a new entry to the index in the right position.
    - needs rearrangement of the index
- Note: this rearrangement is done in the main memory.

# Record Deletion

- Use the techniques for reclaiming space in files when deleting records from the data file.
- We must also delete the corresponding entry from the index in memory.

# Record Updating

There are two cases to consider:

1. The update changes the value of the key field:
  - Treat this as a deletion followed by an insertion
2. The update does not affect the key field:
  - If record size is unchanged, just modify the data record. If record size is changed treat this as a delete/insert sequence.

# **Indexing by Multiple Keys**

- We could build additional indexes for a file to provide multiple views of a data file.
  - e.g. Find all recordings of Beethoven’s work.
- LABEL ID is a **primary key**.
- There may be **other search keys**: title, composer, artist.
- We can build **secondary indexes**.

# Composer index:

Composer	Primary key
Beethoven	ANG3795
Beethoven	DG139201
Beethoven	DG18807
Beethoven	RCA2626
Corea	WAR23699
Dvorak	COL31809
Prokofiev	LON2312

- Note that reference is to the primary key rather than to the byte offset.

# Retrieval using combinations of secondary keys

- Secondary indexes are useful in allowing the following kinds of queries:
  - Find all recordings of Beethoven’s work.
  - Find all recordings titled “Violin concerto”
  - Find all recordings with composer Beethoven and title Symphony No.9.
- Boolean operators “and”, “or” can be used to combine secondary search keys to qualify a request.

# Example

- The last query is executed as follows:

<b>Matches from composer index</b>	<b>Matches from title index</b>	<b>Matched list (logical “and”)</b>
ANG3795	ANG3795	ANG3795
DG139201	COL31809	DG18807
DG18807	DG18807	
RCA2626		

# Problems with simple indexes

If index does not fit in memory:

1. Seeking the index is slow (binary search):
  - We don't want more than 3 or 4 seeks for a search.

N	$\text{Log}(N+1)$
15 keys	4
1000	~10
100,000	~17
1,000,000	~20

2. Insertions and deletions take  $O(N)$  disk accesses.

# Indexes too large to fit into Memory

- Two main alternatives:
  1. Tree-structured (multi-level) index such as B+trees.
  2. Hashed organization (when access speed is a top priority)

# Multilevel Indexing and B+ Trees

# Outline

- Single-level index
- Multi-level index
- B+tree index

All can be classified as:

- Dense vs. sparse index
- Primary vs. secondary index
- Clustered vs. unclustered index

# **Indexed Sequential Access**

Provide a choice between two alternative views of a file:

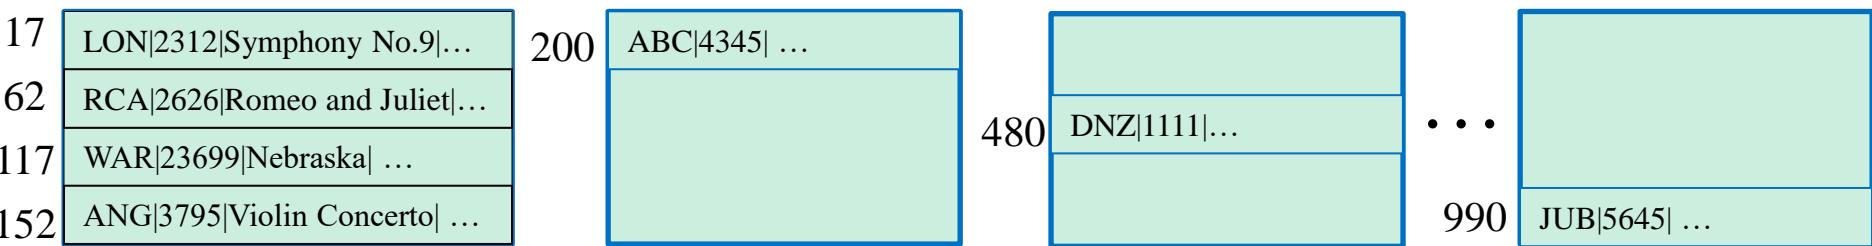
- 1. Indexed:** the file can be seen as a set of records that is indexed by key; or
- 2. Sequential:** the file can be accessed sequentially (physically contiguous records), returning records in order by key.

# Example of applications

- Student record system in a university:
  - Indexed view: access to individual records
  - Sequential view: batch processing when posting grades
- Credit card system:
  - Indexed view: interactive check of accounts
  - Sequential view: batch processing of payments

# Indexing: Basics

- A **pile file** on disk:



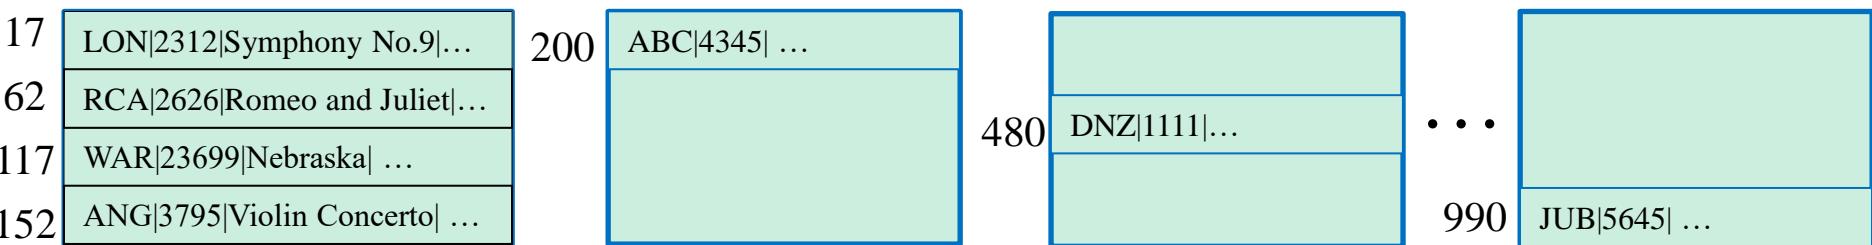
Key	Reference
ABC4345	200
ANG3795	152
DNZ1111	480
JUB5645	990
LON2312	17
RCA2626	62
WAR23699	117

Can it fit in memory?

Most probably Not!

# Indexing: Basics

- A **pile file** on disk:



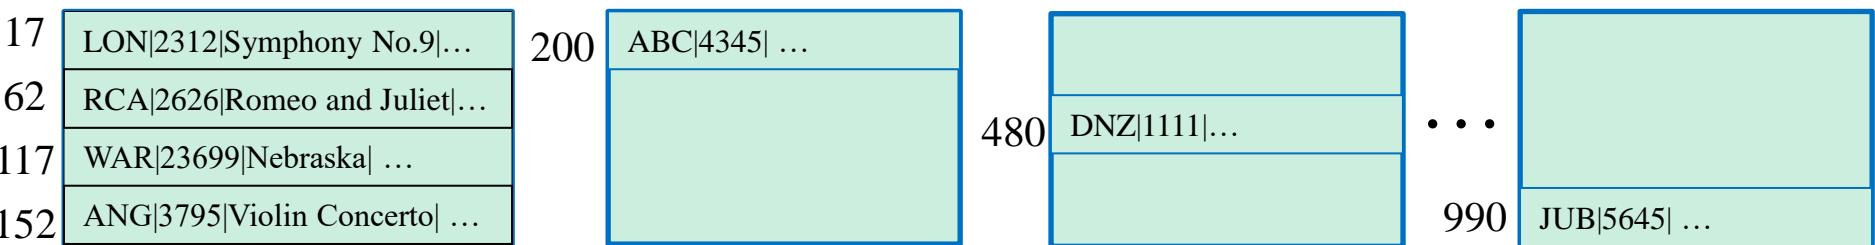
Key	Reference
ABC4345	200
ANG3795	152
DNZ1111	480
JUB5645	990
LON2312	17
RCA2626	62
WAR23699	117

Its index as a **sorted sequential file**

ABC4345	200
ANG3795	152
DNZ1111	480
JUB5645	990
LON2312	17
RCA2626	62
WAR23699	117
...	
ZZZ1111	795

# Indexing: Basics

- A **pile file** on disk:



Key	Reference
ABC4345	200
ANG3795	152
DNZ1111	480
JUB5645	990
LON2312	17
RCA2626	62
WAR23699	117

Its index as a **sorted sequential file**

ABC4345	200
ANG3795	152
DNZ1111	480
JUB5645	990

LON2312	17
RCA2626	62
WAR23699	117
...	

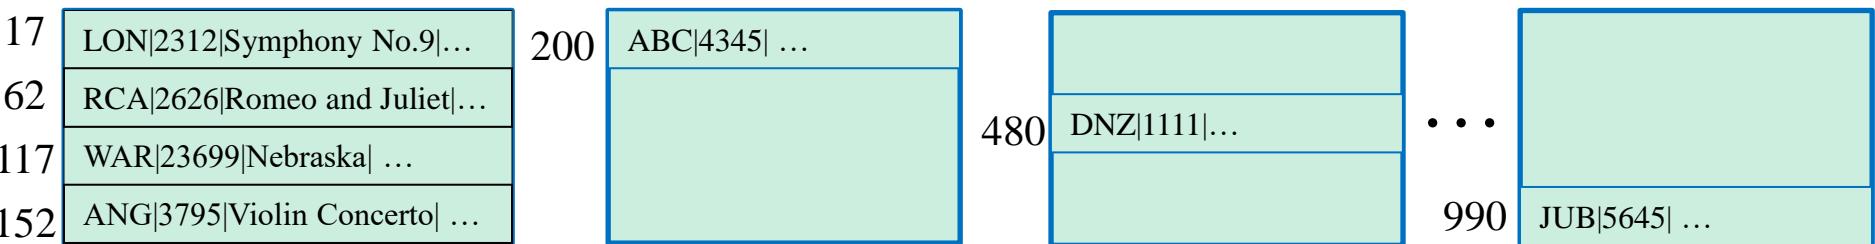
  

ZZZ1111	795

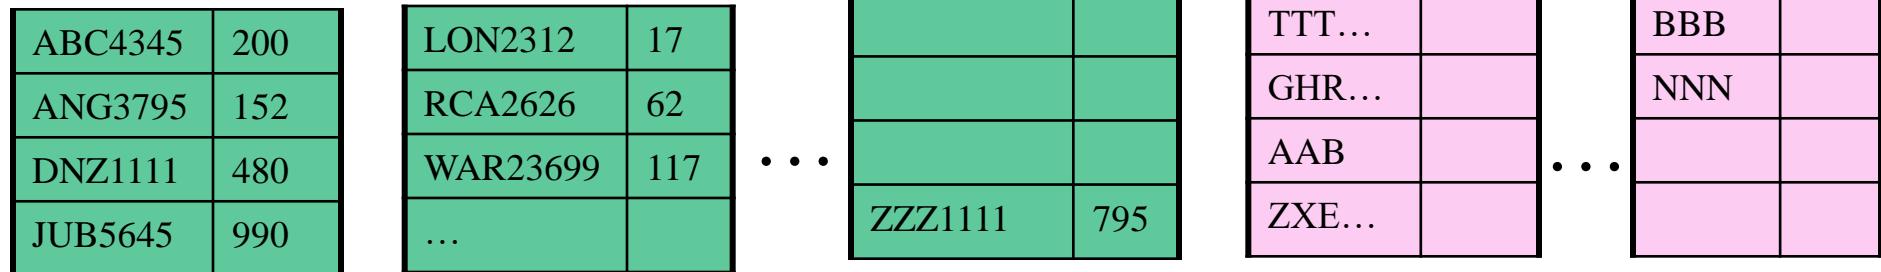
# Indexing: Basics

- A **pile file** on disk:

$$T_F = (b/2) * \text{btt}$$



- Index as a **sorted sequential file**:



Sorted area (x blocks)

$$T_F = \log_2 x * (s + r + \text{btt})$$

Overflow area (y blocks)

$$+ s + r + (y/2) * \text{btt}_{24}$$

# Indexing: Basics

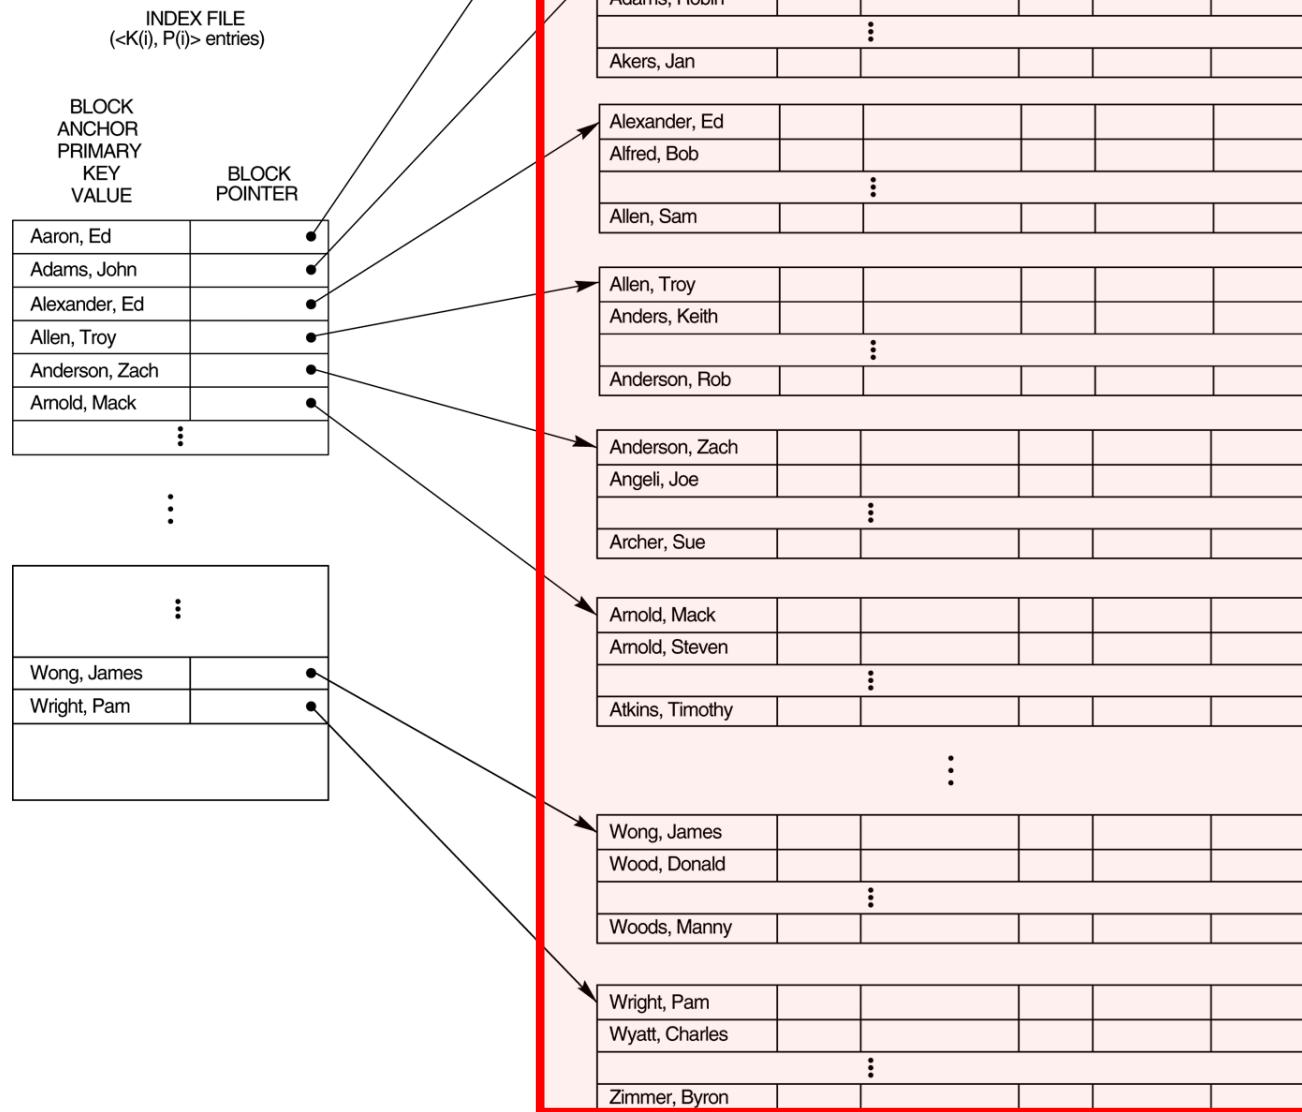
- Remember: We don't want more than 3 or 4 seeks for a search
- So any better way to structure the index?

# The initial idea: Single level index

- The data file is ordered on a ***key field***.
  - The records are grouped into blocks in a sorted way.
- A single level index for these blocks:
  - Includes one index entry *for each block* in the data file; the index entry has the key field value,  $K(i)$ , for the *first record* in the block, which is called the *block anchor*.
  - A similar scheme can use the *last record* in a block.
  - Index is an ordered file with entries (records)  $\langle K(i), P(i) \rangle$ 
    - We can still do binary search
    - Would be smaller than the data file

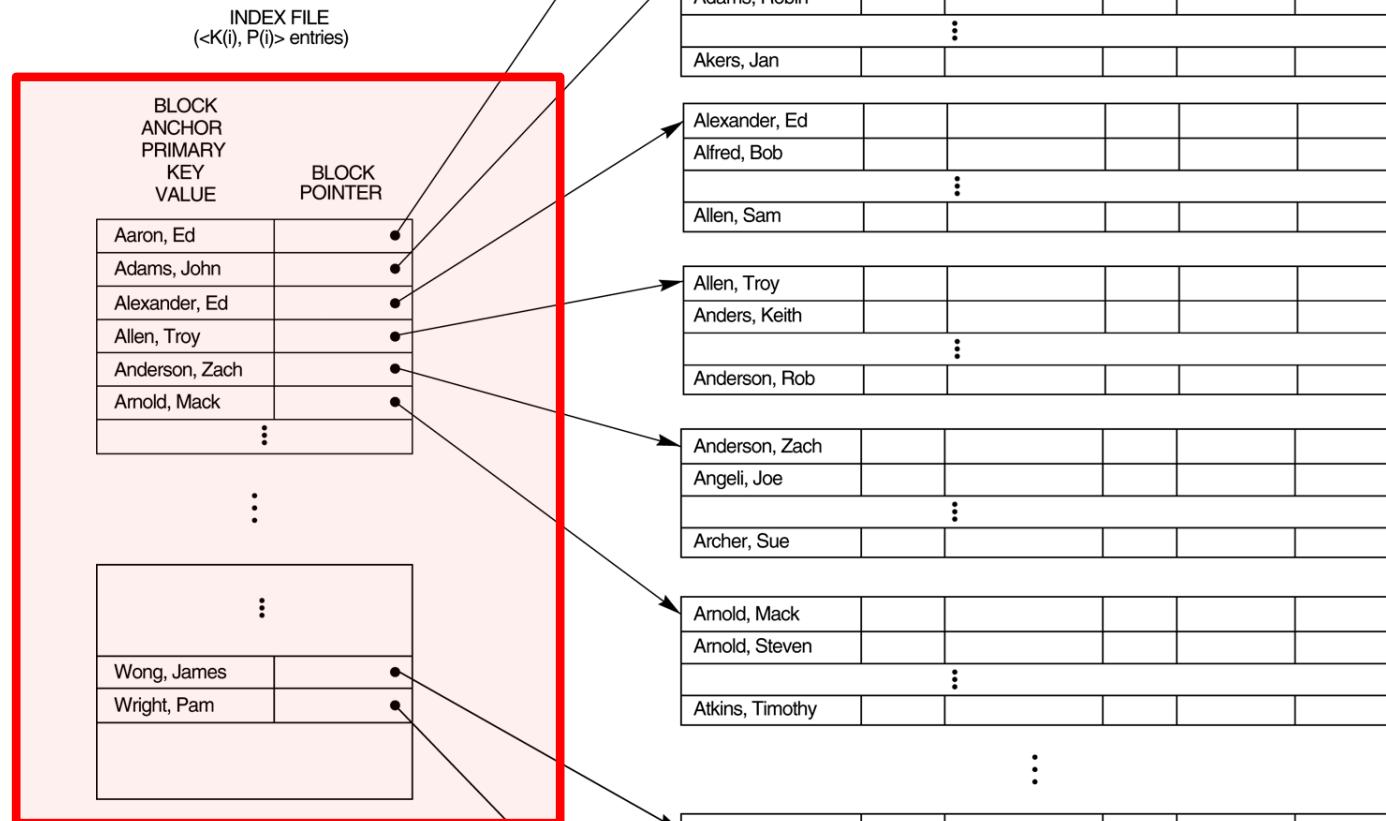
# Single-level index on the ordering key field of the file.

Data file is ordered on the (primary) key field, i.e, Name



# Single-level index on the ordering key field of the file.

Data file is ordered on the (primary) key field, i.e, Name



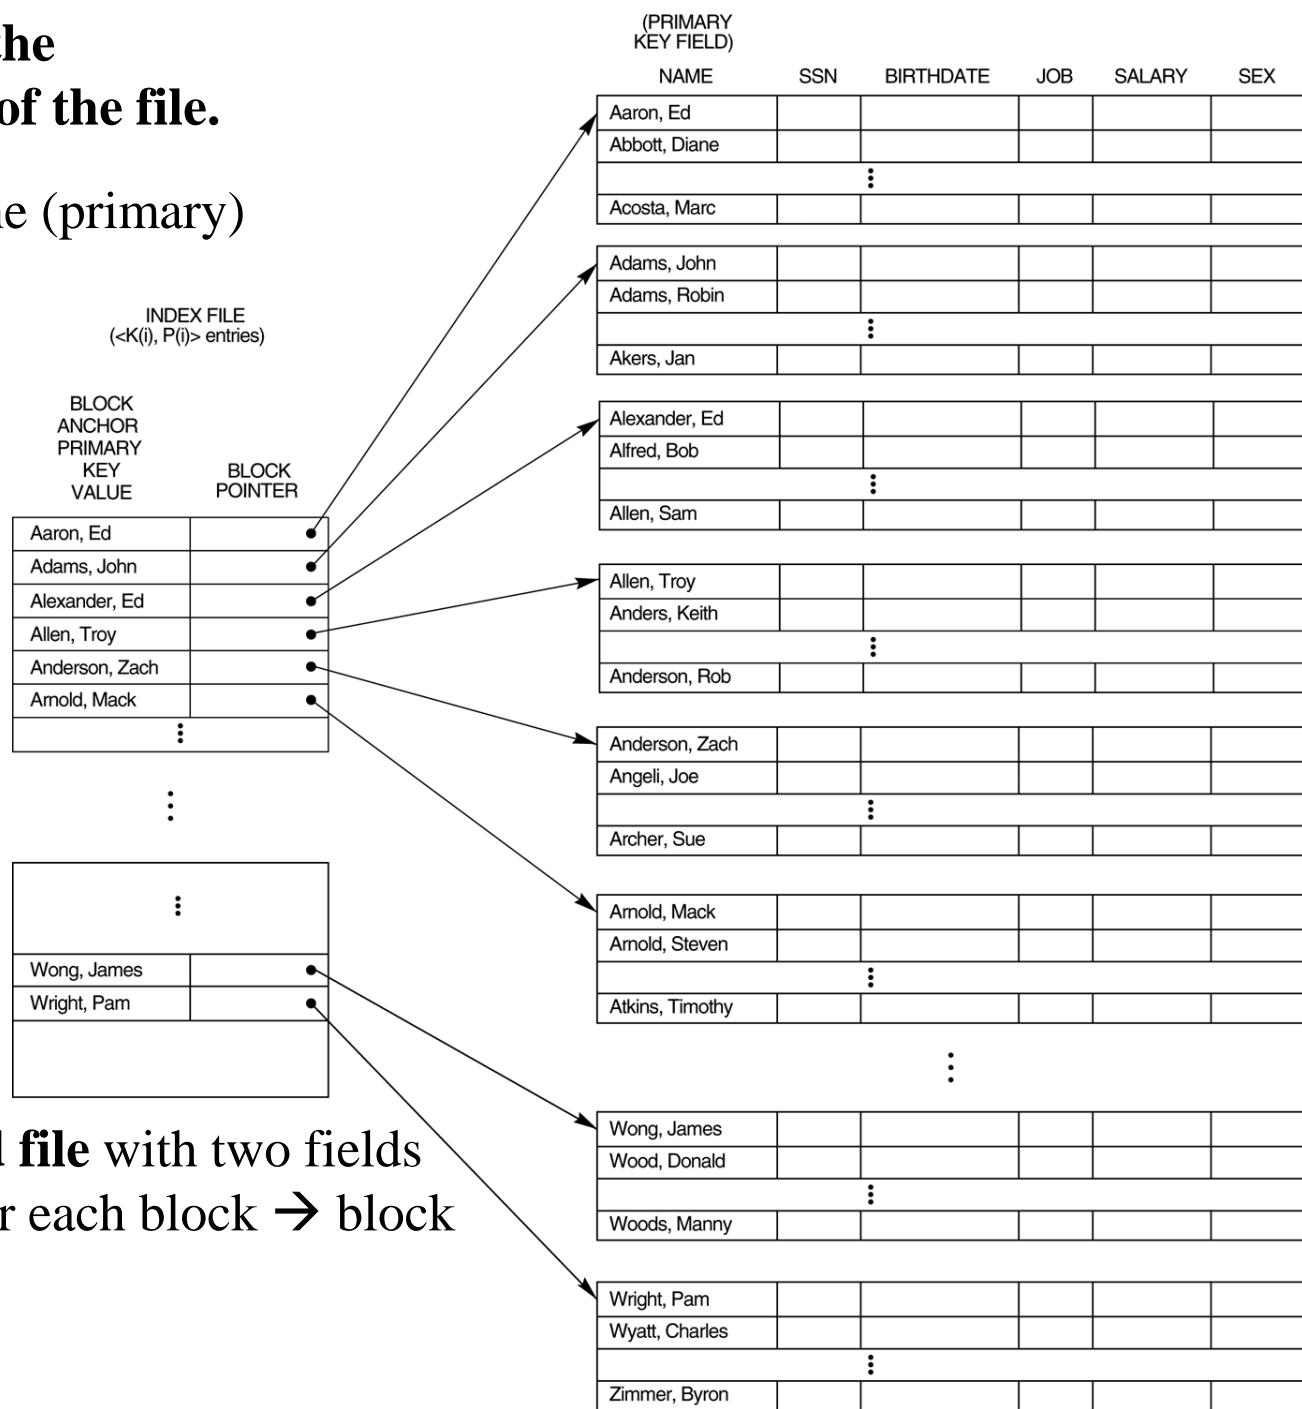
The **index** is an **ordered file** with two fields

$K(i)$ : One index entry for each block  $\rightarrow$  block anchor is PK value

$P(i)$ : Block pointer

# Single-level index on the ordering key field of the file.

Data file is ordered on the (primary) key field, i.e, Name



Select \*  
From Emp E  
Where E.Name =  
**Altman, Sam**

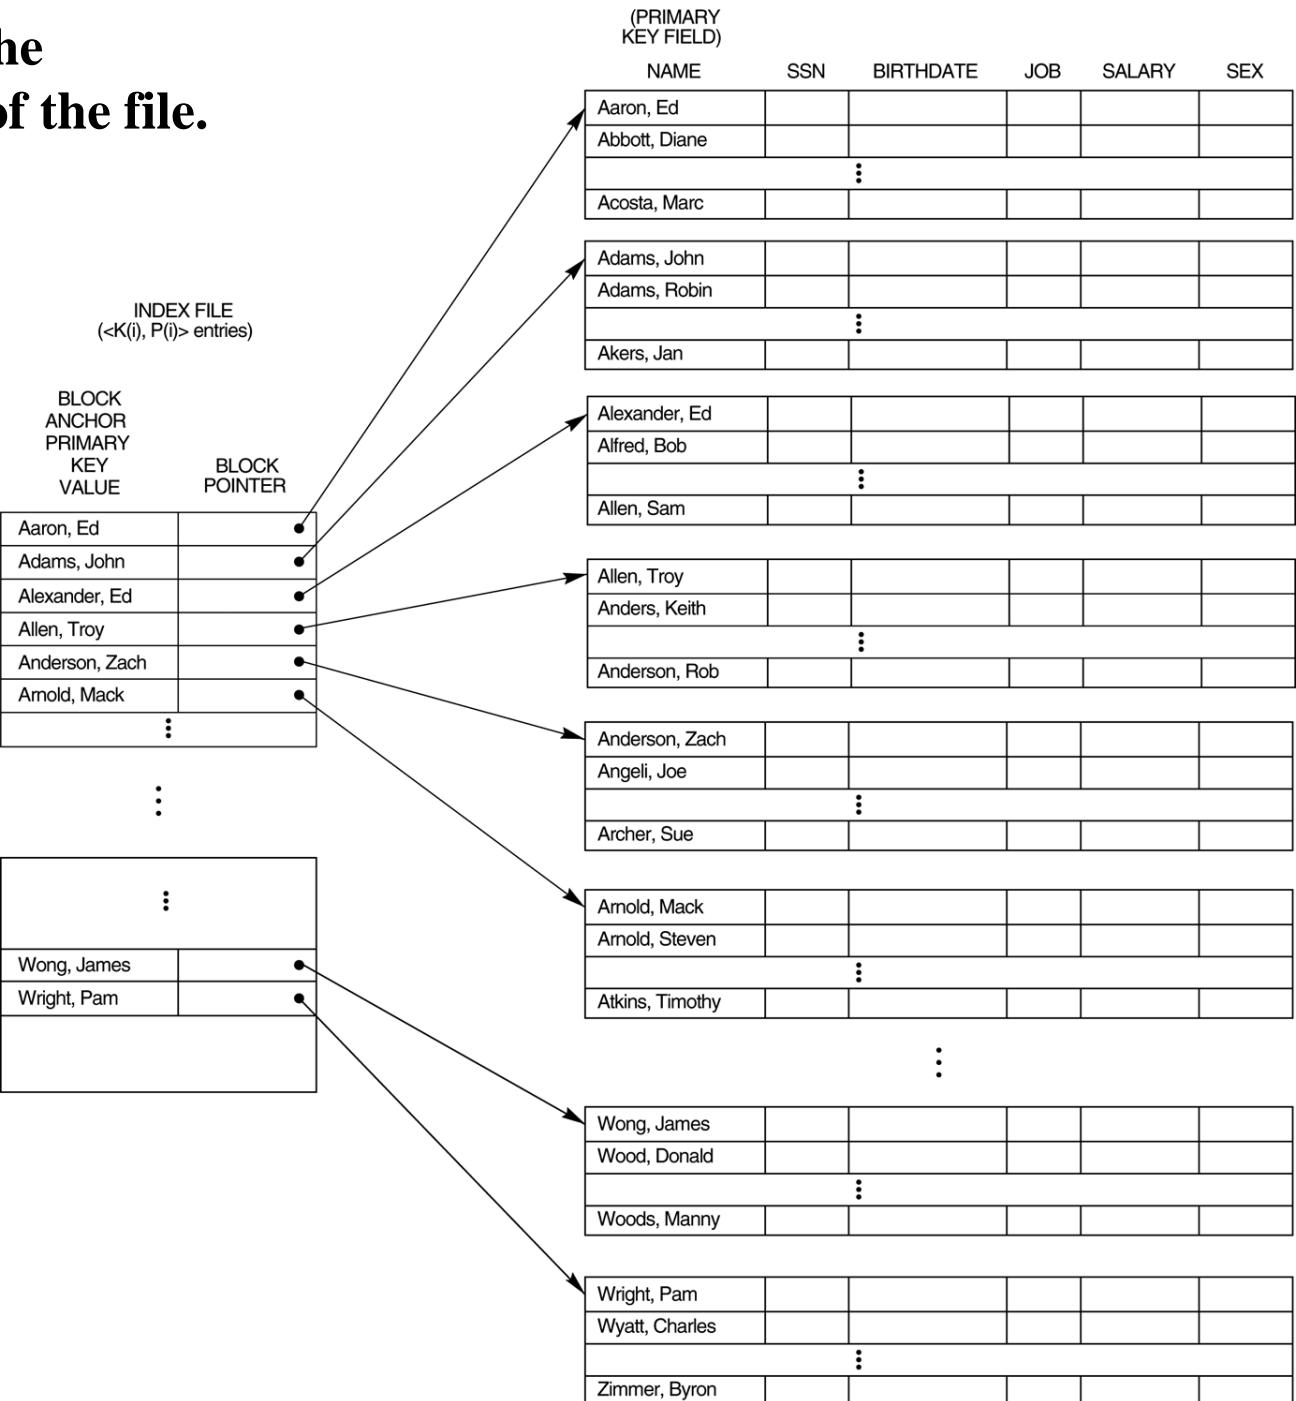
The **index** is an **ordered file** with two fields  
K(i): One index entry for each block → block anchor is PK value  
P(i): Block pointer

**Single-level index** on the ordering **key** field of the file.

This is a **primary index**, because data is ordered on the primary key field (no duplicates).

This is also a **clustered index** because the data is in the same order as the search key.

This is also a sparse  
(nondense) index,  
since it includes an  
entry for each disk  
block of the data file  
(not for each record of  
the data file).



# Important Concepts

- A **primary index** is specified on the *ordering key field* of an ordered file of records
  - (In Raghu's Book: An index on the primary key field is called a primary index)
- If the ordering of the index and data records is the same (or, close), we call this a **clustering index**
  - So, a primary index is also clustering!
  - We can also have clustering index on an *ordering non-key* field
  - But, since a file can have only one physical ordering, it can have **at most one** primary/clustering index

# Important Concepts

- A **dense index** has an index entry for *every* search key value (and hence, every record) in the data file.
- A **sparse (non-dense)** index has index entries for only *some* search values.

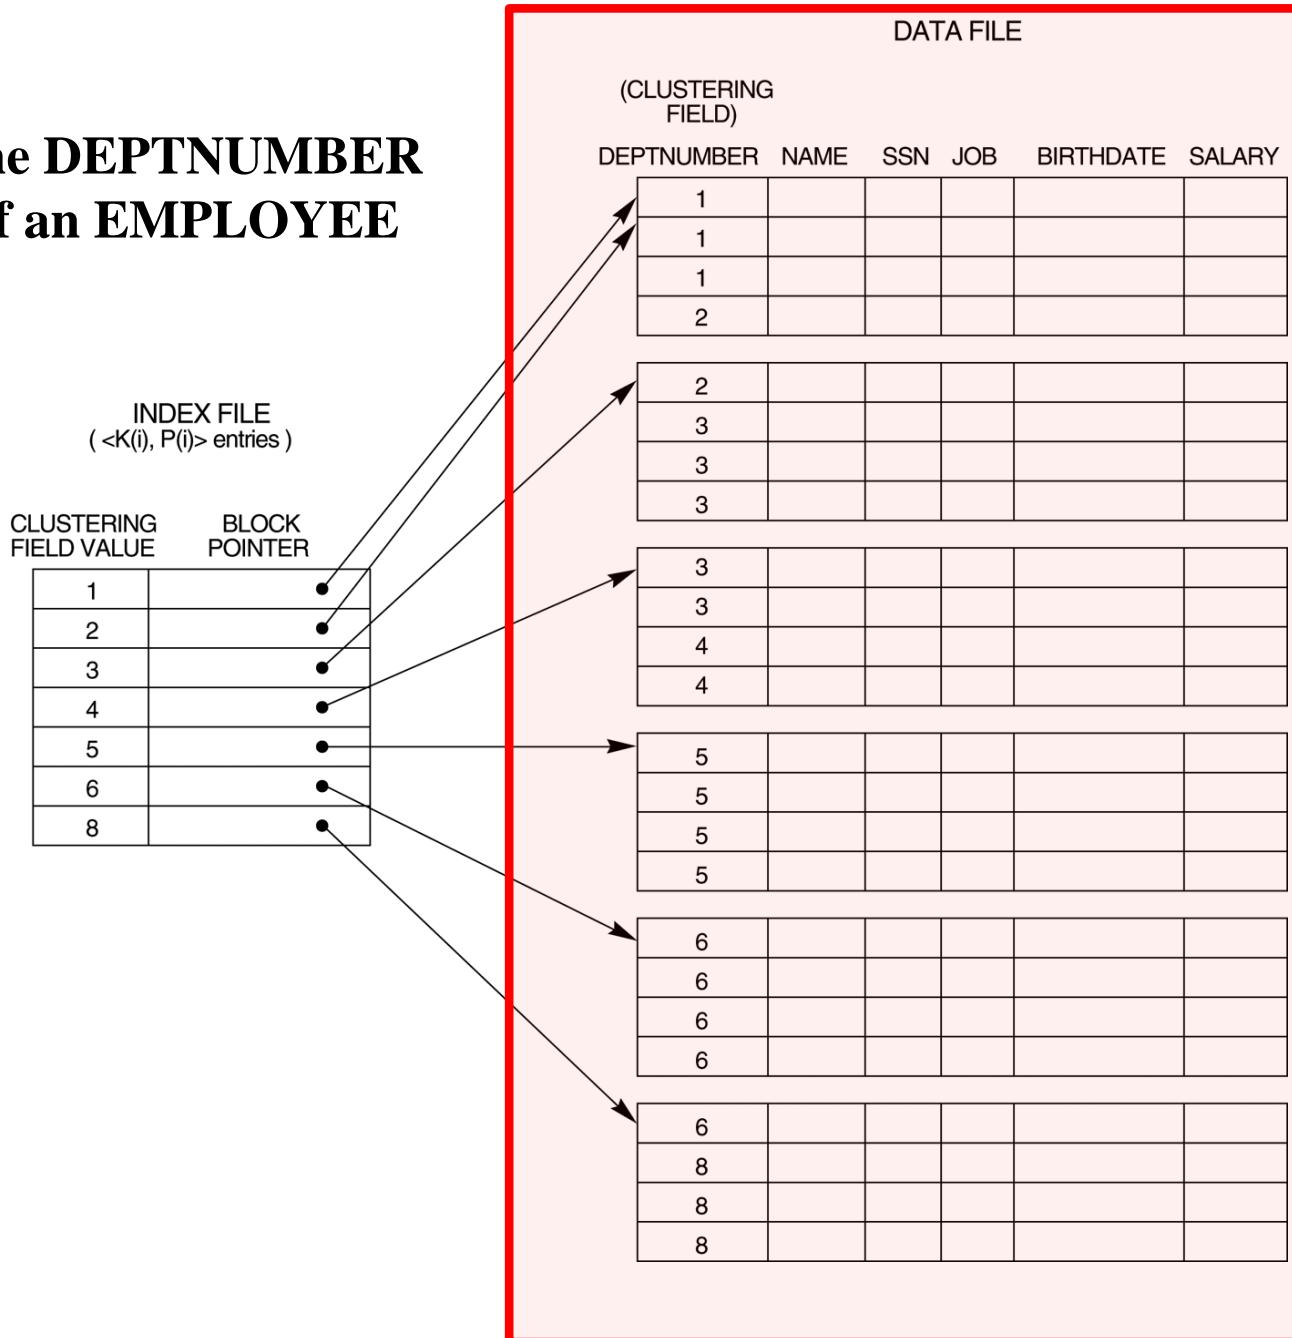
# Another Example of a Clustered Single-Level Index

- The data file is ordered on a *non-key field* (unlike primary index, which requires that the ordering field of the data file have a distinct value for each record).
- Includes one index entry *for each distinct value* of the field; the index entry points to the *first data block* that contains records with that field value.
- It is another example of *sparse (nondense)* index.

## Single-level index

A **clustering index** on the **DEPTNUMBER** ordering **nonkey** field of an **EMPLOYEE** file.

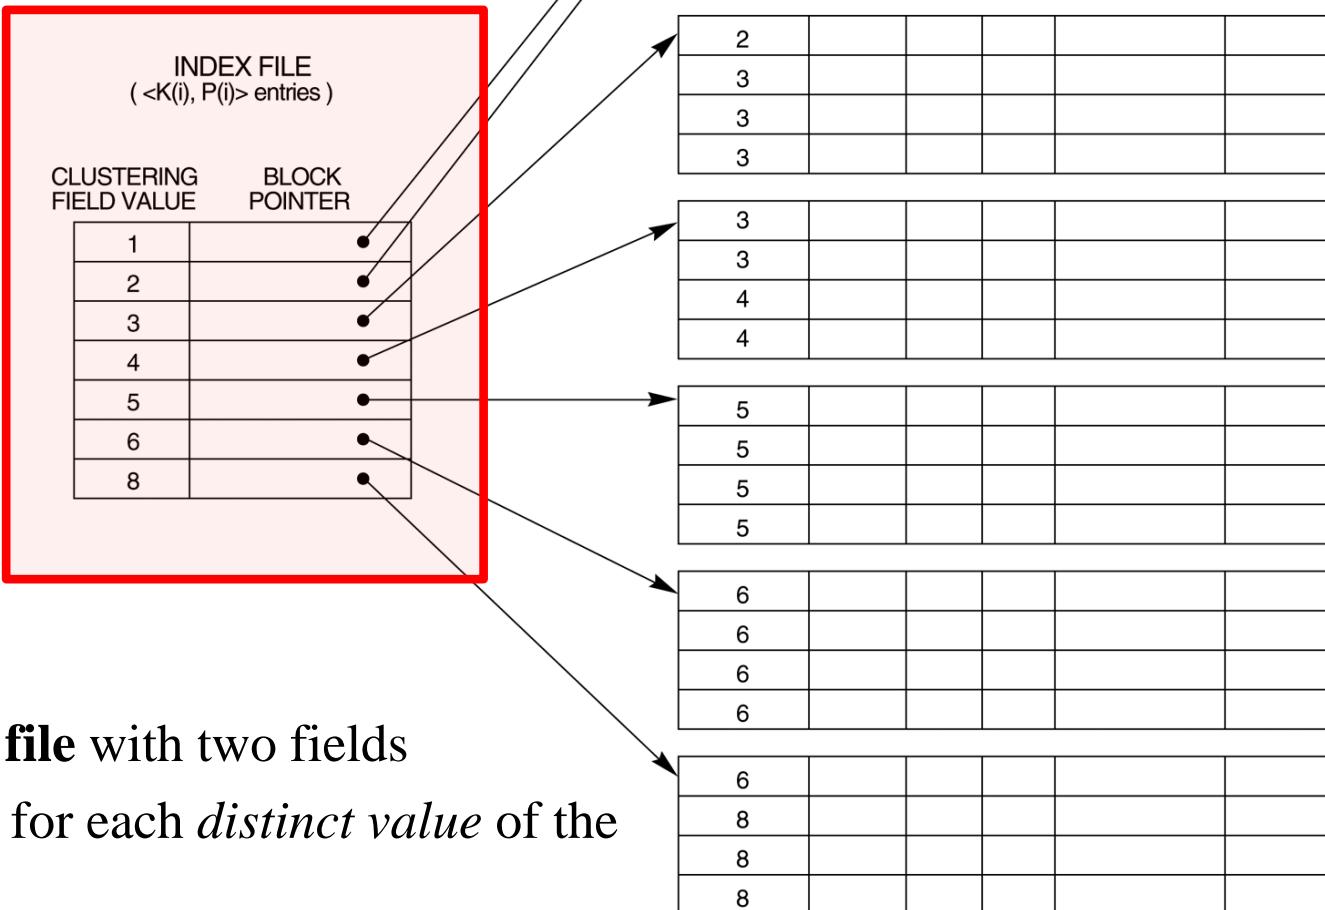
- Data file is ordered on the **non-key field**, i.e, **Dept**



## Single-level index

A **clustering index** on the DEPTNUMBER ordering **nonkey** field of an EMPLOYEE file.

- Data file is ordered on the **non-key field**, i.e, **Dept**



The **index** is an **ordered file** with two fields

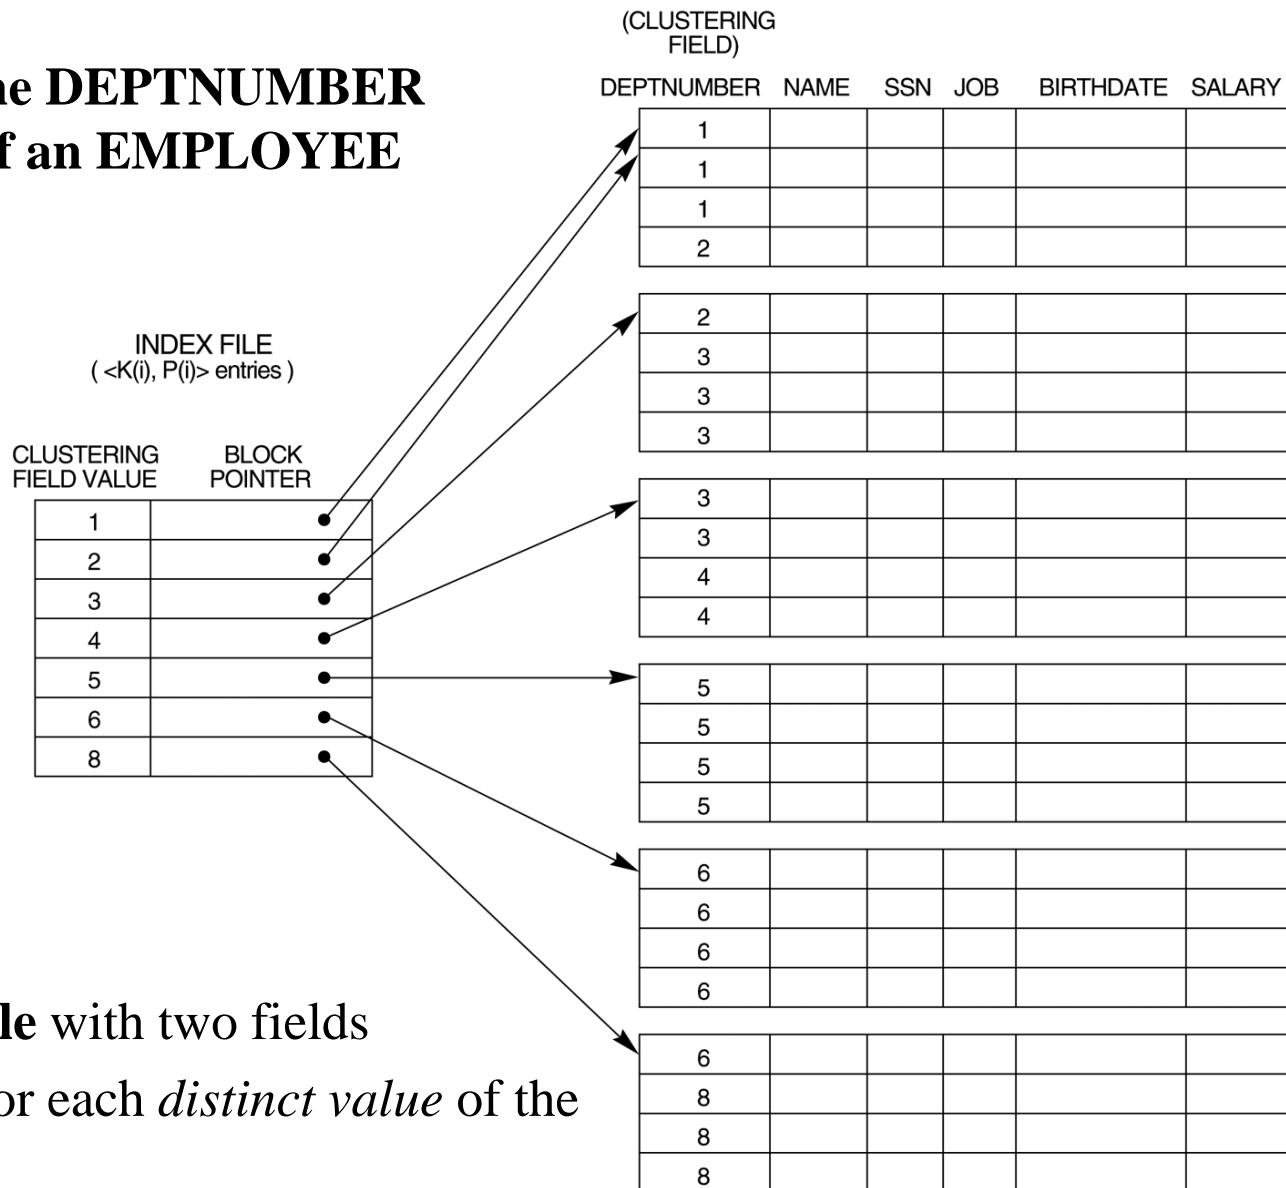
- $K(i)$ : One index entry for each *distinct value* of the field (e.g., Dept)
- $P(i)$ : Block pointer

## Single-level index

A **clustering index** on the DEPTNUMBER ordering **nonkey** field of an EMPLOYEE file.

- Data file is ordered on the **non-key field**, i.e., **Dept**

```
Select *
From Emp E
Where E.Dept = 2
```



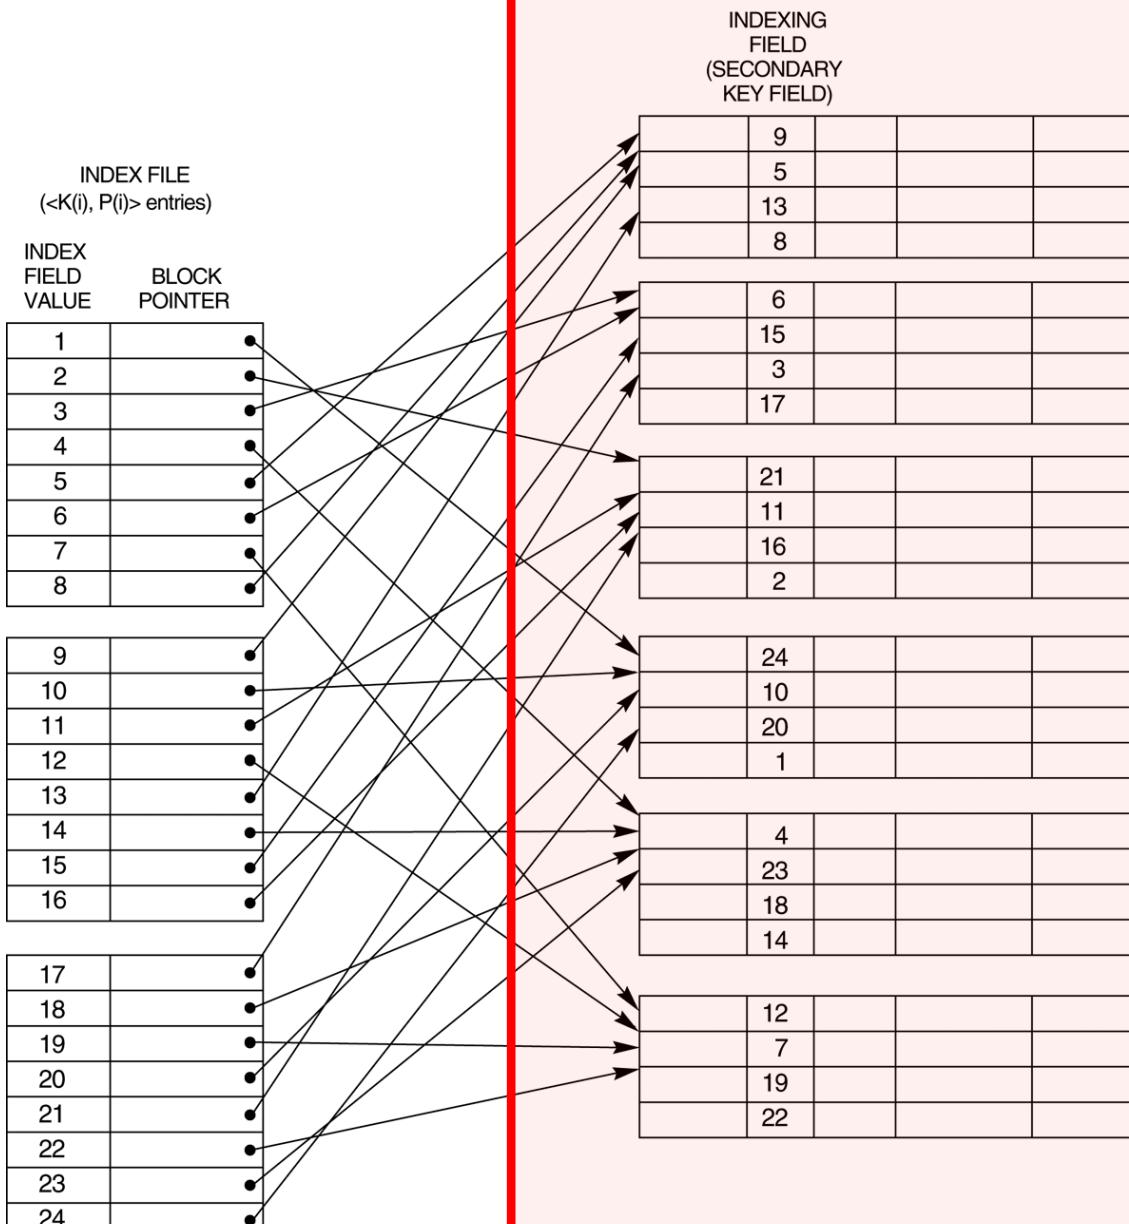
The **index** is an **ordered file** with two fields

- $K(i)$ : One index entry for each *distinct value* of the field (e.g., Dept)
- $P(i)$ : Block pointer

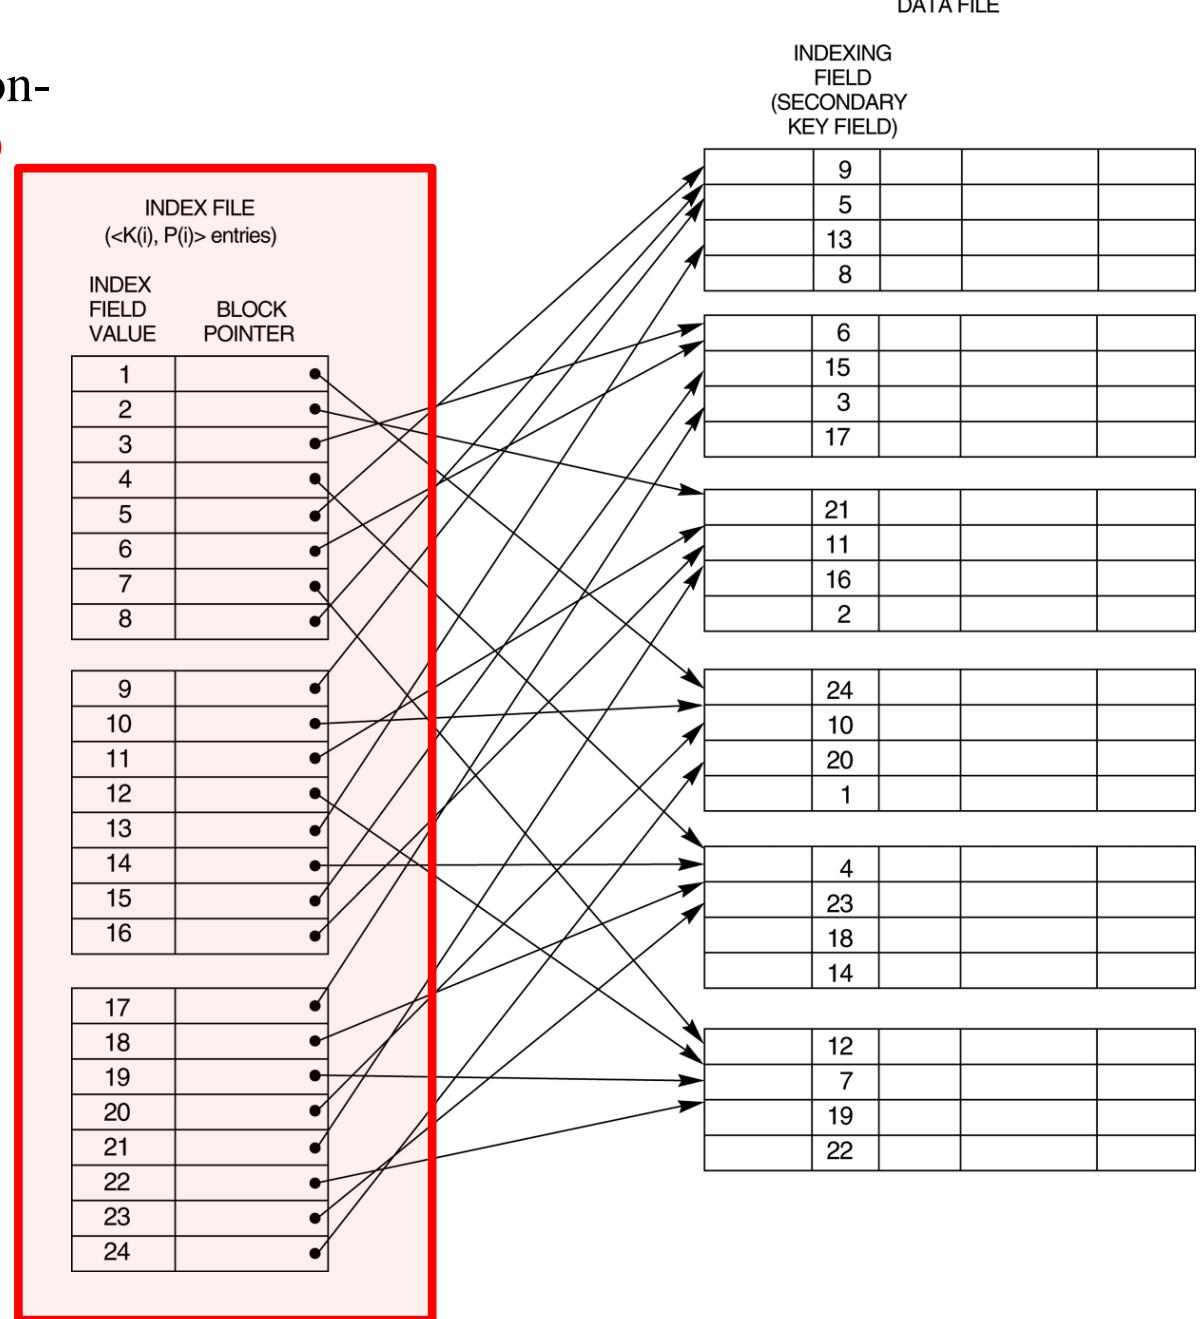
# Single-Level Secondary Index

- A **secondary index** provides a secondary means of accessing a file for which some primary access already exists.
- The secondary index may be on a *non-ordering* field that is either
  - a **candidate key** and has a unique value in every record, or
  - a **nonkey** with duplicate values.
- The **index** is an ordered file with two fields:
  - The first field is the *indexing field*.
  - The second field is either a *block* pointer or a *record* pointer.
- Includes one entry *for each record* in the data file; hence, it is a **dense index**.
- There can be *many* secondary indexes for the same file.

## A **secondary index** on a (non-ordering) **candidate key**, ID field



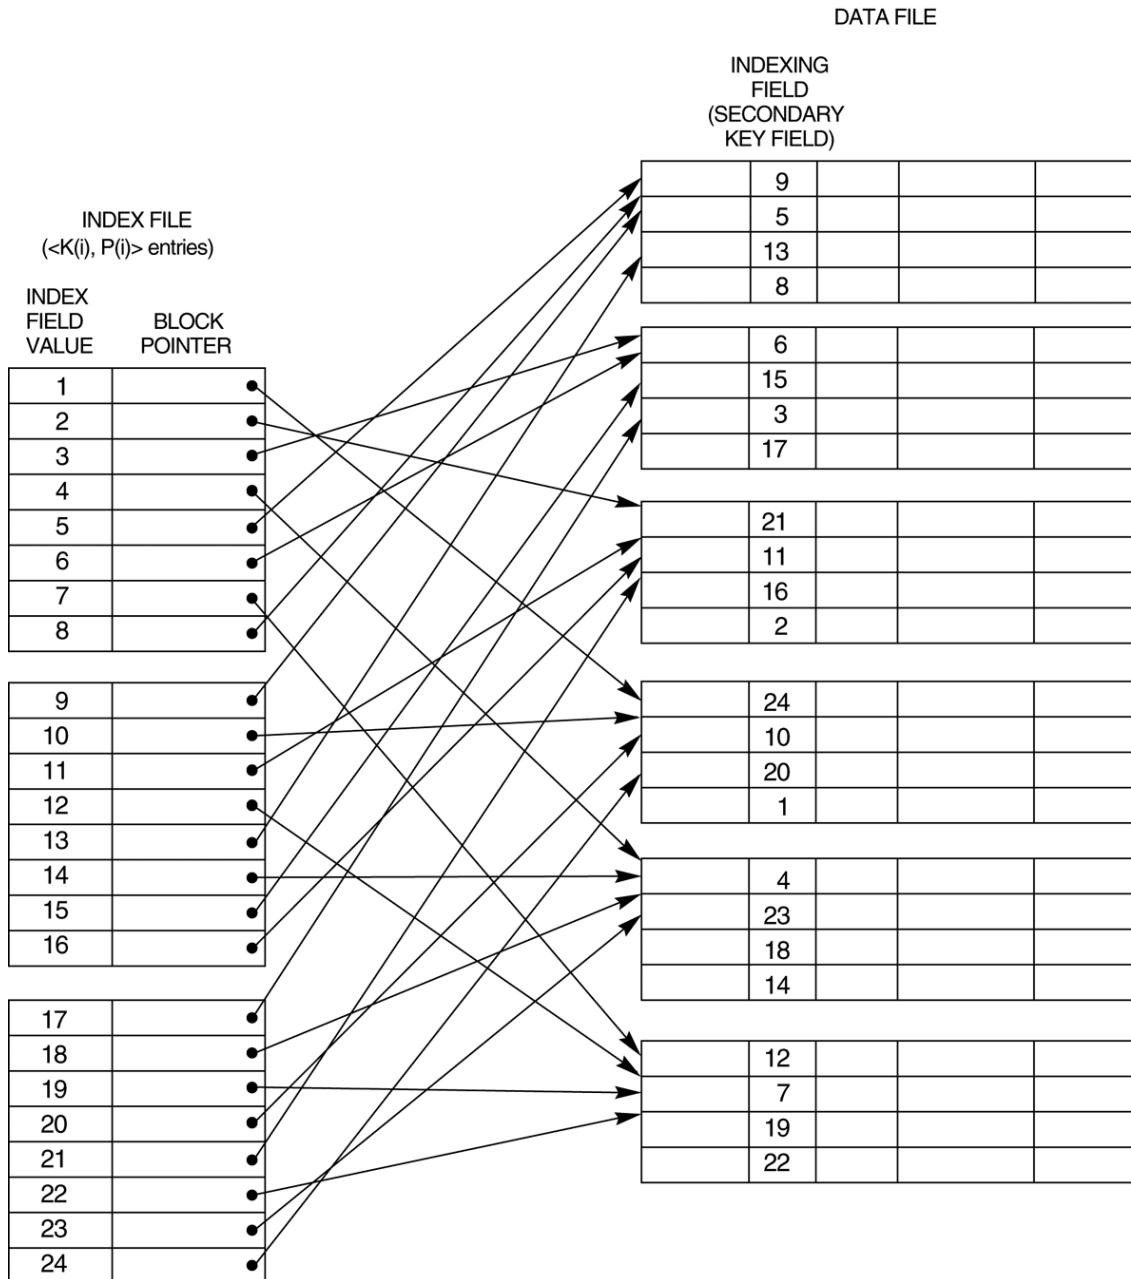
A **secondary index** on a (non-ordering) **candidate key**, ID field



A **secondary index** on a (non-ordering) **candidate key**, ID field

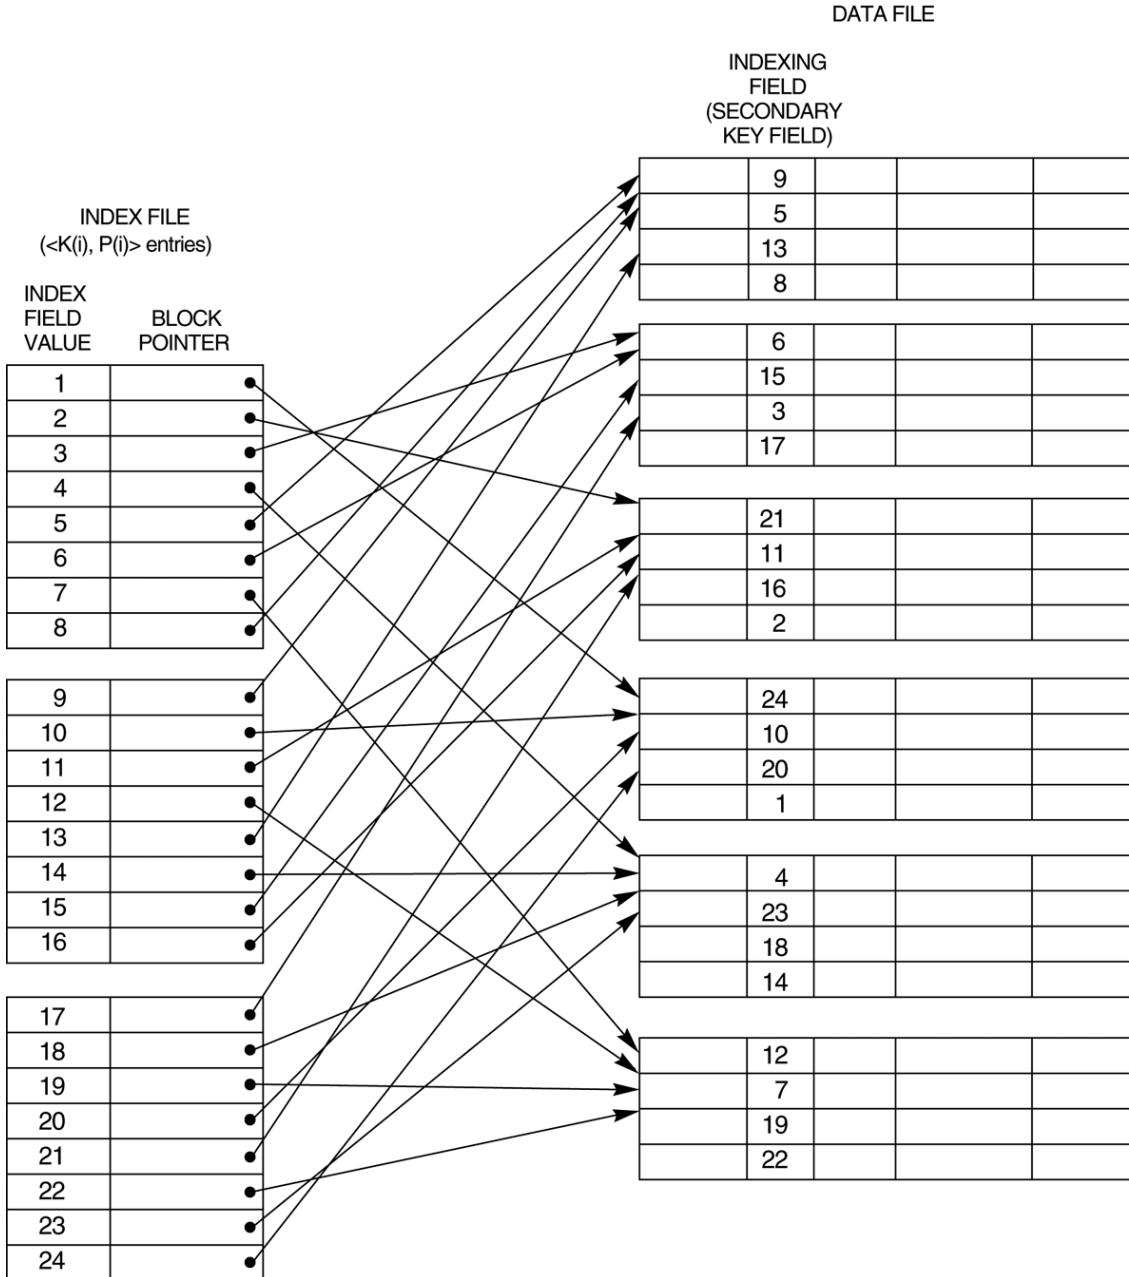
Select \*  
From Emp E  
Where E.ID = 11

The **index** is an **ordered file** with two fields  
**K(i)**: One index entry per record, ie., the candidate key  
**P(i)**: Block pointer



This is a **dense** index.  
No duplicates.

Note that the data file is *not* ordered according to the index field.  
Therefore it is an **unclustered** index



# We can have multiple index on the same file

Secondary index on ID

DATA FILE

INDEXING FIELD  
(SECONDARY KEY FIELD)

INDEX FILE  
( $K(i), P(i)$  entries)

INDEX FIELD  
VALUE      BLOCK  
BLOCK  
POINTER

1	•
2	•
3	•
4	•
5	•
6	•
7	•
8	•

9	•
10	•
11	•
12	•
13	•
14	•
15	•
16	•

17	•
18	•
19	•
20	•
21	•
22	•
23	•
24	•

9			
5			
13			
8			

6			
15			
3			
17			

21			
11			
16			
2			

24			
10			
20			
1			

4			
23			
18			
14			

12			
7			
19			
22			

Data sorted on Name  
Primary index on Name

BLOCK  
ANCHOR  
PRIMARY  
KEY  
VALUE

Aaron, Ed
Adams, John
Alexander, Ed
Allen, Troy
Anderson, Zach
Arnold, Mack

⋮

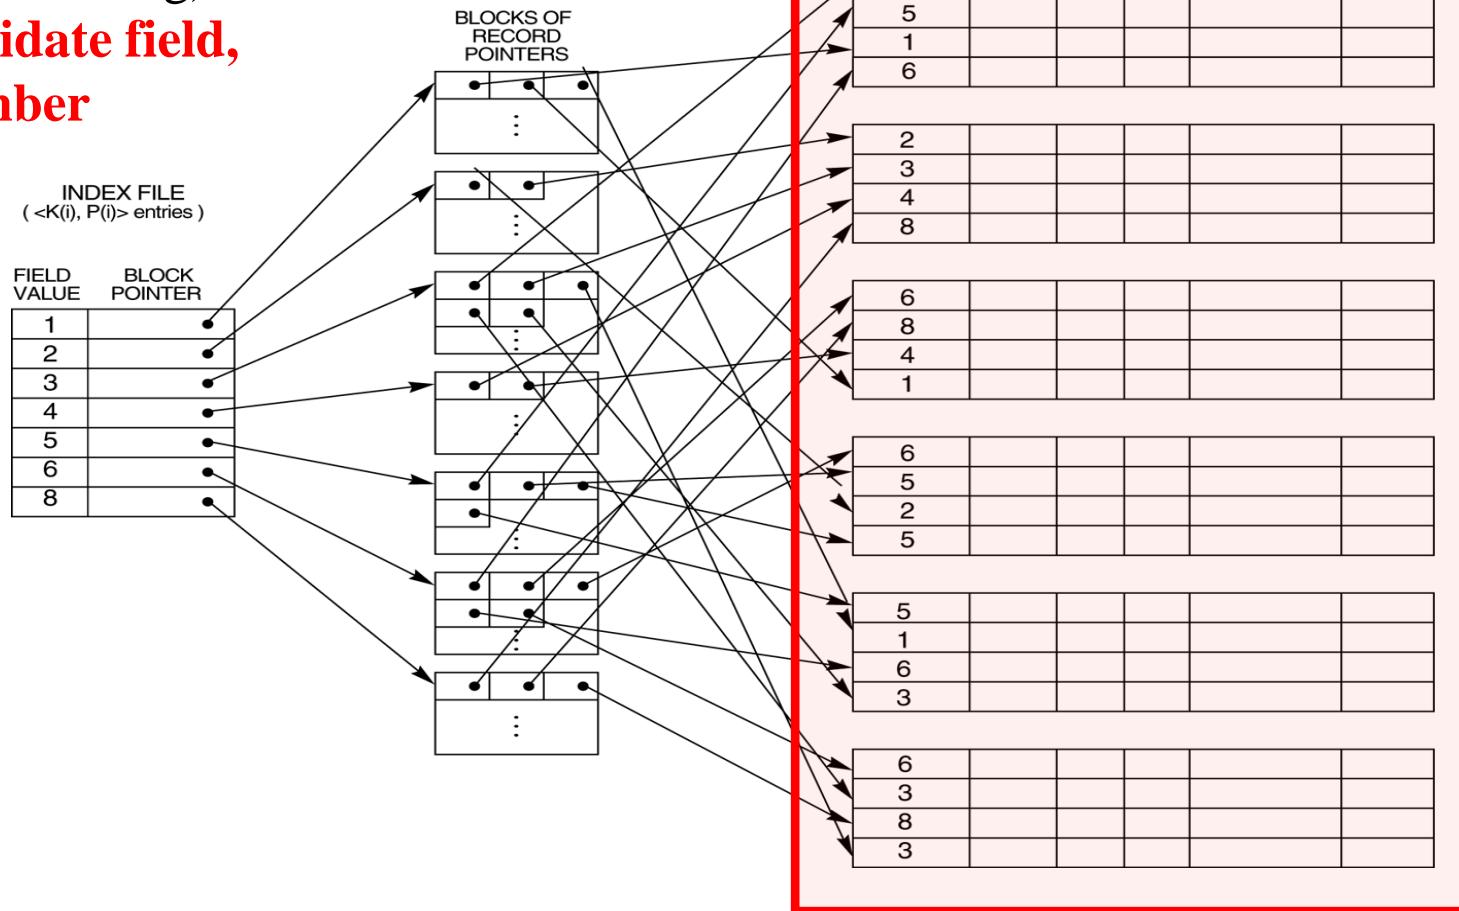
⋮
⋮

Wong, James
Wright, Pam

⋮

A **secondary index (with record pointers)** on a nonkey field implemented using one level of indirection so that index entries are of fixed length and have unique field values. This is an **unclustered** index.

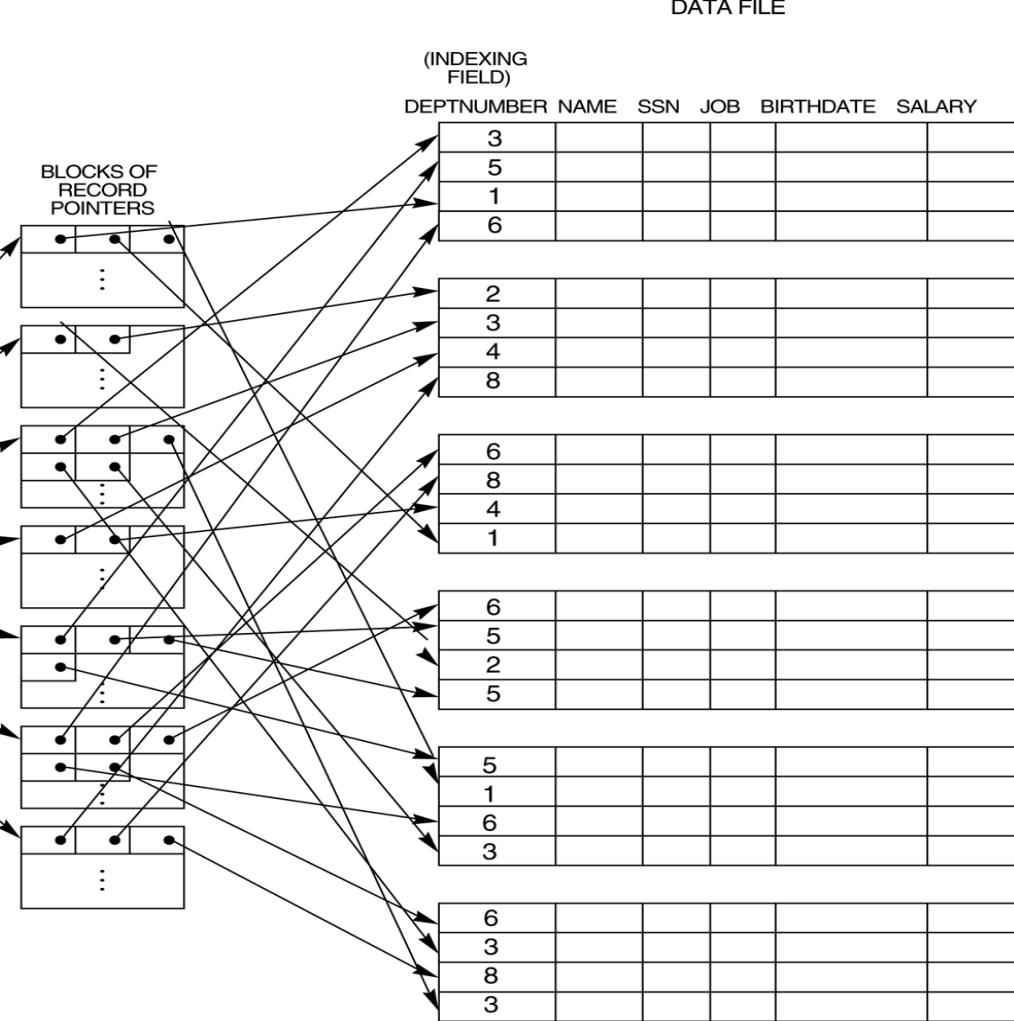
A **secondary index**  
on a (non-ordering)  
**non candidate field,**  
**DeptNumber**



A **secondary index (with record pointers)** on a nonkey field implemented using one level of indirection so that index entries are of fixed length and have unique field values. This is an **unclustered** index.

A **secondary index**  
on a (non-ordering)  
**non candidate field,**  
**DeptNumber**

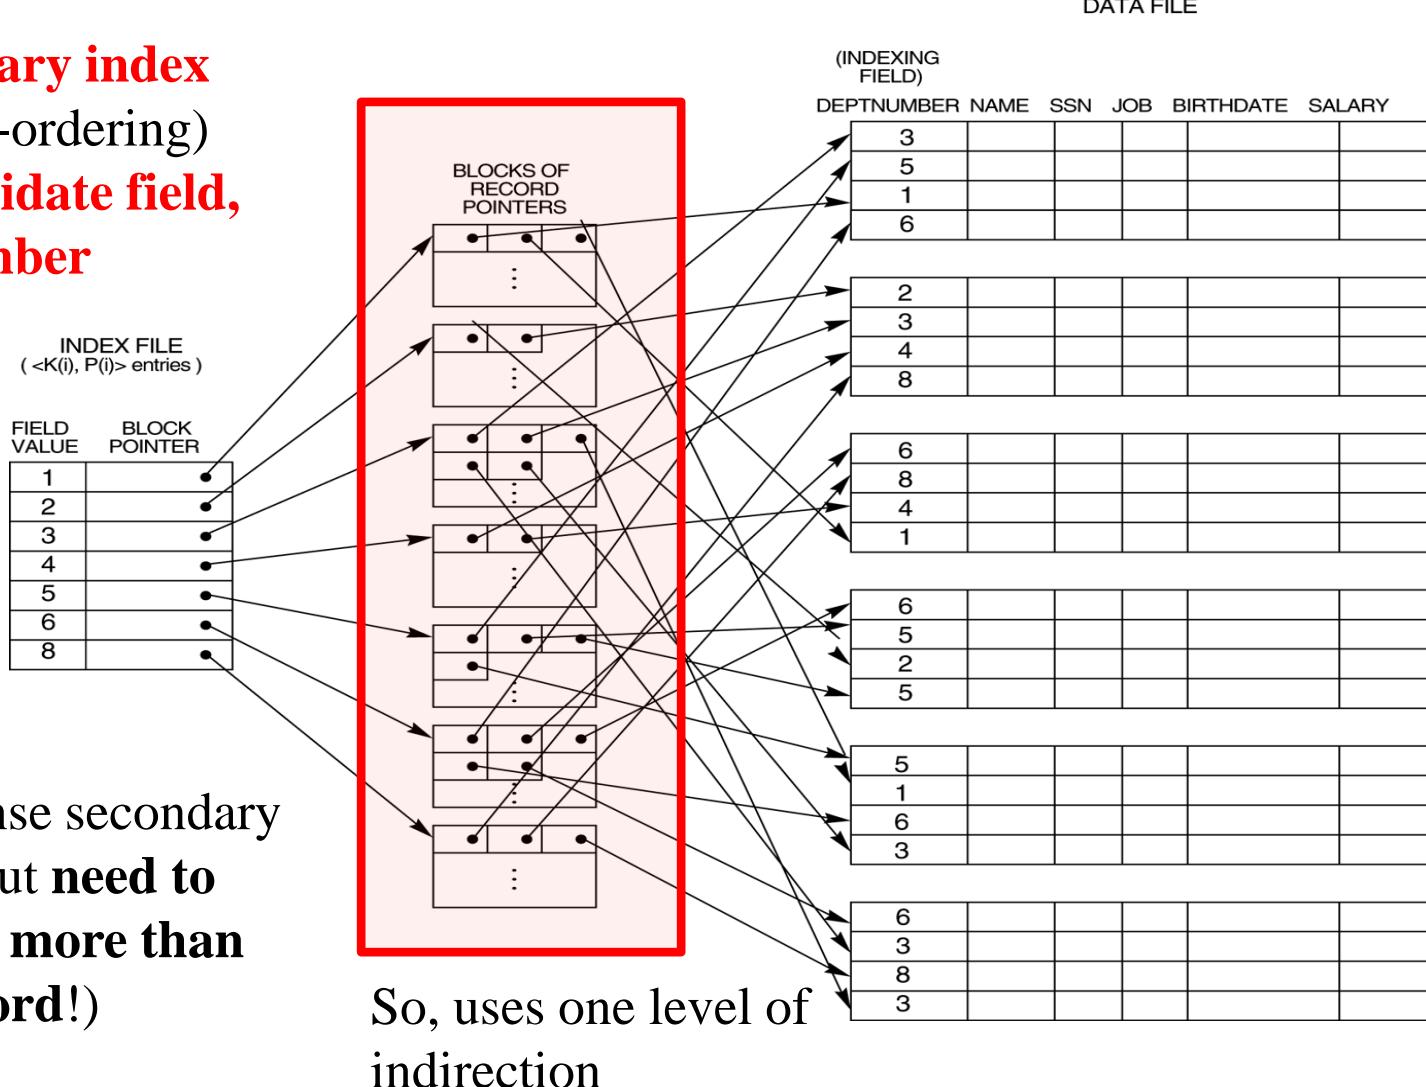
INDEX FILE ( $<K(i), P(i)>$ entries )	
FIELD VALUE	BLOCK POINTER
1	•
2	•
3	•
4	•
5	•
6	•
8	•



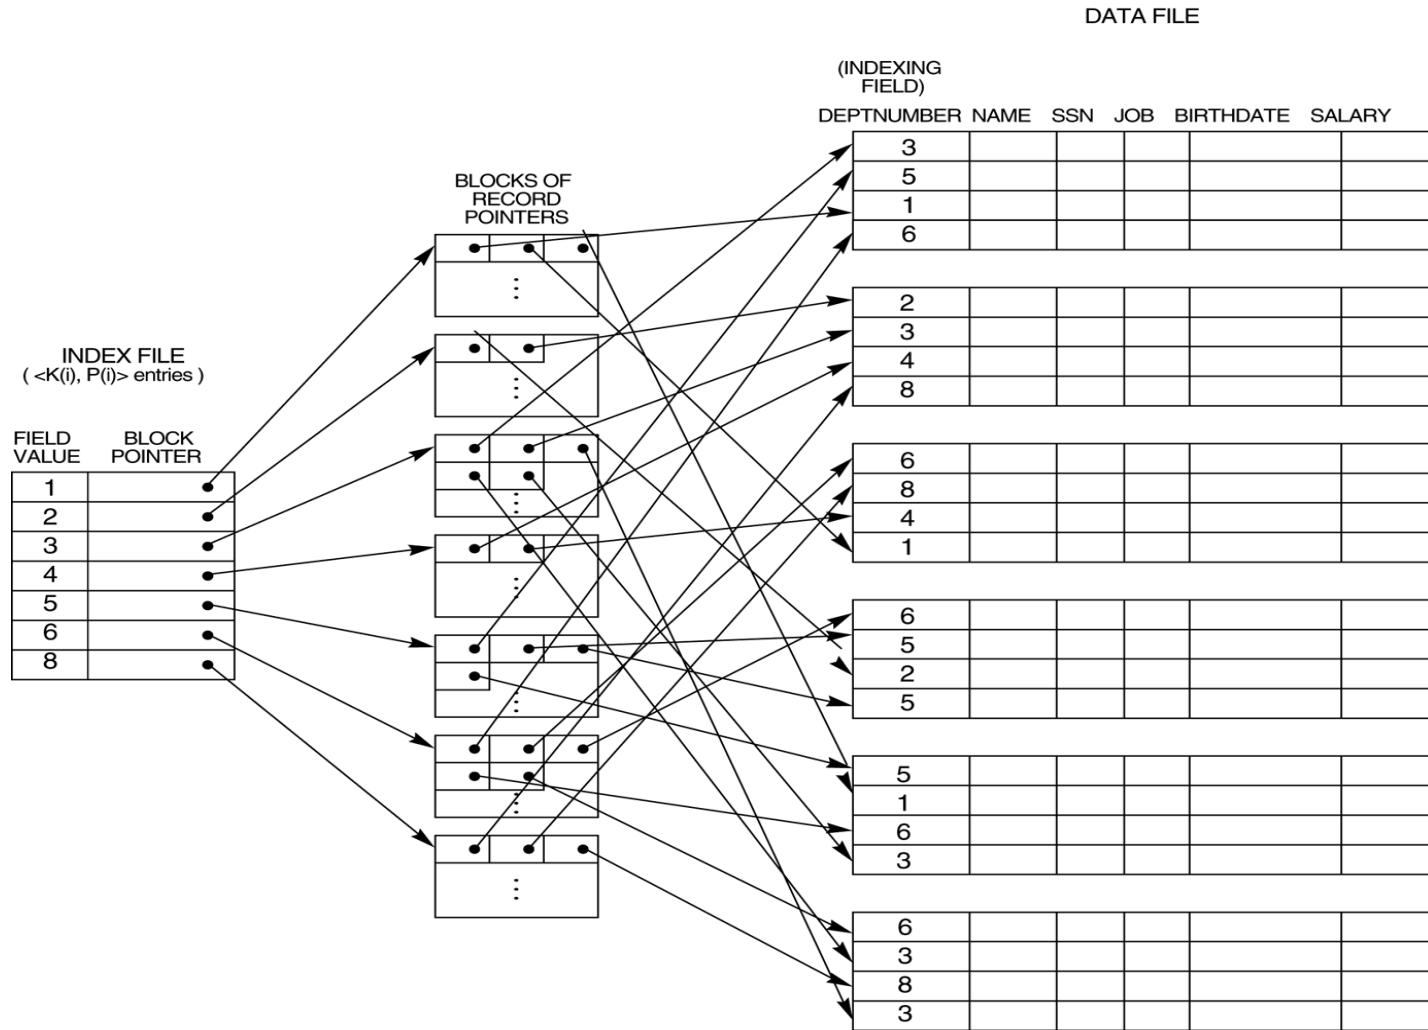
Non-dense secondary  
index (but need to  
point more than one  
record!)

A **secondary index (with record pointers)** on a nonkey field implemented using one level of indirection so that index entries are of fixed length and have unique field values. This is an **unclustered** index.

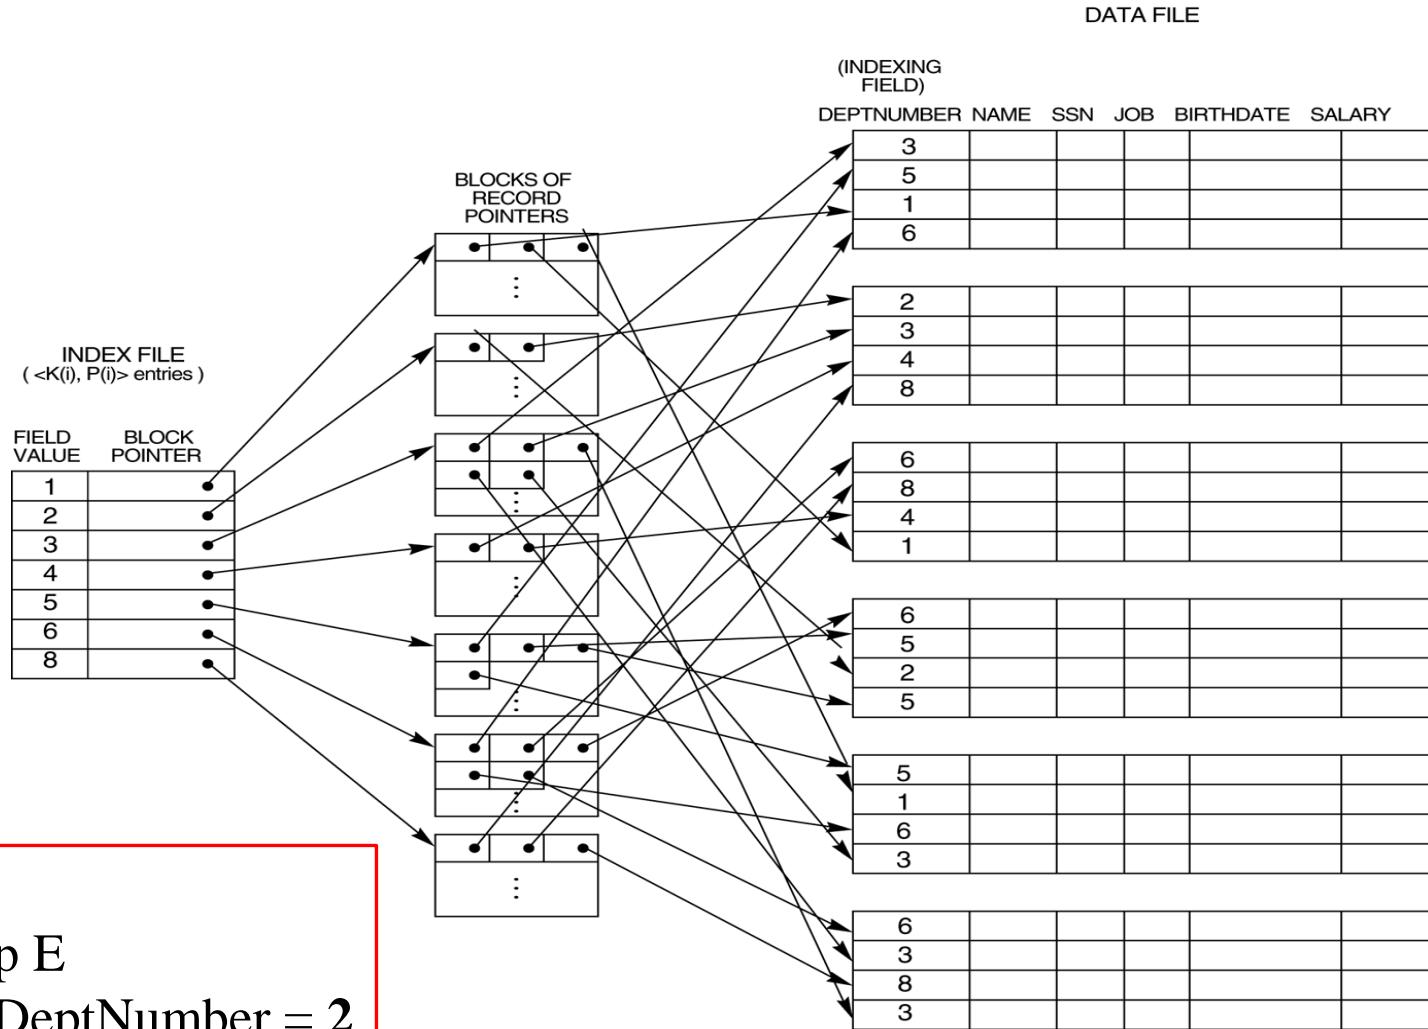
A **secondary index**  
on a (non-ordering)  
**non candidate field,**  
**DeptNumber**



A **secondary index (with record pointers)** on a nonkey field implemented using one level of indirection so that index entries are of fixed length and have unique field values. This is an **unclustered** index.



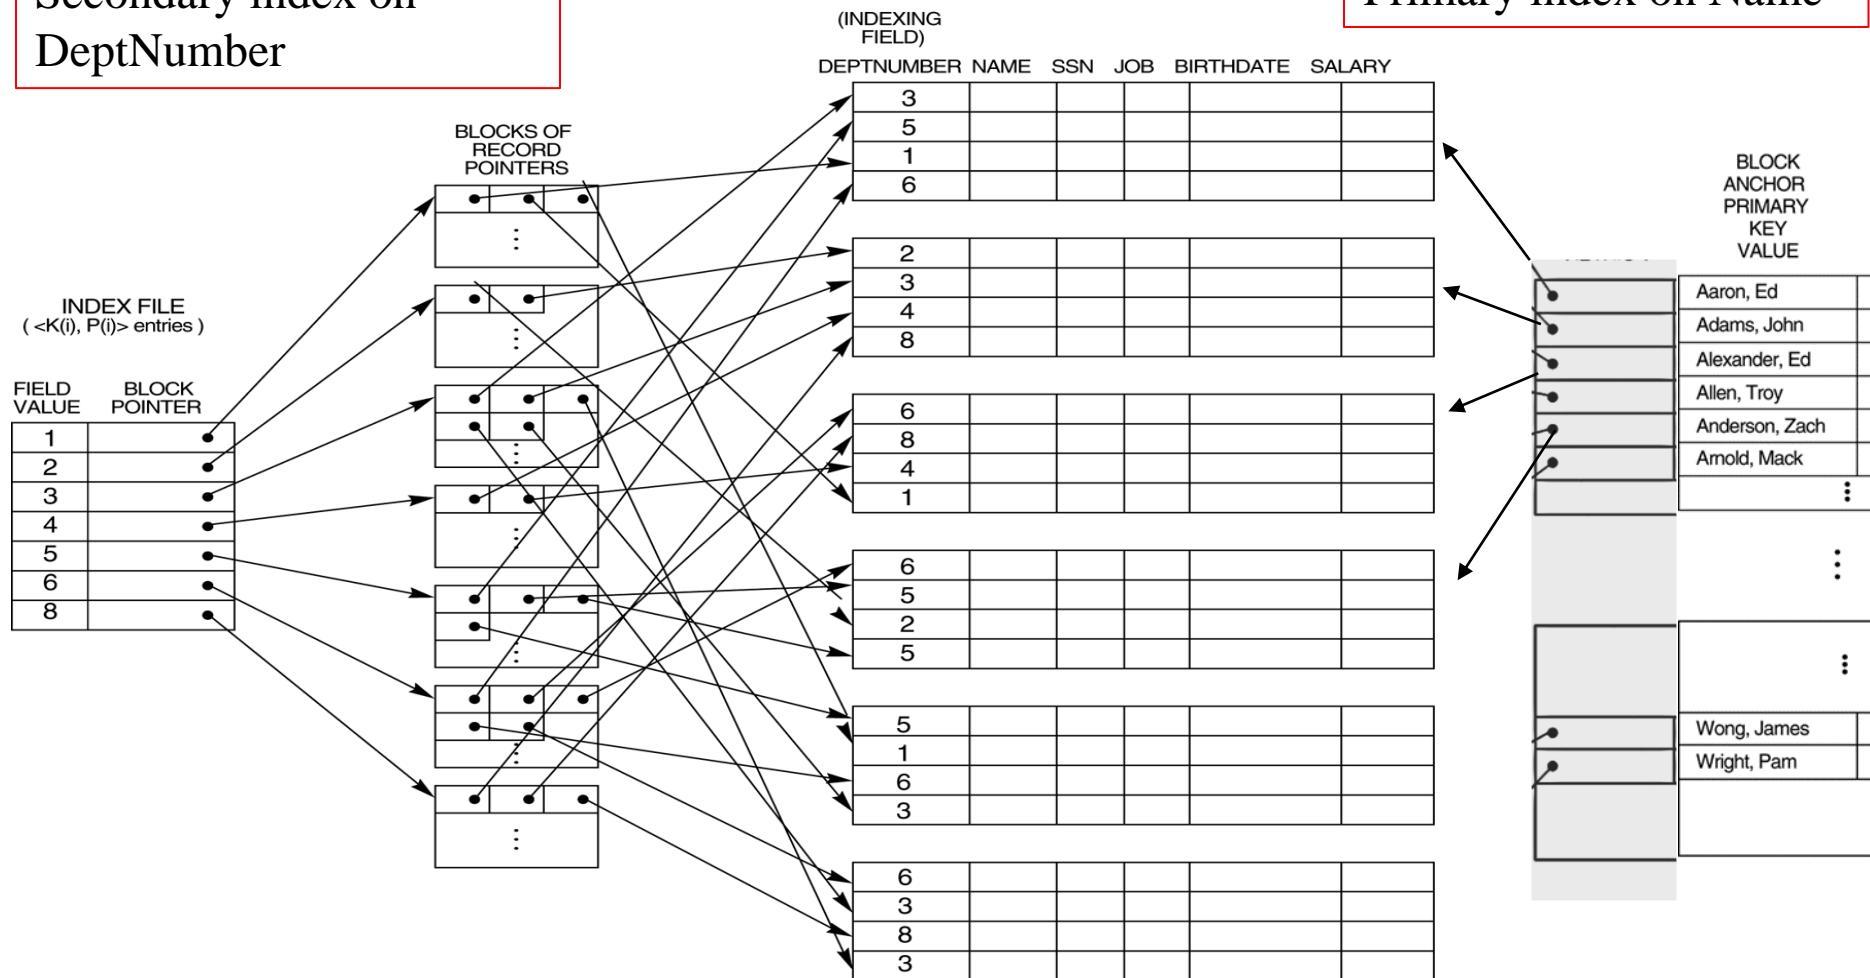
A **secondary index (with record pointers)** on a nonkey field implemented using one level of indirection so that index entries are of fixed length and have unique field values. This is an **unclustered** index.



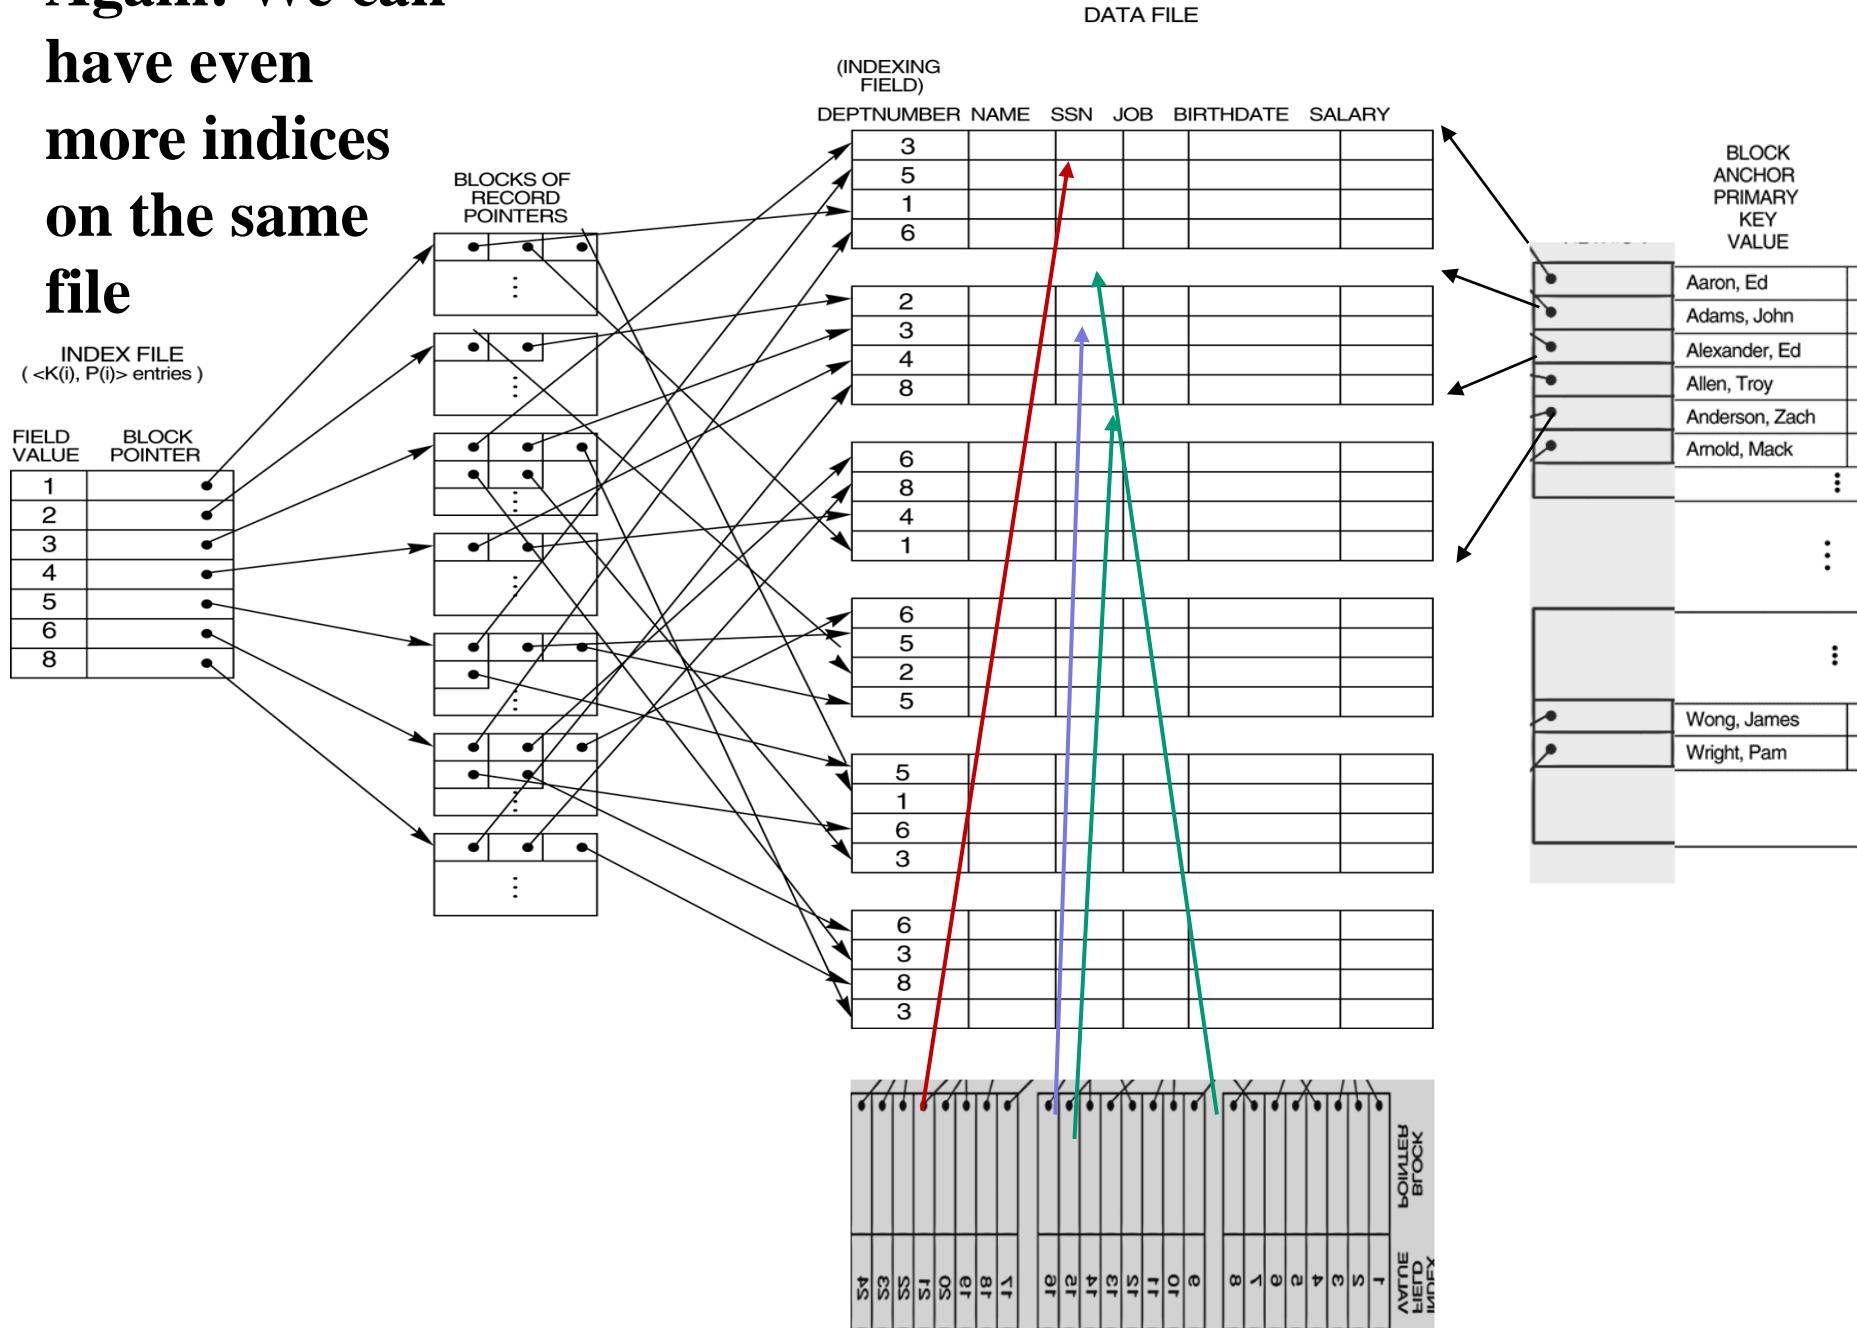
# Again: We can have multiple index on the same file

Secondary index on  
DeptNumber

Data sorted on Name  
Primary index on Name



**Again: We can have even more indices on the same file**



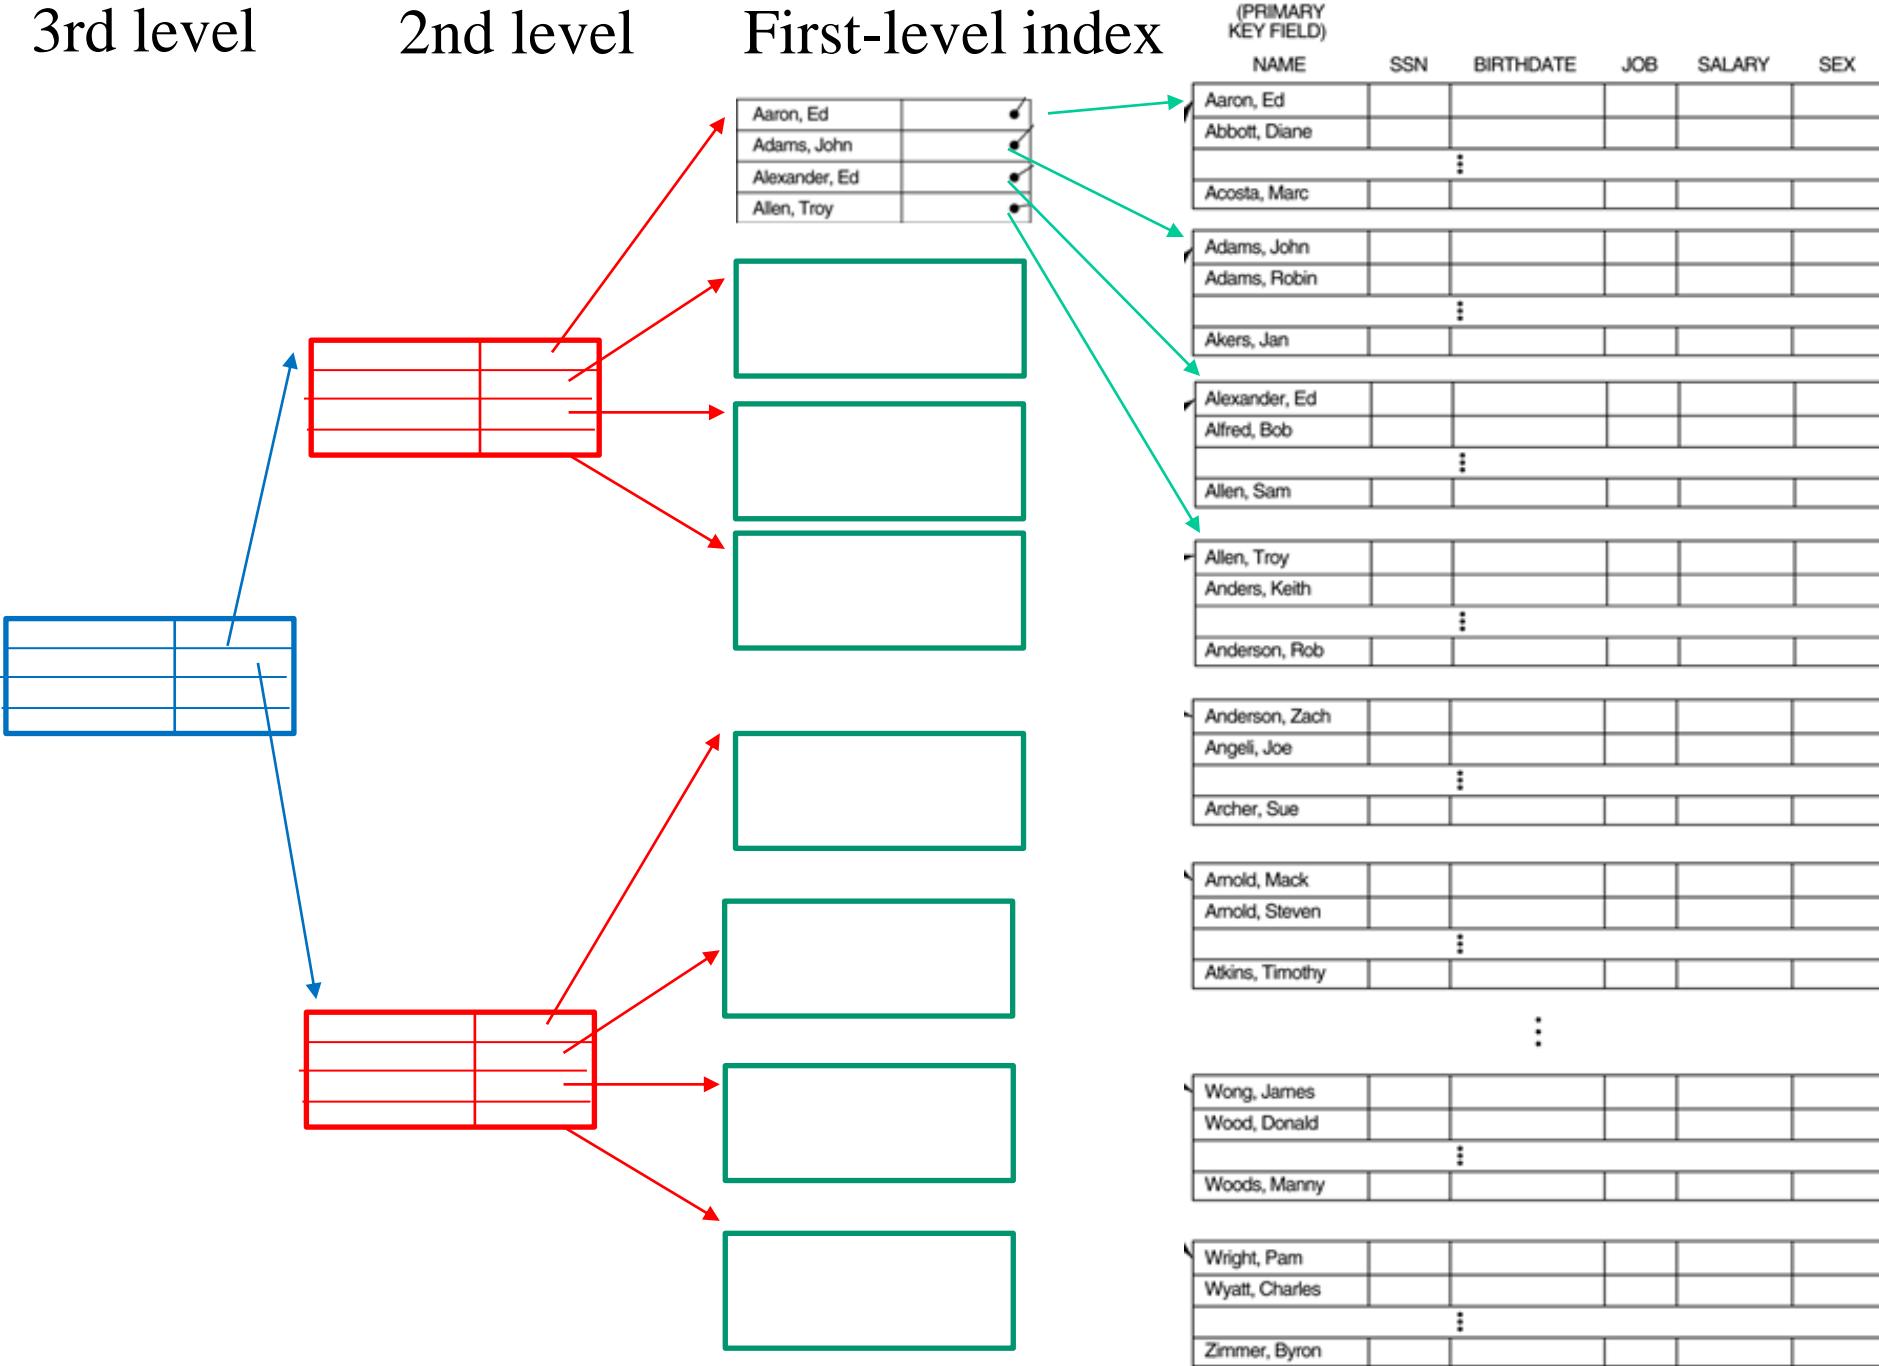
# Multi-Level Indexes

- Because a single-level index is an ordered file, we can create an index *to the index itself*; in this case, the original index file is called the *first-level index* and the index to the index is called the *second-level index*.
- We can repeat the process, creating a third, fourth, ..., top level until all entries of the *top level* fit in one disk block
- A multi-level index can be created for any type of first-level index (primary, secondary, clustering).

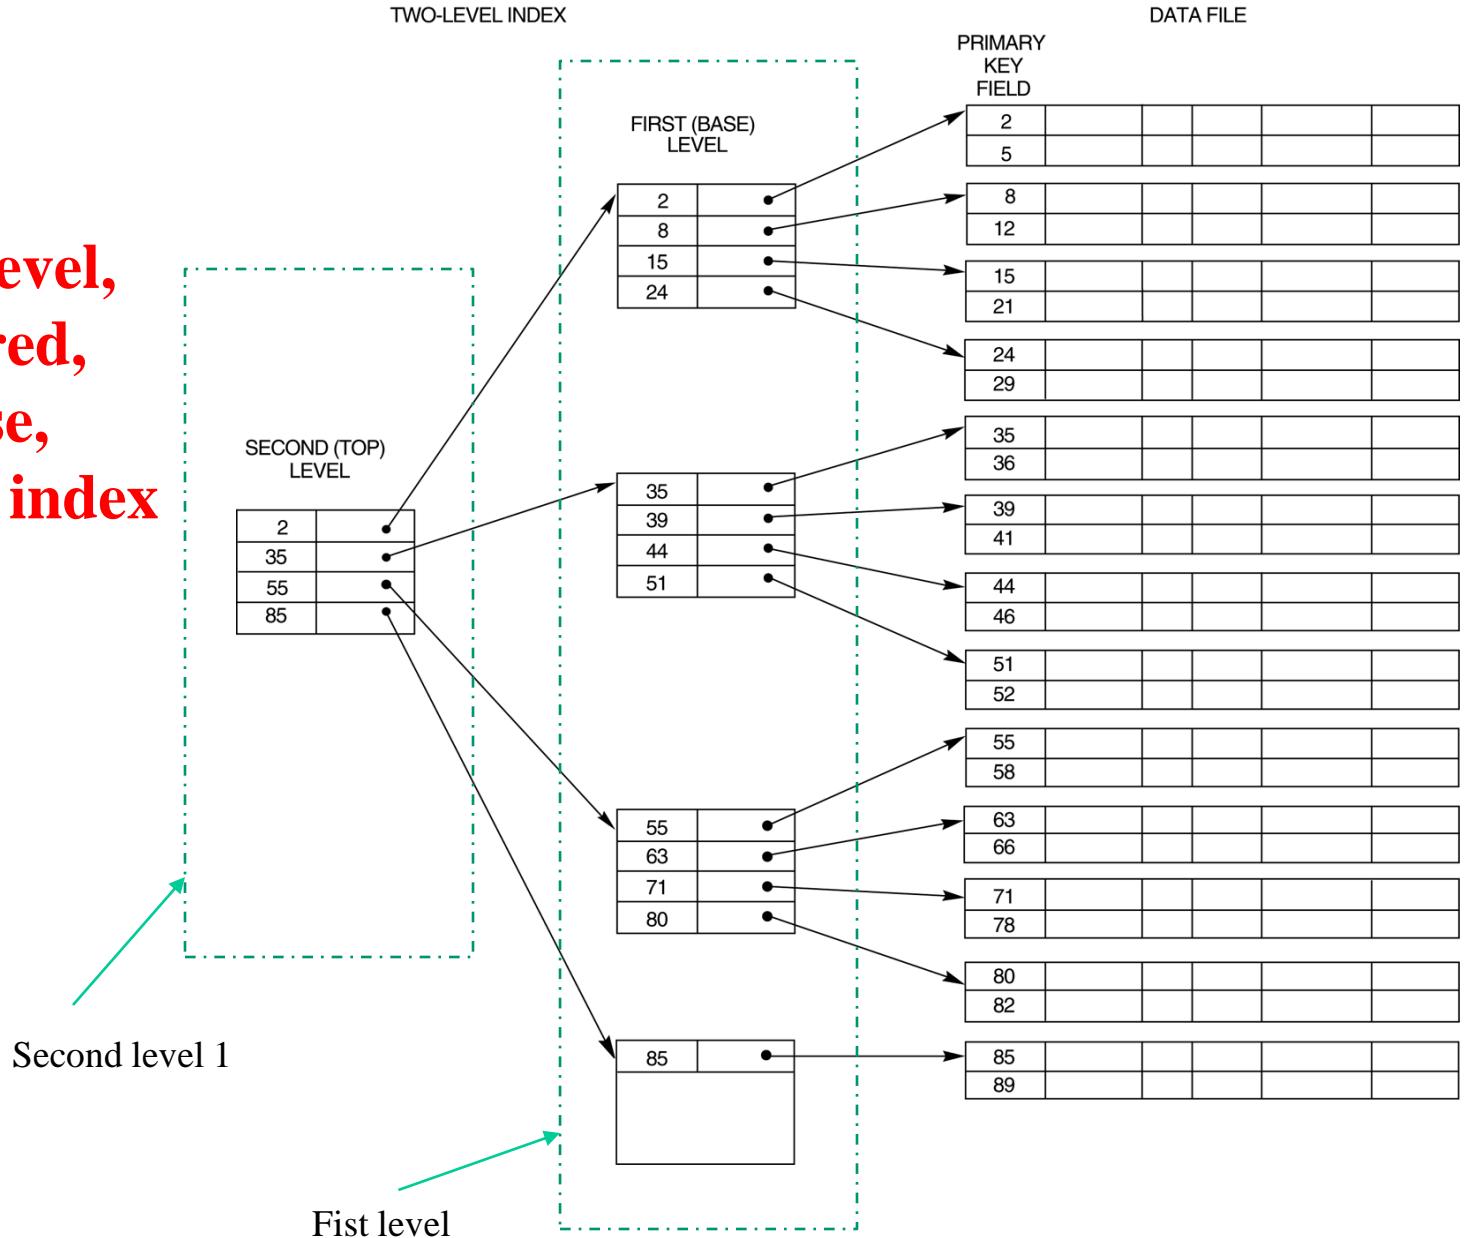
## 3rd level

## 2nd level

## First-level index



**A two-level,  
clustered,  
sparse,  
primary index**



# Multi-Level Indexes

- Such a multi-level index is a form of *search tree*; however,
  - insertion and
  - deletion ofnew index entries is a severe problem because every level of the index is an *ordered file*.
- So this brings us to B+tree index structure.