

# **Relational Model**

## Chapter 3

# Relational Database: Definitions

- Relational database: a set of *relations*
- Relation: made up of 2 parts:
  - **Schema** : specifies name of relation, plus name and type of each column.
    - e.g. **Student**(**sid**: string, **name**: string, **login**: string, **age**: integer, **gpa**: real).
  - **Instance** : a table, with rows and columns.  
#Rows = *cardinality*, #fields = *degree (arity)*.

# Relational Database: Definitions

- We can think of a relation as a **set** of *rows* or *tuples* (i.e., all rows are distinct).
- In relational model, a **field** can only contain atomic values, i.e., lists or sets are **disallowed!**

# Example Instance of Student Relation

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- Cardinality = 3, degree = 5, all rows distinct
- Do all columns in a relation instance have to be distinct?
- While depicting a relation instance, order of rows doesn't matter!

# Alternative Terminology for Relational Model

Formal Terms	Alternative 1	Alternative 2
Relation	Table	File
Tuple	Row	Record
Attribute	Column	Field

# A Relational Database Schema for a University

- **Student**(**sid**: string, **name**: string, **login**: string, **age**: integer, **gpa**:real)
- **Course**(**cid**: string, **cname**:string, **credits**:integer)
- **Enrolled**(**sid**:string, **cid**:string, **grade**:string)
- This is a **conceptual schema**.

# Query Languages

- A major strength of DBMSs:
  - **Data Definition Language** (DDL) is used to set up schema
  - **Data Manipulation Language** (DML) is used to modify and query the data
- SQL:
  - Standard language for creating, manipulating, and querying data in a **relational DBMS**

# Relational Query Languages

- A major strength of the relational model:  
supports simple, powerful *querying* of data.
- Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
  - Precise semantics for relational queries.
  - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.

# The SQL Query Language

- Developed by IBM (system R) in the 1970s
- Need for a standard since it is used by many vendors
- Standards:
  - SQL-86
  - SQL-89 (minor revision)
  - SQL-92 (major revision, current standard)
  - SQL-99 (major extensions)
  - SQL-2003, SQL-2006 (extensions on XML)
  - SQL- 2008
  - SQL-2011

# Creating (Declaring) a Relation

- Simplest form is:

```
CREATE TABLE <name> (  
    <list of elements>  
)
```

- To delete a relation:

```
DROP TABLE <name>;
```

# Elements of Table Declarations

- The most common types are:
  - INT or INTEGER (synonyms).
  - REAL or FLOAT (synonyms).
  - CHAR( $n$ ) = fixed-length string of  $n$  characters.
  - VARCHAR( $n$ ) = variable-length string of up to  $n$  characters.
  - DATE (e.g.: '2007-09-30')
  - TIME (e.g.: '15:30:02.5')
  - Any value can be NULL

# Creating Relations in SQL

- Creates the Student relation. Observe that the type (**domain**) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.
  
- As another example, the Enrolled table holds information about courses that students take.

```
CREATE TABLE Student  
  (sid    CHAR(20),  
   name  CHAR(20),  
   login CHAR(10),  
   age   INTEGER,  
   gpa   REAL)
```

```
CREATE TABLE Enrolled  
  (sid    CHAR(20),  
   cid    CHAR(20),  
   grade CHAR(2))
```

# Adding and Deleting Tuples

- Can **insert** a single tuple using:

```
INSERT INTO Student (sid, name, login, age, gpa)
VALUES (300, 'Smith', 'smith@ee', 18, 3.4)
```

sid	name	login	age	gpa
300	Smith	smith@ee	18	3.4

- \* Powerful variants of these commands are available.

# Adding and Deleting Tuples

- Can **delete** all tuples

```
DELETE  
FROM Student S
```

<b>sid</b>	<b>name</b>	<b>login</b>	<b>age</b>	<b>gpa</b>
300	Smith	smith@ee	18	3.4
100	John	john@ceng	19	1.7
200	Fatma	fatma@ee	18	3.6

- \* Powerful variants of these commands are available.

# Adding and Deleting Tuples

- Can **delete** all tuples satisfying some condition (e.g., name = John):

```
DELETE  
FROM Student S  
WHERE S.name = 'John'
```



sid	name	login	age	gpa
300	Smith	smith@ee	18	3.4
200	Fatma	fatma@ee	19	3.5
200	Fatma	fatma@ee	18	3.6

\* Powerful variants of these commands are available.

# Modifying Tuples

- Can **modify** the column values in an existing row using:

```
UPDATE Student S  
SET      S.age = S.age + 1
```

sid	name	login	age	gpa
300	Smith	smith@ee	19	3.2
100	John	john@ceng	20	1.7
200	Fatma	fatma@ee	19	3.6

# Modifying Tuples

- Can **modify** the column values in an existing row using:

```
UPDATE Student S  
SET      S.age = S.age + 1, S.gpa = S.gpa -1  
WHERE   S.sid = 200
```



sid	name	login	age	gpa
300	Smith	smith@ee	18	3.4
100	John	john@ceng	19	1.7
200	Fatma	fatma@ee	19	2.6

# Modifying Tuples

- Can **modify** the column values in an existing row using:

```
UPDATE Student S  
SET      S.gpa = S.gpa - 0.1  
WHERE   S.gpa >= 3.3
```



sid	name	login	age	gpa
300	Smith	smith@ee	18	3.3
100	John	john@ceng	19	1.7
200	Fatma	fatma@ee	18	3.5

**Remark:** In the examples of DELETE and UPDATE, the WHERE clause is applied *first* to determine the rows to be deleted or updated.

# Integrity Constraints (ICs)

- IC: a condition specified on a DB schema & restricts the data that can be stored in an instance
- condition that must be true for *any* instance of the database; e.g., *domain constraints.*
  - ICs are **specified** when schema is **defined**.
  - ICs are **checked** when relations are **modified**.
- A *legal* instance of a relation is one that satisfies all specified ICs.
  - DBMS should not allow illegal instances.
- If the DBMS checks ICs, stored data is more faithful to real-world meaning.
  - Avoids data entry errors, too!

# Key Constraints

- A set of fields is a (**candidate**) **key** for a relation if :
  1. **uniqueness**: Two distinct tuples cannot have identical values in all the fields of a key, and
  2. **minimality**: No subset of the set of fields in a key is a unique identifier for a tuple
    - If part 2 false, then it is a **superkey**.
    - If there is more than one (candidate) key for a relation, one of the keys is chosen (by DBA) to be the **primary key**.

The DBMS optimizes several tasks for the use of PK (eg., while referring to a tuple from another relation or while accessing the table via an index, etc.)

# Key Constraints: Example

- Student(**sid**: string, **name**: string, **login**: string, **age**: integer, **gpa**:real)
- Based on the application requirements of a university DB:
  - What can be the candidate keys for Student? **sid** or: **login**
  - What can be the superkeys? **sid, name**
  - Which one would you choose as PK?  
**sid**  
**sid, gpa**  
**sid, name, gpa**  
**sid, name, gpa, login, age**

*Remark:* ICs are based upon the semantics of the real-world enterprise that is being described in the database relations.

# Primary and Candidate Keys in SQL

- Possibly many *candidate keys*, one of which is chosen as the *primary key*.

```
CREATE TABLE Student
```

```
  (sid    CHAR(20),  
   name   CHAR(20),  
   login  CHAR(10),  
   age    INTEGER,  
   gpa    REAL,
```

```
  PRIMARY KEY (sid),  
  UNIQUE (login))
```

- For **domain, PK and unique constraints**: If a DB modification operation (add/del/update) violates them , operation is **rejected**

# Primary and Candidate Keys in SQL

- “For a given student and course, there is a single grade.”

sid	cid	grade
53666	Ceng351	A
53666	Ceng351	B-
53666	Ceng331	A

```
CREATE TABLE Enrolled  
(sid CHAR(20)  
cid CHAR(20),  
grade CHAR(2),  
PRIMARY KEY (sid,cid))
```



Good design if *Enrolled* is intended to keep the current semester's data!

An IC can prevent storage of instances that arise in practice! So, DBA must set ICs carefully depending on the nature of the application domain!

# Primary and Candidate Keys in SQL

```
CREATE TABLE Enrolled  
  (sid CHAR(20)  
   cid CHAR(20),  
   grade CHAR(2),  
   PRIMARY KEY (sid),  
   UNIQUE (cid, grade) )
```

- Implication: “Students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade.”

“sid	cid	grade
53666	Ceng351	A
53666	Ceng331	B-
53444	Ceng351	A

- An IC can prevent storage of instances that arise in practice!
- X Bad design (or your university is really weird!!!)**

# PRIMARY KEY vs. UNIQUE

1. There can be only one PRIMARY KEY for a relation, but several UNIQUE attributes.
2. No attribute of a PRIMARY KEY can ever be NULL in any tuple. But attributes declared UNIQUE may have NULL's, and there may be several tuples with NULL.

For unknown/missing/inapplicable values

# Foreign Keys, Referential Integrity

- *Foreign key* : Set of fields in one relation that is used to ‘refer’ to a tuple in another relation.  
(Must correspond to primary key of the second relation.) Like a ‘logical pointer’.
- e.g. *sid* is a foreign key referring to *Student*:
  - Enrolled(*sid*: string, *cid*: string, *grade*: string)
  - If all foreign key constraints are enforced,  
*referential integrity* is achieved, i.e., no dangling references.

# Foreign keys...

Enrolled(*sid*: string, *cid*: string, ...)

Student(*sid*:string, ...)

Primary key

Course(*cid*:string, ...)

Primary key



# Foreign Keys in SQL

- Only students listed in the Student relation should be allowed to enroll for courses.

```
CREATE TABLE Enrolled
  (sid CHAR(20), cid CHAR(20), grade CHAR(2),
   PRIMARY KEY (sid,cid),
   FOREIGN KEY (sid) REFERENCES Student )
```

Enrolled

<i>sid</i>	<i>cid</i>	<i>grade</i>
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Student


<i>sid</i>	<i>name</i>	<i>login</i>	<i>age</i>	<i>gpa</i>
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

# Enforcing Referential Integrity

- Consider *Student* and *Enrolled*; *sid* in *Enrolled* is a foreign key that references *Student*.
- What should be done if an *Enrolled* tuple with a non-existent *student id* is **inserted**? (*Reject it!*)
- What should be done if a *Student* tuple is **deleted**?
  - Disallow deletion of a *Student* tuple that is referred to.
  - Set *sid* in *Enrolled* tuples that refer to it to a *default sid*.
  - (In SQL, also: Set *sid* in *Enrolled* tuples that refer to it to a special value *null*, denoting ‘*unknown*’ or ‘*inapplicable*’.)
  - Also delete all *Enrolled* tuples that refer to it.
- Similar if primary key of *Student* tuple is **updated**.


# Handling Foreign Key Violations

- Deletion from *Student*:
  - NO ACTION: Reject if row(s) in *Enrolled* references row to be deleted (default response)




## Handling Foreign Key Violations (cont'd)

- Deletion from *Student* (cont'd):
  - **SET NULL**: Set value of foreign key in referencing row(s) in A to **null**



## Handling Foreign Key Violations (cont'd)

- Deletion from *Student* (cont'd):
  - **SET DEFAULT**: Set value of foreign key in referencing row(s) in *Enrolled* to default value (y) which must exist in *Student*



## Handling Foreign Key Violations (cont'd)


- Deletion from *Student* (cont'd):
  - **SET DEFAULT**: Set value of foreign key in referencing row(s) in *Enrolled* to default value (y) which must exist in *Student*

Enrolled

	y

Student


y	



Default  
value

# Handling Foreign Key Violations (cont'd)

- Deletion from *Student* (cont'd):
  - **CASCADE**: Delete referencing row(s) in *Enrolled* as well



# Referential Integrity in SQL/92

- SQL/92 supports all 4 options on deletes and updates.
  - Default is NO ACTION (*delete/update is rejected*)
  - SET NULL / SET DEFAULT (sets foreign key value of referencing tuple)
  - CASCADE (also delete all tuples that refer to deleted tuple)

# Which of these make sense for Enrolled?

- SQL/92 supports all 4 options on deletes and updates:

NO ACTION / SET NULL / SET DEFAULT / CASCADE

CREATE TABLE *Enrolled*

(*sid* CHAR(20) **DEFAULT** '9999',

*cid* CHAR(20),

*grade* CHAR(2),

**PRIMARY KEY** (*sid,cid*),

**FOREIGN KEY** (*sid*) REFERENCES *Student* **ON DELETE** [REMOVED]  
[REMOVED] **ON UPDATE** [REMOVED]

**FOREIGN KEY** (*cid*) REFERENCES *Course* **ON DELETE** NO ACTION  
**ON UPDATE** CASCADE )

# Foreign Keys: Remarks

- NULL can appear in a FK (unless FK is also the part of the PK in the child table, of course)
- Assume in Student relation, we also keep student's advisor id, which refers to Instructor relation (with PK field **iid**):

Student

sid	name	login	age	gpa	iid
53666	Jones	jones@cs	18	3.4	1
53688	Smith	Smith@eecs	18	3.2	Null
53650	Smith	Smith@math	19	3.8	1

Instructor

iid	name	salary
1	Altingovde	
2	Karagoz	
3	Chair	

**CREATE TABLE** *Student*

(*sid* CHAR(20),  
*name* CHAR(20),  
*login* CHAR(10),  
*age* INTEGER,  
*gpa* REAL,  
*iid* CHAR(20) DEFAULT '3',

**PRIMARY KEY** (*sid*),

**FOREIGN KEY** (*iid*) **REFERENCES** *Instructor* **ON DELETE** [REDACTED]

**ON UPDATE** [REDACTED])

## Student

<b>sid</b>	<b>name</b>	<b>login</b>	<b>age</b>	<b>gpa</b>	<b>iid</b>
53666	Jones	jones@cs	18	3.4	1
53688	Smith	Smith@eecs	18	3.2	Null
53650	Smith	Smith@math	19	3.8	1

<b>iid</b>	<b>name</b>	<b>salary</b>
1	Altingovde	
2	Karagoz	
3	Chair	

# Foreign Keys: Remarks

- NULL can appear in a FK (unless FK is also the part of the PK in the child table, of course)
- The field names in the FK may be different than that of the referred PK, but the number of fields should be same and field types should be compatible

CREATE TABLE *Enrolled*

(*student-id* CHAR(20) DEFAULT '9999',

*cid* CHAR(20),

*grade* CHAR(2),

PRIMARY KEY (*student-id,cid*),

FOREIGN KEY (*student-id*) REFERENCES *Student(sid)*

ON DELETE NO ACTION

CENG 351 File Structures and Data Management

39

ON UPDATE CASCADE

# Foreign Keys: Remarks

- A FK can refer to the same relation
  - In Student relation, there could be the partner field that is intended to store the sid of another student...
  - More examples will be discussed in the ER week...

# Where do ICs Come From?

- ICs are based upon the semantics of the real-world enterprise that is being described in the database relations.
- We can check a database instance to see if an IC is violated, but we can not infer that an IC is true by looking at an instance.
  - An IC is a statement about *all possible* instances!
  - For example, we know *name* is not a key, but the assertion that *sid* is a key is given to us.
- Key and foreign key ICs are the most common; more general ICs supported too.