

SQL: Structured Query Language

Part 2

Chapter 5

Nested Queries: IN

IN: Allows us to test whether a value is in a given set of elements
(usually generated by another SQL query)

Find names of sailors who've ids 1 or 2 or 3 or 4 or 5

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid IN (1, 2, 3, 4, 5)
```


sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
9	Mike	4	57
4	Mary	1	17
22	Jake	10	57
3	Nancy	8	27

Nested Queries: IN

IN: Allows us to test whether a value is in a given set of elements
(usually generated by another SQL query)

Find names of sailors who've ids 1 or 2 or 3 or 4 or 5

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid IN (1, 2, 3, 4, 5)
```



sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
9	Mike	4	57
4	Mary	1	17
22	Jake	10	57
3	Nancy	8	27

sname
Fred
Jim
Mary
Nancy

Nested Queries: IN

Find names of sailors who've reserved boat #102:

```
SELECT S.sname  
FROM   Sailors S  
WHERE  S.sid IN  
       (SELECT R.sid  
        FROM   Reserves R  
        WHERE  R.bid=102)
```

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20
3	104	11/20

Nested Queries: IN

Find names of sailors who've reserved boat #102:

```
SELECT S.sname  
FROM   Sailors S  
WHERE  S.sid IN  
       (SELECT R.sid  
        FROM   Reserves R  
        WHERE  R.bid=102)
```

sid
1
2

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20 5
3	104	11/20

Find sid's of sailors who've reserved a red **and** a green boat

Reserves

sid	bid	day
1	103	12/9/2015
2	102	13/9/2015
2	103	1/1/2020
3	101	1/1/2020
3	102	5/1/2020

Boats

bid	bname	color
101	Nina	red
102	Pinta	green
103	Santa Maria	blue

Find sid's of sailors who've reserved a red **and** a green boat

(sid of Sailors who reserve red boat
AND sid IN

(sid of Sailors who
reserve green boat)

Reserves

sid	bid	day
1	103	12/9/2015
2	102	13/9/2015
2	103	1/1/2020
3	101	1/1/2020
3	102	5/1/2020

Boats

bid	bname	color
101	Nina	red
102	Pinta	green
103	Santa Maria	blue

Find sid's of sailors who've reserved a red **and** a green boat

(sid of Sailors who reserve red boat
AND sid **IN**

```
(SELECT R2.sid
  FROM Boats B2,Reserves R2
 WHERE R2.bid=B2.bid
   AND B2.color='green')
```

sid
2
3

Reserves

sid	bid	day
1	103	12/9/2015
2	102	13/9/2015
2	103	1/1/2020
3	101	1/1/2020
3	102	5/1/2020

Boats

bid	bname	color
101	Nina	red
102	Pinta	green
103	Santa Maria	blue



Find sid's of sailors who've reserved a red **and** a green boat


```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
      AND B.color='red'  
      AND R.sid IN
```

sid
2
3

```
(SELECT R2.sid  
FROM Boats B2,Reserves R2  
WHERE R2.bid=B2.bid  
      AND B2.color='green')
```

Reserves

sid	bid	day
1	103	12/9/2015
2	102	13/9/2015
2	103	1/1/2020
3	101	1/1/2020
3	102	5/1/2020



Boats

bid	bname	color
101	Nina	red
102	Pinta	green
103	Santa Maria	blue ₉

Find sid's of sailors who've reserved a red **and** a green boat

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
      AND B.color='red'  
      AND R.sid IN
```

sid
2
3

(~~SELECT R2.sid
FROM Boats B2,Reserves R2
WHERE R2.bid=B2.bid
 AND B2.color='green'~~)

Reserves

sid	bid	day
1	103	12/9/2015
2	102	13/9/2015
2	103	1/1/2020
3	101	1/1/2020
3	102	5/1/2020

Boats

bid	bname	color
101	Nina	red
102	Pinta	green
103	Santa Maria	blue

Find sid's of sailors who've reserved a red **and a green boat**

→ INTERSECT can be re-written using IN

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid
      AND B.color='red'
      AND R.sid IN (SELECT R2.sid
                      FROM Boats B2,Reserves R2
                      WHERE R2.bid=B2.bid
                          AND B2.color='green')
```

Nested Queries: NOT IN

Find names of sailors who've not reserved boat #102:

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20 ¹²
3	104	11/20

Nested Queries: NOT IN

Find names of sailors who've not reserved boat #102:

```
SELECT S.sname  
FROM Sailors S  
WHERE S.sid NOT IN  
(SELECT R.sid  
FROM Reserves R  
WHERE R.bid=102)
```

sid
1
2

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27



sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20
3	104	11/20

Nested Queries: NOT IN

Find names of sailors who've not reserved boat #102:

```
SELECT S.sname
FROM Sailors S
WHERE S.sid NOT IN
  (SELECT R.sid
   FROM Reserves R
   WHERE R.bid=102)
```

→ EXCEPT can be re-written using NOT IN

Nested Queries with Correlation

EXISTS: Allows us to test whether a set is NON-EMPTY

Find names of sailors who've reserved boat #102:

```
SELECT S.sname
FROM   Sailors S
WHERE EXISTS
       (SELECT *
        FROM  Reserves R
        WHERE R.bid=102 AND S.sid=R.sid)
```

- Subquery must be recomputed for each Sailors tuple.
 - Think of subquery as a function call that runs a query

Nested Queries with Correlation

EXISTS: Allows us to test whether a set is NON-EMPTY

Find names of sailors who've reserved boat #102:

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20
3	104	11/20 16

Nested Queries with Correlation

EXISTS: Allows us to test whether a set is NON-EMPTY

Find names of sailors who've reserved boat #102:

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS
  (SELECT *
   FROM Reserves R
   WHERE R.bid=102 AND S.sid=R.sid)
```

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27


sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20
3	104	11/20 17

Nested Queries with Correlation

EXISTS: Allows us to test whether a set is NON-EMPTY

Find names of sailors who've reserved boat #102:

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS
  (SELECT *
   FROM Reserves R
   WHERE R.bid=102 AND S.sid=R.sid)
```



sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27


sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20
3	104	11/20 18

Nested Queries with Correlation

EXISTS: Allows us to test whether a set is NON-EMPTY

Find names of sailors who've reserved boat #102:

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS
  (SELECT *
   FROM Reserves R
   WHERE R.bid=102 AND S.sid=R.sid)
```



sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20
3	104	11/20 19


Nested Queries with Correlation

Find sailors with at most one reservation for boat #102

Nested Queries with Correlation

Find sailors with at most one reservation for boat #102

```
SELECT S.sname
FROM   Sailors S
WHERE  UNIQUE
       (SELECT R.bid
        FROM Reserves R
        WHERE R.bid=102 AND S.sid=R.sid)
```



Nested Queries with Correlation

Find sailors with at most one reservation for boat #102

```
SELECT S.sname
FROM   Sailors S
WHERE  UNIQUE
       (SELECT R.bid
        FROM Reserves R
        WHERE R.bid=102 AND S.sid=R.sid)
```


UNIQUE checks for duplicate tuples. When applied to a subquery, it is **TRUE**:

- if no row appears twice,
- if the answer is empty set

Nested Queries with Correlation

Finds sailors with at most one reservation for boat #102

```
SELECT S.sname
FROM Sailors S
WHERE UNIQUE
  (SELECT R.bid
   FROM Reserves R
   WHERE R.bid=102 AND S.sid=R.sid)
```




sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20
3	104	11/20
2	102	11/20

Nested Queries with Correlation

Finds sailors with at most one reservation for boat #102

```
SELECT S.sname  
FROM Sailors S  
WHERE UNIQUE  
(SELECT R.bid  
  FROM Reserves R  
 WHERE R.bid=102 AND S.sid=R.sid)
```



The diagram shows three green arrows pointing to the first three rows of the 'Sailors' table. The table has columns: sid, sname, rating, and age. The rows are: (1, Fred, 7, 22), (2, Jim, 2, 39), and (3, Nancy, 8, 27).

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20
3	104	11/20
2	102	11/20


Nested Queries with Correlation

Finds sailors with at most one reservation for boat #102

```
SELECT S.sname  
FROM Sailors S  
WHERE UNIQUE
```

```
(SELECT R.bid  
FROM Reserves R  
WHERE R.bid=102 AND S.sid=R.sid)
```

sname
Fred
Nancy



sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

sid	bid	day
1	102	9/12
2	102	9/13
2	101	10/20
3	104	11/20
2	102	11/20

In the subquery, what happens if we write * instead of *R.bid*?

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- Also available: $op\text{ ANY}$, $op\text{ ALL}$ $>, <, =, \geq, \leq, \neq$

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- Also available: *op ANY, op ALL >,<,=,≥,≤,≠*

Find sailors whose rating is greater than some sailor called Jim:

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Jim	1	17
22	Fred	7	50
3	Nancy	1	21

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- Also available: *op ANY*, *op ALL* $>$, $<$, $=$, \geq , \leq , \neq

Find sailors whose rating is greater than some sailor called Jim:

```
SELECT *
FROM   Sailors S
WHERE  S.rating > ANY
       (SELECT s2.rating
        FROM   Sailors s2
        WHERE  s2.sname='Jim')
```

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Jim	1	17
22	Fred	7	50
3	Nancy	1	21

If the subquery returns an empty set, comparison returns FALSE²⁸

More on Set-Comparison Operators

Find sailors whose rating is greater than every sailor called Jim.

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Jim	1	17
22	Fred	7	50
3	Nancy	1	21

More on Set-Comparison Operators

Find sailors whose rating is greater than every sailor called Jim.

```
SELECT *
FROM   Sailors S
WHERE  S.rating > ALL
          (SELECT s2.rating
           FROM   Sailors s2
           WHERE  s2.sname='Jim')
```

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Jim	1	17
22	Fred	7	50
3	Nancy	1	21

If the subquery returns an **empty set**, comparison returns **TRUE**

More on Set-Comparison Operators

Find sailors with highest rating.

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Jim	1	17
22	Fred	7	50
3	Nancy	1	21

More on Set-Comparison Operators

Find sailors with highest rating.

```
SELECT *
FROM Sailors S
WHERE S.rating >= ALL (SELECT S2.rating
                        FROM Sailors S2)
```

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Jim	1	17
22	Fred	7	50
3	Nancy	1	21

Note: **IN** equivalent to **= ANY**
NOT IN equivalent to **<> ALL**

Division in SQL

Find sailors who've reserved all boats.

$$\rho (Tempsids, (\pi_{sid,bid} \text{Reserves}) / (\pi_{bid} \text{Boats}))$$

$$\pi_{sname} (Tempsids \bowtie Sailors)$$

Division in SQL

Find sailors who've reserved all boats.

$$\rho (Tempsids, (\pi_{sid,bid} \text{Reserves}) / (\pi_{bid} \text{Boats}))$$
$$\pi_{sname} (Tempsids \bowtie Sailors)$$

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS ( SELECT B.bid
                    FROM Boats B
                    WHERE NOT EXISTS ( SELECT R.bid
                                         FROM Reserves R
                                         WHERE R.bid=B.bid
                                              AND R.sid=S.sid ) )
```

Division in SQL

Find sailors who've reserved all boats.

$$\rho \left(Tempsids, (\pi_{sid,bid} Reserves) / (\pi_{bid} Boats) \right)$$

$$\pi_{sname} \quad (Tempsids \quad \bowtie \quad Sailors \quad)$$

```

SELECT S.sname      Sailors S such that ...
FROM Sailors S
WHERE NOT EXISTS ( SELECT B.bid
                    FROM Boats B   there is no boat B
                    WHERE NOT EXISTS ( SELECT R.bid
                                         FROM Reserves R
                                         WHERE R.bid=B.bid
                                         AND R.sid=S.sid ) )
without a Reserves tuple showing S reserved B 35  

(i.e., there is no reservation tuple)

```

Division in SQL

Find sailors who've reserved all boats.

```
SELECT S.sname  
FROM Sailors S  
WHERE NOT EXISTS ( SELECT B.bid  
                   FROM Boats B  
                   WHERE NOT EXISTS ( SELECT R.bid  
                           FROM Reserves R  
                           WHERE R.bid=B.bid  
                                 AND R.sid=S.sid ) )
```

What is the green subquery finding?

Division in SQL

Find sailors who've reserved all boats.

```
SELECT S.sname  
FROM Sailors S  
WHERE NOT EXISTS ( SELECT B.bid  
                   FROM Boats B  
                   WHERE NOT EXISTS ( SELECT R.bid  
                           FROM Reserves R  
                           WHERE R.bid=B.bid  
                                 AND R.sid=S.sid ) )
```

What is the green subquery finding?

Ans: The boats that are **not reserved** by the given sailor!

sid	name
1	fred
2	wilma

bid	color
101	red
102	green

sid	bid
1	101
1	102
2	102

```

SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS ( SELECT B.bid
                    FROM Boats B
                    WHERE NOT EXISTS ( SELECT R.bid
                                        FROM Reserves R
                                        WHERE R.bid=B.bid
                                              AND R.sid=S.sid ) )
    
```


sid	name
1	fred
2	wilma

bid	color
101	red
102	green

sid	bid
1	101
1	102
2	102

Answer:

```
SELECT S.sname  
FROM Sailors S  
WHERE NOT EXISTS ( SELECT B.bid  
Empty → returns True
```



sname

fred

sid	name
1	fred
2	wilma

bid	color
101	red
102	green

sid	bid
1	101
1	102
2	102

SELECT S.sname
 FROM Sailors S
 WHERE NOT EXISTS (SELECT B.bid
 FROM Boats B
 WHERE NOT EXISTS (SELECT R.bid
 FROM Reserves R
 WHERE R.bid=B.bid
 AND R.sid=S.sid))

Non Empty → returns False

bid
101

Answer:

sname
fred

S.sid=2, R.bid=101 → Empty
 S.sid=2, R.bid=102 → Non empty

Division using EXCEPT

Find sailors who've reserved all boats.

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
    ((SELECT B.bid
      FROM Boats B)
     EXCEPT
     (SELECT R.bid
      FROM Reserves R
      WHERE R.sid=S.sid))
```

sid	name
1	fred
2	wilma

bid	color
101	red
102	green

sid	bid
1	101
1	102
2	102

Division using EXCEPT

Find sailors who've reserved all boats.

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
    ((SELECT B.bid
      FROM Boats B)
     EXCEPT
     (SELECT R.bid
      FROM Reserves R
      WHERE R.sid=S.sid))
```

sname
fred

bid
101
102

bid
101
102

sid	name
1	fred
2	wilma

bid	color
101	red
102	green

sid	bid
1	101
1	102
2	102



Division using EXCEPT

Find sailors who've reserved all boats.

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
    ((SELECT B.bid
      FROM Boats B)
     EXCEPT
     (SELECT R.bid
      FROM Reserves R
      WHERE R.sid=S.sid))
```

sid	name
1	fred
2	wilma

bid	color
101	red
102	green

sid	bid
1	101
1	102
2	102

sname
fred

Aggregate Operators

- Significant extension of relational algebra.

```
COUNT (*)
COUNT ( [DISTINCT] A)
SUM ( [DISTINCT] A)
AVG ( [DISTINCT] A)
MAX (A)
MIN (A)
```



single column

Aggregate Operators

```
SELECT COUNT (*)
FROM Sailors S
```

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Mary	1	17
22	Fred	7	50
3	Nancy	2	21

```
SELECT COUNT (DISTINCT S.rating)
FROM Sailors S
WHERE S.sname='Fred'
```

```
SELECT AVG (S.age)
FROM Sailors S
WHERE S.rating=7
```

```
SELECT AVG ( DISTINCT S.age)
FROM Sailors S
WHERE S.rating=7
```

Aggregate Operators

Find the names of the sailors with the highest rating.

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Mary	1	17
22	Fred	7	50
3	Nancy	2	21

Aggregate Operators

Find the names of the sailors with the highest rating.

```
SELECT S.sname  
FROM Sailors S  
WHERE S.rating=  
      (SELECT MAX(S2.rating)  
       FROM Sailors S2)
```

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Mary	1	17
22	Fred	7	50
3	Nancy	2	21

Find name and age of the oldest sailor(s)

Find name and age of the oldest sailor(s)

```
SELECT S.sname, MAX (S.age)  
FROM Sailors S
```

Find name and age of the oldest sailor(s)

~~SELECT S.sname, MAX(S.age)
FROM Sailors S~~



- This query is illegal:
 - If the SELECT clause uses an aggregate operation, then it must use *only* aggregate operations
(unless the query contains **GROUP BY** clause)

Find name and age of the oldest sailor(s)

- Correct way to write the query

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age =  
    (SELECT MAX (S2.age)  
     FROM Sailors S2)
```

ORDER BY

- Orders the results of a query by the specified fields
- If included in a query, it should be the last part
- Can order by multiple fields by separating the fields with commas, the field that comes before others has precedence in ordering

```
SELECT s.rating, s.age
```

```
FROM Sailors s
```

```
ORDER BY s.rating ASC, s.age DESC
```

First orders by rating in ascending order, in case of ties between ratings, orders by age in descending order

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

ORDER BY

- Orders the results of a query by the specified fields
- If included in a query, it should be the last part
- Can order by multiple fields by separating the fields with commas, the field that comes before others has precedence in ordering

SELECT s.rating, s.age

FROM Sailors s

ORDER BY s.rating ASC, s.age DESC

First orders by rating in ascending order, in case of ties between ratings, orders by age in descending order

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0



rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0
10	16.0