

SQL: Structured Query Language

Part 3

GROUP BY and HAVING

- So far, we've applied aggregate operators to all (qualifying) tuples. Sometimes, we want to apply them to each of several *groups* of tuples.
- Consider: *Find the age of the youngest sailor for each rating level.*
 - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
 - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (!):

For $i = 1, 2, \dots, 10$:

```
SELECT MIN (S.age)
FROM Sailors S
WHERE S.rating = i
```

Find the age of the youngest sailor for each rating level with at least 2 sailors

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
GROUP BY S.rating
HAVING COUNT(*) >1
```

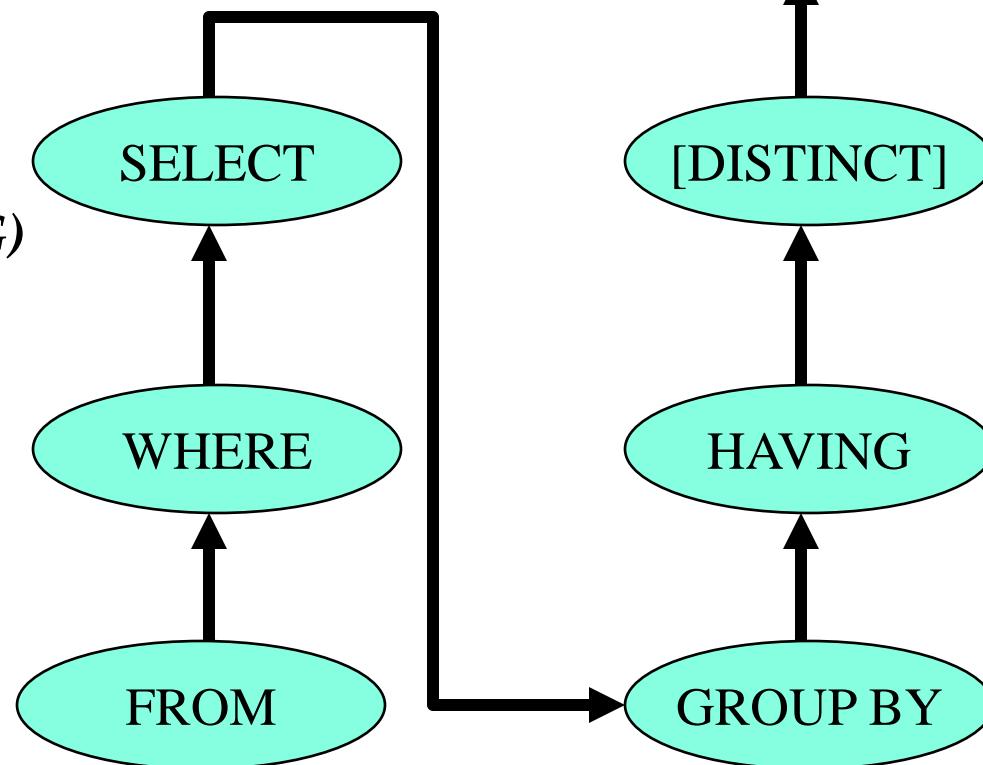
Conceptual Evaluation

```
SELECT      [DISTINCT] target-list  
FROM        relation-list  
WHERE       qualification  
GROUP BY   grouping-list  
HAVING     group-qualification
```

*Project away columns
(just keep those used in
SELECT, GBY, HAVING)*

*Apply selections
(eliminate rows)*

*Relation
cross-product*



*Eliminate
duplicates*

*Eliminate
groups*

*Form groups
& aggregate*

Find the age of the youngest sailor for each rating level with at least 2 sailors

Find the age of the youngest sailor for each rating level with at least 2 sailors

```
SELECT S.rating, MIN (S.age)
```

```
FROM Sailors S
```

```
GROUP BY S.rating
```

```
HAVING COUNT(*) >1
```

Find the age of the youngest sailor for each rating level with at least 2 sailors

```
SELECT S.rating, MIN (S.age)
```

```
FROM Sailors S
```

```
GROUP BY S.rating
```

```
HAVING COUNT(*) >1
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

sid	sname	rating	age
29	brutus	1	33.0
22	dustin	7	45.0
64	horatio	7	35.0
31	lubber	8	55.5
58	rusty	10	35.0
71	zorba	10	16.0

Expressions in *group-qualification* must have a single value per group!

One answer tuple is generated **per** qualifying group.

How about this change on the query?

```
SELECT S.rating, MIN (S.age)  
FROM Sailors S  
GROUP BY S.rating  
HAVING S.rating > 8
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

sid	sname	rating	age
29	brutus	1	33.0
22	dustin	7	45.0
64	horatio	7	35.0
31	lubber	8	55.5
58	rusty	10	35.0
71	zorba	10	16.0

Expressions in *group-qualification* must have a single value per group!

One answer tuple is generated **per** qualifying group.

How about this change on the query?

```
SELECT S.rating, MIN (S.age)  
FROM Sailors S  
GROUP BY S.rating  
HAVING S.age > 30
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

sid	sname	rating	age
29	brutus	1	33.0
22	dustin	7	45.0
64	horatio	7	35.0
31	lubber	8	55.5
58	rusty	10	35.0
71	zorba	10	16.0

Expressions in *group-qualification*
must have a single value per group!

One answer tuple is generated **per**
qualifying group.

How about this change on the query?

```
SELECT S.rating, MIN (S.age)  
FROM Sailors S  
GROUP BY S.rating  
HAVING S.age > 30
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

sid	sname	rating	age
29	brutus	1	33.0
22	dustin	7	45.0
64	horatio	7	35.0
31	lubber	8	55.5
58	rusty	10	35.0
71	zorba	10	16.0

Expressions in *group-qualification* must have a single value per group!

One answer tuple is generated ^{per} **qualifying group.**

Find the age of the youngest sailor for each rating level with at least 2 sailors

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
GROUP BY S.rating
HAVING COUNT(*) >1
```

Equivalently...

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
GROUP BY S.rating
HAVING 1 < (SELECT COUNT (*)
               FROM Sailors S2
              WHERE S.rating=S2.rating)
```

- Shows HAVING clause can also contain a subquery.

Find the age of the youngest sailor for each rating level with at least 2 sailors

```
SELECT S.rating, MIN (S.age)  
FROM Sailors S  
GROUP BY S.rating  
HAVING COUNT(*) >1
```

Equivalently...

```
SELECT S.rating, MIN (S.age)  
FROM Sailors S  
GROUP BY S.rating  
HAVING 1 < (SELECT COUNT(*)  
              FROM Sailors S2  
              WHERE S.rating=S2.rating)
```

sid	sname	rating	age
29	brutus	1	33.0
22	dustin	7	45.0
64	horatio	7	35.0
31	lubber	8	55.5
58	rusty	10	35.0
71	zorba	10	16.0

- Shows HAVING clause can also contain a subquery.

Queries With GROUP BY and HAVING

```
SELECT      [DISTINCT] target-list
FROM        relation-list
WHERE       qualification  $\subseteq .$ 
GROUP BY    grouping-list  $\subseteq .$ 
HAVING     group-qualification
```

- The *target-list* contains (i) attribute names (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
 - The attribute list (i) must be **a subset of *grouping-list***. Intuitively, each answer tuple corresponds to a *group*, and these attributes must have **a single value per group**. (A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.)

Conceptual Evaluation

- The cross-product of *relation-list* is computed, tuples that fail *qualification* are discarded, ‘*unnecessary*’ fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- The *group-qualification* is then applied to eliminate some groups. Expressions in *group-qualification* must have a *single value per group!*
 - In effect, an attribute in *group-qualification* that is not an argument of an aggregate op also appears in *grouping-list*.
- One answer tuple is generated **per qualifying group**.

For each red boat, find the number of reservations for this boat

```
SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

- Grouping over a join of two relations.

For each red boat, find the number of reservations for this boat

Reserves

sid	bid	date
1	101	1/1/2017
1	102	15/1/2017
2	108	3/3/2016
1	101	5/6/2016
1	108	4/4/2017

Boats

bid	color
101	red
102	green
108	red

```
SELECT B.bid, COUNT (*) AS scount
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid
```



sid	R.bid	date	B.bid	color
1	101	...	101	red
1	102		102	green
2	108		108	red
1	101		101	red
1	108		108	red

sid R.bid B.bid color

1	101	...	101	red
1	101	...	101	red
2	108	...	108	red
1	108	...	108	red

For each red boat, find the number of reservations for this boat

```
SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

- Grouping over a join of two relations.
- What do we get if we remove *B.color=‘red’* from the WHERE clause and add a HAVING clause with this condition?

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 sailors (of any age)

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 sailors (of any age)

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT (*)
               FROM Sailors S2
              WHERE S.rating=S2.rating )
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 sailors (of any age)

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
```



sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Rating	age
1	33.0
7	45.5
7	35.0
8	55.5
10	35.0

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 sailors (of any age)

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT (*)
               FROM Sailors S2
              WHERE S.rating=S2.rating )
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Rating	age
1	33.0
7	45.5
7	35.0
8	55.5
10	35.0

- Shows HAVING clause can also contain a subquery.

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 sailors (of any age)

```

SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT(*)
               FROM Sailors S2
               WHERE S.rating=S2.rating )
    
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Rating	age
1	33.0
7	45.5
7	35.0
8	55.5
10	35.0

For rating=10
Count(*)
2

For rating=7
Count(*)
2

rating	
7	35.0
10	35.0

- Shows HAVING clause can also contain a subquery.

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

- Only S.rating and S.age are mentioned in the SELECT, GROUP BY or HAVING clauses; other attributes ‘unnecessary’.

Rating	age
7	45.0
8	55.5
7	35.0
1	33.0
10	35.0

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Rating	age
1	33.0
7	45.5
7	35.0
8	55.5
10	35.0

rating	
7	35.0

Answer relation

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

SELECT Temp.rating, Temp.minage
FROM

WHERE Temp.rcount > 1

Assume a TEMP storing number
& min age of sailors ≥ 18 per rating

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0



rating	rcount	minage
1	1	33.0
7	2	35.0
8	1	55.5
10	1	35.0

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

SELECT Temp.rating, Temp.minage

FROM

(SELECT S.rating, COUNT(*) AS rcount, MIN(S.age) AS minage
FROM Sailors S
WHERE S.age ≥ 18
GROUP BY S.rating) AS Temp

WHERE Temp.rcount > 1

Assume a TEMP storing number
& min age of sailors ≥ 18 per rating

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0



rating	rcount	minage
1	1	33.0
7	2	35.0
8	1	55.5
10	1	35.0

RECALL: Find name and age of the oldest sailor(s)

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age =  
    (SELECT MAX (S2.age)  
     FROM Sailors S2)
```

Find those ratings for which the average age is the minimum over all ratings

Find those ratings for which the average age is the minimum over all ratings

```
SELECT S.rating  
FROM Sailors S  
WHERE S.age = (SELECT MIN (AVG (S2.age)) FROM Sailors S2)
```

???

Find those ratings for which the average age is the minimum over all ratings

```
SELECT S.rating  
FROM Sailors S  
WHERE S.age = (SELECT MIN (AVG (S2.age)) FROM Sailors S2)
```



- Aggregate operations **cannot be nested!** WRONG:

Find those ratings for which the average age is the minimum over all ratings

- Correct solution (in SQL/92):

```
SELECT Temp.rating, Temp.avgage  
FROM
```

```
WHERE Temp.avgage = (SELECT MIN(Temp.avgage)  
                      FROM Temp)
```

OK to define once
on pen&paper

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

rating	avgage
1	33.0
7	40.0
8	55.5
10	25.5

Find those ratings for which the average age is the minimum over all ratings

- Correct solution (in SQL/92):

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG (S.age) AS avgage
      FROM Sailors S
     GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN(Temp.avgage)
                      FROM Temp)
```

OK to define once
on pen&paper

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

rating	avgage
1	33.0
7	40.0
8	55.5
10	25.5

Find those ratings for which the average age is the minimum over all ratings

- How about....?

```
SELECT Temp.rating, MIN(Temp.avgage)
FROM (SELECT S.rating, AVG(S.age) AS avgage
      FROM Sailors S
     GROUP BY S.rating) AS Temp
GROUP BY Temp.rating
```

Temp	rating	avgage
	1	33.0
	7	40.0
	8	55.5
	10	25.5

Find those ratings for which the average age is the minimum over all ratings

- How about....?

```
SELECT Temp.rating, MIN(Temp.avgage)
FROM (SELECT S.rating, AVG(S.age) AS avgage
      FROM Sailors S
     GROUP BY S.rating) AS Temp
GROUP BY Temp.rating
```

Temp	rating	avgage
	1	33.0
	7	40.0
	8	55.5
	10	25.5

If the SELECT clause uses an aggregate operation, then it must use *only* aggregate operations unless the query contains **GROUP BY** clause

```
SELECT S.sname, MAX (S.age)
FROM Sailors S
```

Find those ratings for which the average age is the minimum over all ratings

```
SELECT Temp.rating, Temp.avgage  
FROM (SELECT S.rating, AVG (S.age) AS avgage      What if you  
      FROM Sailors S  
      GROUP BY S.rating) AS Temp  
WHERE Temp.avgage = (SELECT MIN(Temp.avgage)  
                      FROM Temp)
```

Find those ratings for which the average age is the minimum over all ratings

```
SELECT Temp.rating, Temp.avgage  
FROM (SELECT S.rating, AVG(S.age) AS avgage  
      FROM Sailors S  
     GROUP BY S.rating) AS Temp  
WHERE Temp.avgage = (SELECT MIN(Temp.avgage)  
                      FROM Temp)
```

rating	avgage
1	33.0
7	40.0
8	55.5
10	25.5

```
SELECT      S.rating, AVG(S.age) AS avgage  
FROM        Sailors S  
GROUP BY   S.rating  
HAVING     AVG(S.age) <= ALL  
              (SELECT AVG(S2.age)  
               FROM Sailors S2  
              GROUP BY S2.rating)
```

Find those ratings for which the average age is the minimum over all ratings

```
SELECT Temp.rating, Temp.avgage  
FROM (SELECT S.rating, AVG (S.age) AS avgage  
      FROM Sailors S  
     GROUP BY S.rating) AS Temp  
WHERE Temp.avgage = (SELECT MIN (Temp.avgage)  
                      FROM Temp)
```

rating	avgage
1	33.0
7	40.0
8	55.5
10	25.5

```
SELECT      S.rating, AVG (S.age) AS avgage  
FROM        Sailors S  
GROUP BY    S.rating  
HAVING     AVG (S.age) <= ALL  
              (SELECT  AVG(S2.age)          33.0  
               FROM    Sailors S2         40.0  
              GROUP BY S2.rating)      55.5  
                                         25.5
```

Null Values

- Field values in a tuple are sometimes *unknown* (e.g., a rating has not been assigned) or *inapplicable* (e.g., no spouse's name).
 - SQL provides a special value *null* for such situations.
- The presence of *null* complicates many issues.
 - E.g.: comparison

→ Special operators “IS NULL” and “IS NOT NULL”.

sid	sname	rating	age
21	Dan	NULL	38

Null Values: 3 valued logic

- Is $rating > 8$ true or false when $rating$ is equal to *null*? What about AND, OR and NOT connectives?
- We need a 3-valued logic (true, false and *unknown*).

sid	sname	rating	age
21	Dan	NULL	38

- $rating > 8$ can be **true**, **false** or **unknown**
- $rating > 8$ OR $age < 40$???
- $rating > 8$ OR $age > 40$???
- $rating > 8$ AND $age > 40$???

- Meaning of constructs must be defined carefully. (e.g., WHERE clause **eliminates rows** that **don't evaluate to true**.)

Null Values: 3 valued logic

- Is $rating > 8$ true or false when $rating$ is equal to *null*? What about AND, OR and NOT connectives?
- We need a 3-valued logic (true, false and *unknown*).

sid	sname	rating	age
21	Dan	NULL	38

- $rating > 8$ can be true, false or unknown
- $rating > 8$ OR $age < 40$???
- $rating > 8$ OR $age > 40$???
- $rating > 8$ AND $age > 40$???

OR TABLE		
T	T	T
T	F	T
F	T	T
T	Unknown	T
Unknown	T	T
F	F	F
F	Unknown	Unknown
Unknown	F	Unknown
Unknown	Unknown	Unknown

- Meaning of constructs must be defined carefully. (e.g., WHERE clause **eliminates rows** that **don't evaluate to true**.)

Null Values: Aggregate Operators

- COUNT(*) → counts also the null values
- All others (as below): **discards** NULL values

COUNT ([DISTINCT] A)

SUM ([DISTINCT] A)

AVG ([DISTINCT] A)

MAX (A)

MIN (A)

Outer Joins

- Left outer join for S and R:
 - Each S row without a matching R rows appears (once) in the result, with R columns including NULLs
 - `SELECT S.sid, R.bid`
`FROM Sailors S NATURAL LEFT OUTER JOIN Reserves R`

Sailors

sid	sname	rating	age
22	Dustin	7	45.0
31	lubber	8	55.5
71	Zorba	10	16.0

Reserves



sid	bid	date
22	101	...
31	102	
71	null	

Modifying Tables – Insert

- Inserting a single row into a table
 - Attribute list can be omitted if it is the same as in CREATE TABLE
 - NULL and DEFAULT values can be specified

```
INSERT INTO Transcript(StudId, CrsCode, Semester, Grade)
VALUES (12345, 'CSE305', 'S2000', NULL)
```

Bulk Insertion

- Insert the rows output by a SELECT

```
CREATE TABLE DeansList (  
    StudId      INTEGER,  
    Credits     INTEGER,  
    CumGpa      REAL,  
    PRIMARY KEY  (StudId)
```

```
INSERT INTO DeansList (StudId, Credits, CumGpa)  
SELECT      T.StudId, 3 * COUNT (*), AVG(T.Grade)  
FROM        Transcript T  
GROUP BY    T.StudId  
HAVING     AVG (T.Grade) > 3.5 AND  
                COUNT(*) > 30
```

Recall: Transcript(*StudId*, *CrsCode*, *Semester*, *Grade*)

Modifying Tables – Delete

- Similar to SELECT except:
 - No project list in DELETE clause
 - **No Cartesian product** in FROM clause (only 1 table name)
 - Rows satisfying WHERE clause (general form, including subqueries, allowed) are deleted instead of output

```
DELETE FROM Transcript T  
WHERE T.Grade IS NULL AND T.Semester <> 'F2020'
```

Modifying Data - Update

```
UPDATE Employee E  
SET      E.Salary = E.Salary * 1.05  
WHERE   E.Department = 'R&D'
```

- Updates rows in a **single table**
- All rows satisfying WHERE clause (general form, including subqueries, allowed) are updated

Test first!

UPDATE Employee E

SET E.Salary = E.Salary * 5

WHERE E.Salary <= (SELECT AVG(E2.salary)
FROM Employee E2)



eid	ename	salary	age
1	Fred	1000	20
2	Jim	2000	39
9	Mike	500	20
4	Mary	3000	17
22	Fred	500	50
3	Nancy	1400	21

$$\text{Avg: } 8400/6 = 1400$$

AND THAT IS HOW YOU WRITE A SQL QUERY

ANY QUESTIONS?