

Relational Algebra

Relational Query Languages

- Query languages: Allow manipulation and retrieval of data from a database.
- Relational model supports simple, powerful query languages:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- Query Languages != programming languages!
 - QLs not expected to be “Turing complete”.
 - QLs not intended to be used for complex calculations.
 - QLs support easy, efficient access to large data sets.

Formal Relational Query Languages

Two mathematical Query Languages form the basis for “real” languages (e.g. SQL), and for implementation:

- Relational Algebra: More operational, very useful for representing execution plans.
- Relational Calculus: Lets users describe what they want, rather than how to compute it. (Non-operational, declarative.)
 - *Understanding Algebra & Calculus is key to*
 - *understanding SQL, query processing!*

Preliminaries

- A query is applied to *relation instances*, and the result of a query is also a relation instance.
 - *Schemas* of input relations for a query are fixed
 - The schema for the *result* of a given query is also fixed!
Determined by definition of query language constructs.
- Positional vs. named-field notation:
 - Positional notation easier for formal definitions, named-field notation more readable.
 - Both used in SQL

Example Instances

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

- “Sailors” and “Reserves” relations for our examples.
- We’ll use positional or named field notation, assume that names of fields in query results are ‘inherited’ from names of fields in query input relations.

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Relational Algebra

- Basic operations:
 - Projection (π) Deletes unwanted columns from relation.
 - Selection (σ) Selects a subset of rows from relation.
 - Cross-product (\times) Allows us to combine two relations.
 - Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
 - Union (\cup) Tuples in reln. 1 and in reln. 2.
 - Renaming (ρ)
- Additional operations:
 - Intersection (\cap), join (\bowtie), division ($/$), $:$ Can be represented over basic operations, but (very!) useful.
- Since each operation returns a relation, operations can be *composed!* (Algebra is “closed”.)

Projection (π)

- Deletes attributes that are not in *projection list*.
- *Schema* of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

$\pi_{sname, rating}(S2)$

Projection (π)

- Projection operator *eliminates duplicates*! (Why??)
 - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it. (Why not?)

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

age
35.0
55.5

$\pi_{age}(S2)$

Selection (σ)

- Selects rows that satisfy *selection condition*.
- No duplicates in result! (Why?)
- *Schema* of result identical to schema of (only) input relation.

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

$$\sigma_{rating > 8}(S2)$$

Selection (σ)

- *Result* relation can be the *input* for another relational algebra operation! (*Operator composition.*)

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sname	rating
yuppy	9
rusty	10

$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$

Selection (σ)

Selection Condition:

- Comparisons in the syntax
attribute *op* constant, attribute *op* attribute
 $op \in \{=, \neq, <, \leq, >, \geq\}$
- Can be combined with Boolean operators, \vee , \wedge

Union, Intersection, Set-Difference

- All of these operations (\cup , \cap , $-$) take two input relations, which must be union-compatible:
 - Same number of fields.
 - ‘Corresponding’ fields have the same type.
- What is the *schema* of result?

$$S1 \cup S2$$

$$S1 \cap S2$$

$$S1 - S2$$

Union

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$S1 \cup S2$

Intersection

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$S1 \cap S2$

Set Difference

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating	age
22	dustin	7	45.0

$S1 - S2$

Cross-Product (X)

- S1 X R1 : Each row of S1 is paired with each row of R1.
- *Result schema* has one field per field of S1 and R1, with field names 'inherited' if possible.

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

X

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Cross-Product (X)

- $S1 \times R1$: Each row of S1 is paired with each row of R1.
- *Result schema* has one field per field of S1 and R1, with field names 'inherited' if possible.
 - *Conflict*: Both S1 and R1 have a field called *sid*.
((sid), sname, rating, age, (sid), bid, day)
- Renaming operator: $\rho (C(1 \rightarrow sid1, 5 \rightarrow sid2), S1 \times R1)$
(sid1, sname, rating, age, sid2, bid, day)

Joins

- Condition Join: $R \bowtie_c S = \sigma_c (R \times S)$

$$\sigma_{S1.sid < R1.sid} (S1 \times R1)$$

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

X

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Joins

- Condition Join: $R \bowtie_c S = \sigma_c (R \times S)$

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

$S1 \bowtie R1$
 $S1.sid < R1.sid$

S1.sid	sname	rating	age	R1.sid	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

- *Result schema* same as that of cross-product.
- Fewer tuples than cross-product, might be able to compute more efficiently
- Sometimes called a *theta-join*.

Joins

- Equi-Join: A special case of condition join where the condition c contains only *equalities*.

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

$$S1 \bowtie_{sid} R1$$

- *Result schema* similar to cross-product, **but only one copy of fields** for which equality is specified.
- Natural Join: Equijoin on *all* common fields.

$$S1 \bowtie R1$$

Division

- Not supported as a primitive operator, but useful for expressing queries like:

Find sailors who have reserved all boats.

- Let A have 2 fields, x and y ; B have only field y :
 - $A/B = \{ \langle x \rangle \mid \exists \langle x, y \rangle \in A \ \forall \langle y \rangle \in B \}$
 - i.e., A/B contains all x tuples (sailors) such that for every y tuple (boat) in B , there is an xy tuple in A .
 - Or: If the set of y values (boats) associated with an x value (sailor) in A contains all y values in B , the x value is in A/B .
- In general, x and y can be any lists of fields; y is the list of fields in B , and $x \cup y$ is the list of fields of A .

Examples of Division A/B

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

A

pno
p2

B1

sno
s1
s2
s3
s4

A/B1

pno
p2
p4

B2

sno
s1
s4

A/B2

pno
p1
p2
p4

B3

sno
s1

A/B3

Expressing A/B Using Basic Operators

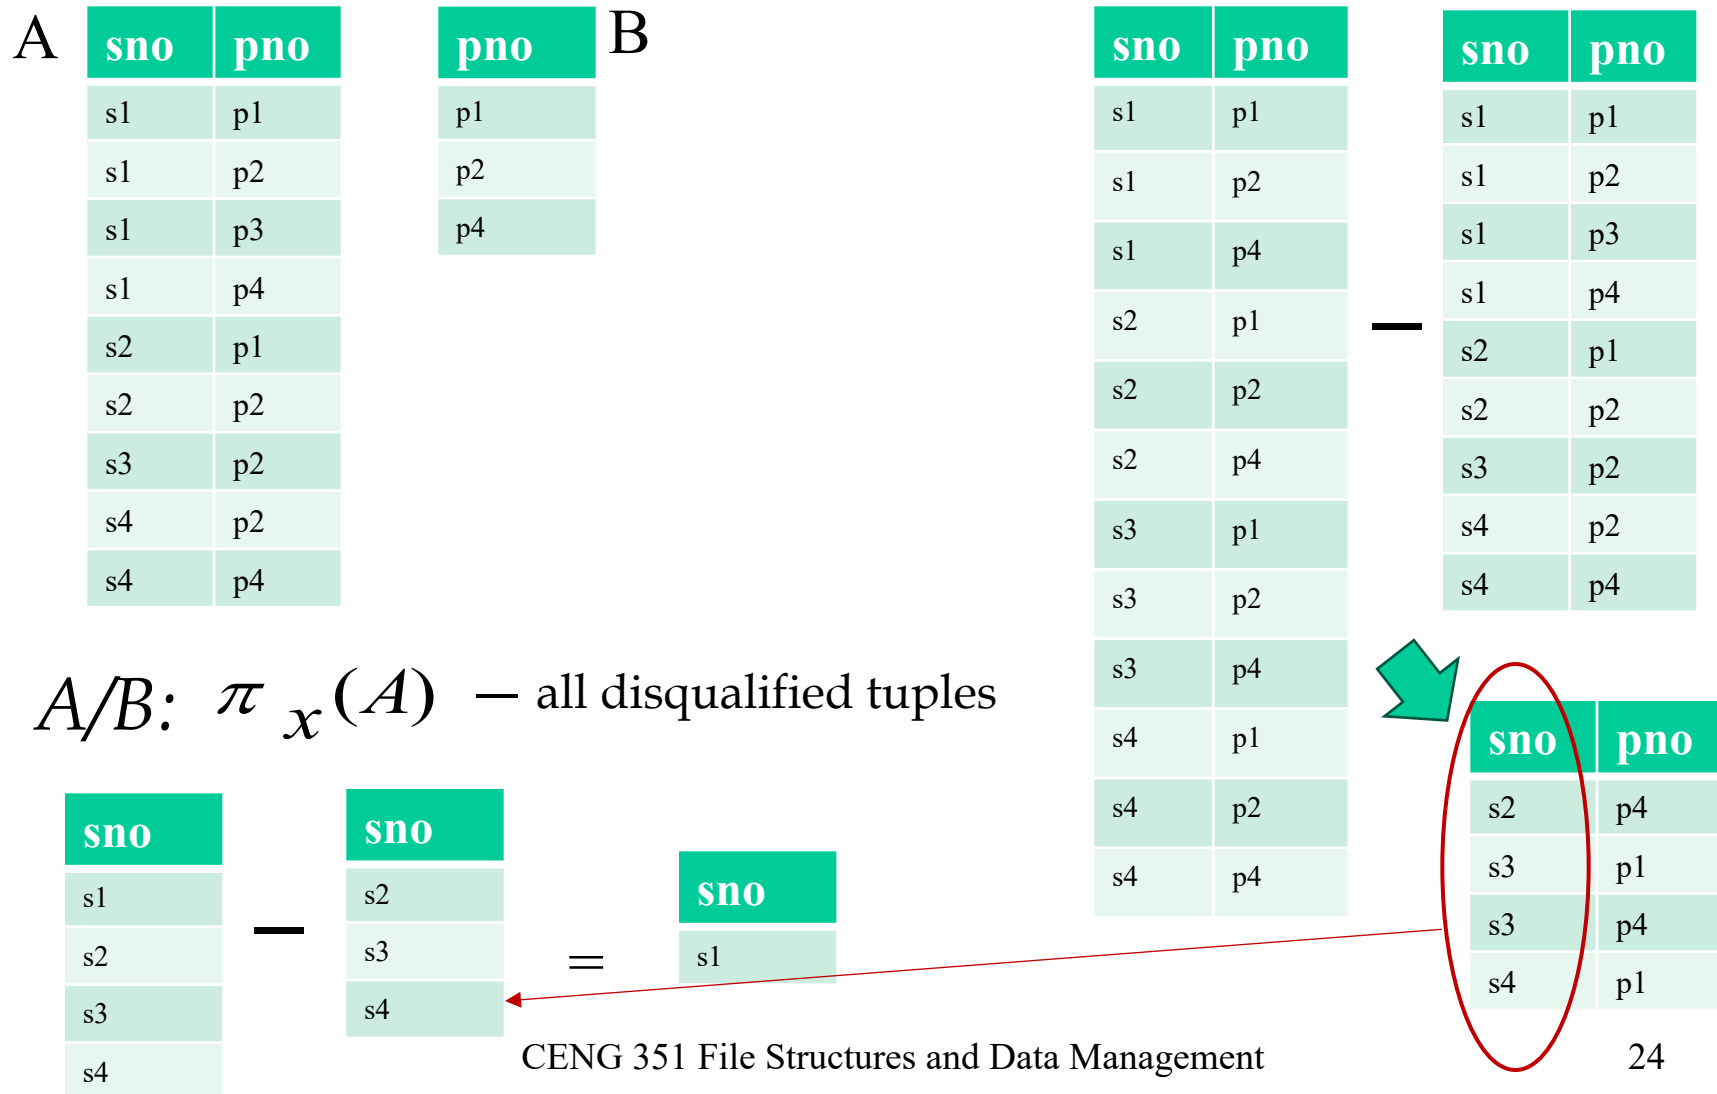
- Division is not an essential op; just a useful shorthand.
 - (Also true of joins, but joins are so common that systems implement joins specially.)
- *Idea:* For A/B , compute all x values that are not 'disqualified' by some y value in B .
 - x value is *disqualified* if by attaching y value from B , we obtain an xy tuple that is not in A .

Disqualified x values: $\pi_x ((\pi_x(A) \times B) - A)$

A/B : $\pi_x(A) -$ all disqualified tuples

Expressing A/B Using Basic Operators

Disqualified x values: $\pi_x((\pi_x(A) \times B) - A)$



<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Figure 4.15 An Instance *S3* of Sailors

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Figure 4.16 An Instance *R2* of Reserves

<i>bid</i>	<i>bname</i>	<i>color</i>
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Figure 4.17 An Instance *B1* of Boats

Find names of sailors who've reserved boat #103

Find names of sailors who've reserved boat #103

➤ Solution 1:

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Find names of sailors who've reserved boat #103

➤ Solution 1:

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Find names of sailors who've reserved boat #103

➤ Solution 1:

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$

<i>sid</i>	<i>bid</i>	<i>day</i>		<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	101	10/10/98		22	Dustin	7	45.0
22	102	10/10/98		29	Brutus	1	33.0
22	103	10/8/98		31	Lubber	8	55.5
22	104	10/7/98		32	Andy	8	25.5
31	102	11/10/98		58	Rusty	10	35.0
31	103	11/6/98		64	Horatio	7	35.0
31	104	11/12/98		71	Zorba	10	16.0
64	101	9/5/98		74	Horatio	9	35.0
64	102	9/8/98		85	Art	3	25.5
74	103	9/8/98		95	Bob	3	63.5

Find names of sailors who've reserved boat #103

➤ Solution 1:

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$

<i>sid</i>	<i>bid</i>	<i>day</i>		<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	101	10/10/98		22	Dustin	7	45.0
22	102	10/10/98		29	Brutus	1	33.0
22	103	10/8/98	→	31	Lubber	8	55.5
22	104	10/7/98		32	Andy	8	25.5
31	102	11/10/98		58	Rusty	10	35.0
31	103	11/6/98	→	64	Horatio	7	35.0
31	104	11/12/98		71	Zorba	10	16.0
64	101	9/5/98		74	Horatio	9	35.0
64	102	9/8/98		85	Art	3	25.5
74	103	9/8/98	→	95	Bob	3	63.5

<i>sid</i>	<i>bid</i>	<i>day</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	103	10/8/98	Dustin	7	45.0
31	103	11/6/98	Lubber	8	55.5
74	103	9/8/98	Horatio	9	35.0

Find names of sailors who've reserved boat #103

➤ Solution 1:

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98



sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

sname

Dustin

Lubber

Horatio

CENG 3

sid	bid	day	sname	rating	age
22	103	10/8/98	Dustin	7	45.0
31	103	11/6/98	Lubber	8	55.5
74	103	9/8/98	Horatio	9	35.0

Find names of sailors who've reserved boat #103

➤ Solution 1:

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$

→ Notice that RA expression described a step-by-step procedure for computing the desired answer.

→ Thus, an RA expression is like a recipe, or a plan, for evaluating a query, and relational systems in fact use algebra expressions to represent *query evaluation plans*.

Find names of sailors who've reserved boat #103

➤ Solution 1:

$$\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$$

➤ Solution 2:

$$\begin{aligned} & \rho(Temp1, \sigma_{bid=103} Reserves) \\ & \rho(Temp2, Temp1 \bowtie Sailors) \\ & \pi_{sname}(Temp2) \end{aligned}$$

➤ Solution 3:

$$\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$$

Find names of sailors who've reserved a red boat

- Information about boat color is only available in Boats; so need an extra join:

Find names of sailors who've reserved a red boat

- Information about boat color is only available in Boats; so need an extra join:

$$\pi_{sname} ((\sigma_{color = 'red'} Boats) \bowtie Reserves \bowtie Sailors)$$

Find names of sailors who've reserved a red boat

- Information about boat color is only available in Boats; so need an extra join:


$$\pi_{sname} ((\sigma_{color = 'red'} Boats) \bowtie Reserves \bowtie Sailors)$$

<i>bid</i>	<i>bname</i>	<i>color</i>
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Find names of sailors who've reserved a red boat

- Information about boat color is only available in Boats; so need an extra join:

$$\pi_{sname} ((\sigma_{color = 'red'} Boats) \bowtie Reserves \bowtie Sailors)$$



Find names of sailors who've reserved a red boat

- Information about boat color is only available in Boats; so need an extra join:

$$\pi_{sname} ((\sigma_{color = 'red'} Boats) \bowtie Reserves \bowtie Sailors)$$

<i>bid</i>	<i>bname</i>	<i>color</i>
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

⋈

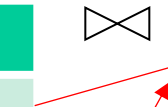
<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

bid	bname	color	sid	day
102	Interlake	red	22	10/10/98
102	Interlake	red	31	11/10/98
102	Interlake	red	64	9/8/98
104	Marine	red	22	10/7/98
194	Marine	red	31	11/12/98

Find names of sailors who've reserved a red boat

- Information about boat color is only available in Boats; so need an extra join:

$$\pi_{sname} ((\sigma_{color = 'red'} Boats) \bowtie Reserves \bowtie Sailors)$$

bid	bname	color	sid	day		sid	sname	rating	age
102	Interlake	red	22	10/10/98		22	Dustin	7	45.0
102	Interlake	red	31	11/10/98		29	Brutus	1	33.0
102	Interlake	red	64	9/8/98		31	Lubber	8	55.5
104	Marine	red	22	10/7/98		32	Andy	8	25.5
194	Marine	red	31	11/12/98		58	Rusty	10	35.0
						64	Horatio	7	35.0
						71	Zorba	10	16.0
						74	Horatio	9	35.0
					85	Art	3	25.5	
					95	Bob	3	63.5	

Find names of sailors who've reserved a red boat

- Information about boat color is only available in Boats; so need an extra join:

$$\pi_{sname} ((\sigma_{color = 'red'} Boats) \bowtie Reserves \bowtie Sailors)$$

bid	bname	color	sid	day
102	Interlake	red	22	10/10/98
102	Interlake	red	31	11/10/98
102	Interlake	red	64	9/8/98
104	Marine	red	22	10/7/98
194	Marine	red	31	11/12/98



sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

bid	bname	color	sid	day	sname	rating	age
102	Interlake	red	22	10/10/98	Dustin	7	45.0
102	Interlake	red	31	11/10/98	Lubber	8	55.5
102	Interlake	red	64	9/8/98	Horatio	7	35.0
104	Marine	red	22	10/7/98	Dustin	7	45.0
194	Marine	red	31	11/12/98	Lubber	8	55.5

Find names of sailors who've reserved a red boat

- Information about boat color is only available in Boats; so need an extra join:

$$\pi_{sname} ((\sigma_{color = 'red'} Boats) \bowtie Reserves \bowtie Sailors)$$

bid	bname	color	sid	day
102	Interlake	red	22	10/10/98
102	Interlake	red	31	11/10/98
102	Interlake	red	64	9/8/98
104	Marine	red	22	10/7/98
194	Marine	red	31	11/12/98



sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

sname
Dustin
Lubber
Horatio



bid	bname	color	sid	day	sname	rating	age
102	Interlake	red	22	10/10/98	Dustin	7	45.0
102	Interlake	red	31	11/10/98	Lubber	8	55.5
102	Interlake	red	64	9/8/98	Horatio	7	35.0
104	Marine	red	22	10/7/98	Dustin	7	45.0
194	Marine	red	31	11/12/98	Lubber	8	55.5

Find names of sailors who've reserved a red boat

- Information about boat color is only available in Boats; so need an extra join:

$$\pi_{sname} ((\sigma_{color = 'red'} Boats) \bowtie Reserves \bowtie Sailors)$$

- A more efficient solution:

$$\pi_{sname} (\pi_{sid} ((\pi_{bid} \sigma_{color = 'red'} Boats) \bowtie Res) \bowtie Sailors)$$

- A query optimizer can find this given the first solution!

Find sailors who've reserved a red or a green boat

- Can identify all red or green boats, then find sailors who've reserved one of these boats:

$$\rho \text{ (Tempboats, } (\sigma_{color='red' \vee color='green'} \text{ Boats}))$$
$$\pi_{sname} \text{ (Tempboats } \bowtie \text{ Reserves } \bowtie \text{ Sailors)}$$

- Can also define Tempboats using union! (How?)
- What happens if \vee is replaced by \wedge in this query?

Find sailors who've reserved a red and a green boat

- Previous approach won't work! Must identify sailors who've reserved red boats, sailors who've reserved green boats, then find the intersection (note that *sid* is a key for Sailors):

$$\begin{aligned} & \rho (Tempred, \pi_{sid} ((\sigma_{color = 'red', Boats}) \bowtie Reserves)) \\ & \rho (Tempgreen, \pi_{sid} ((\sigma_{color = green, Boats}) \bowtie Reserves)) \\ & \pi_{sname} ((Tempred \cap Tempgreen) \bowtie Sailors) \end{aligned}$$

Find the names of sailors who've reserved all boats

- Uses division; schemas of the input relations to / must be carefully chosen:

$$\rho (Tempsids, (\pi_{sid, bid} Reserves) / (\pi_{bid} Boats))$$

$$\pi_{sname} (Tempsids \bowtie Sailors)$$

- To find sailors who've reserved all 'Interlake' boats:

$$/ \pi_{bid} (\sigma_{bname='Interlake'} Boats)$$

.....

Find sailors (sids) who've reserved at least 2 boats

Reserves

sid	bid
1	101
1	102
2	102

R1

sid	bid
1	101
1	102
2	102

R2

sid	bid
1	101
1	102
2	102

Find sailors (sids) who've reserved at least 2 boats

R1	sid	bid
	1	101
	1	102
	2	102

R2	sid	bid
	1	101
	1	102
	2	102

R1.sid	R1.bid	R2.sid	R2.bid
1	101	1	101
1	101	1	102
1	101	2	102

Find sailors (sids) who've reserved at least 2 boats

R1

sid	bid
1	101
1	102
2	102

R2

sid	bid
1	101
1	102
2	102

R1.sid	R1.bid	R2.sid	R2.bid
1	101	1	101
1	101	1	102
1	101	2	102
1	102	1	101
1	102	1	102
1	102	2	102
2	102	1	101
2	102	1	102
2	102	2	102

Find sailors (sids) who've reserved at least 2 boats

R1

sid	bid
1	101
1	102
2	102

R2

sid	bid
1	101
1	102
2	102

R1.sid	R1.bid	R2.sid	R2.bid
1	101	1	101
1	101	1	102
1	102	1	101
1	102	1	102
2	102	2	102

- STUDENT (ID, Name, Major, Bdate)
- COURSE (Course_number, Cname, Dept)
- ENROLLED (ID, Course_number, Quarter, Grade)
- BOOK_ADOPTION (Course_number, Quarter, Book_isbn)
- TEXT (Book_isbn, Book_title, Publisher, Author)

List the departments, which have all of their adopted books published by ‘Pearson Publishing’.

- STUDENT (ID, Name, Major, Bdate)
- COURSE (Course_number, Cname, Dept)
- ENROLLED (ID, Course_number, Quarter, Grade)
- BOOK_ADOPTION (Course_number, Quarter, Book_isbn)
- TEXT (Book_isbn, Book_title, Publisher, Author)

Provide a list of textbooks (include Course_number, Book_isbn, Book_title) for courses offered by the ‘CS’ department that have used more than one book.